

**Name:** Deep Pawar(A20545137)  
**Professor:** Joseph Rosen  
**Institute:** Illinois Institute of Technology

# CSP 554: Big Data Technologies

Fall 2024 - Assignment 2

- Questions and Answers:**

e. Create an S3 bucket of whatever name you choose. For purposes of this discussion I will assume it is called mybucket (but of course your bucket must have a name unique across all of AWS)

f. Using the techniques we learned during assignment #1 upload A12345678.txt to the S3 bucket into an object called A12345678.txt (of course use your id). Note throughout the rest of this document we will still call this object by the name “myid.txt”.

**Ans:**

- Bucket Name:** a20545137—csp554
- Object Name:** A20545137.txt

The screenshot displays the AWS Management Console interface for an S3 upload. At the top, a green banner confirms 'Upload succeeded'. Below this, the 'Upload: status' section provides a summary of the upload process. A table shows the upload details: 1 file (21.0 B) was successfully uploaded to the destination s3://a20545137--csp554. The 'Files and folders' section shows a table with one entry: A20545137.txt, which is a text/plain file of 21.0 B, successfully uploaded.

Name	Folder	Type	Size	Status	Error
A20545137.txt	-	text/plain	21.0 B	Succeeded	-



```
hadoop@ip-172-31-24-248:~  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$ hadoop fs -mkdir /user/csp554-2  
[hadoop@ip-172-31-24-248 ~]$ hadoop fs -ls /user  
Found 8 items  
drwxr-xr-x   - hadoop hdfsadmingroup    0 2024-09-09 18:20 /user/csp554  
drwxr-xr-x   - hadoop hdfsadmingroup    0 2024-09-09 18:22 /user/csp554-2  
drwxrwxrwx   - hadoop hdfsadmingroup    0 2024-09-09 17:18 /user/hadoop  
drwxr-xr-x   - mapred mapred            0 2024-09-09 17:18 /user/history  
drwxrwxrwx   - hdfs hdfsadmingroup      0 2024-09-09 17:18 /user/hive  
drwxrwxrwx   - hue hue                 0 2024-09-09 17:18 /user/hue  
drwxrwxrwx   - oozie oozie              0 2024-09-09 17:20 /user/oozie  
drwxrwxrwx   - root hdfsadmingroup      0 2024-09-09 17:18 /user/root  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$
```

[illegible]

15. (2 points) Copy the object myid.txt you uploaded to an S3 bucket into the Hadoop master node Linux file system. The actual object includes your student id as above.

Note, Amazon EMR and Hadoop provide a variety of file systems that you can use with EMR. You specify which file system to use with a file system prefix. For example, `s3://myawsbucket/path` references an Amazon S3 bucket using EMRFS (EMR file system). See: <https://docs.aws.amazon.com/emr/latest/ManagementGuide/emr-plan-file-systems.html>

The way you do this would be as follows to copy an object from an S3 bucket to the Linux file system of the Hadoop master node.

```
aws s3 cp s3://mybucket/myid.txt /home/hadoop/myid.txt
```

The above is an AWS CLI (command line interpreter) command. For more information about how to use the CLI to manipulate S3 buckets see: <https://docs.aws.amazon.com/cli/latest/reference/s3/index.html>

After you executed the above command perform an “ls /home/hadoop” and take a screen snapshot of names of the files or directories that are listed and include it in your assignment submission.

**Ans:**

- **Command Used:**

- i. To copy an object from an S3 bucket to the Linux file system of the Hadoop master node:** `aws s3 cp s3://a20545137--csp554/A20545137.txt /home/hadoop/A20545137.txt`
- ii. To list the files and directories:** `ls /home/hadoop/`

[illegible]

[illegible]



18. (2 points) Execute a command to remove the myid.txt file in the hdfs directory /user/csp554-2

Clue: look up about how to use the “rm” command in the file system shell document.

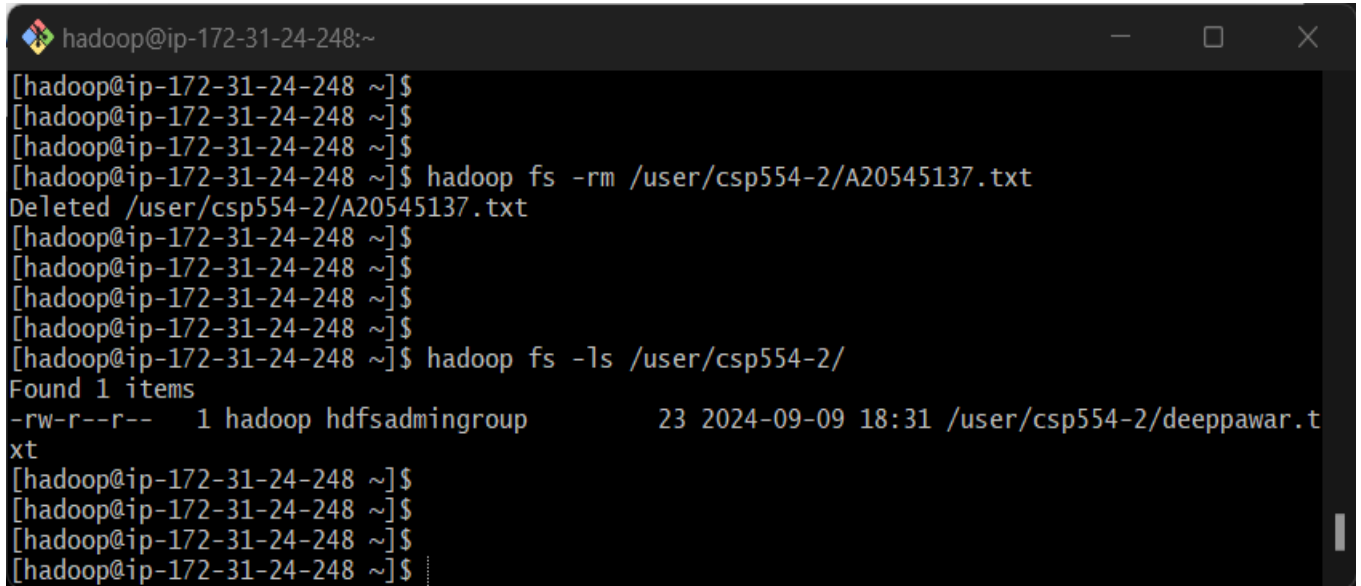
Write down the command you executed, then list the content of the /user/csp554-2 HDFS directory and take a screen snapshot of the listed content of the directory and include it in your assignment submission.

**Ans:**

- **Command Used:**

- i. **To remove the myid.txt (A20545137.txt) file in the hdfs directory:** `hadoop fs -rm /user/csp554-2/A20545137.txt`

- ii. **To list the files and directories:** `hadoop fs -ls /user/csp554-2/`

A terminal window with a dark background and light text. The window title is 'hadoop@ip-172-31-24-248:~'. The terminal shows a series of commands and their outputs. The first command is 'hadoop fs -rm /user/csp554-2/A20545137.txt', which results in 'Deleted /user/csp554-2/A20545137.txt'. The second command is 'hadoop fs -ls /user/csp554-2/', which results in 'Found 1 items' followed by a line of file details: '-rw-r--r-- 1 hadoop hdfsadmin group 23 2024-09-09 18:31 /user/csp554-2/deeppawar.txt'. The terminal ends with several empty prompts.

```
hadoop@ip-172-31-24-248:~$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$ hadoop fs -rm /user/csp554-2/A20545137.txt  
Deleted /user/csp554-2/A20545137.txt  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$ hadoop fs -ls /user/csp554-2/  
Found 1 items  
-rw-r--r-- 1 hadoop hdfsadmin group 23 2024-09-09 18:31 /user/csp554-2/deeppawar.t  
xt  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$  
[hadoop@ip-172-31-24-248 ~]$
```