Name: Deep Pawar(A20545137)
Professor: Joseph Rosen
Institute: Illinois Institute of Technology

CSP 554: Big Data Technologies

Fall 2024 - Assignment 3

- **Questions and Answers:**

4) Create a new EMR cluster the same as you did previously. Since you already have a security key (".pem" or ".cer" file) just use that one during cluster creation. Or, if you deleted your security key, just create a new one.

Ans:

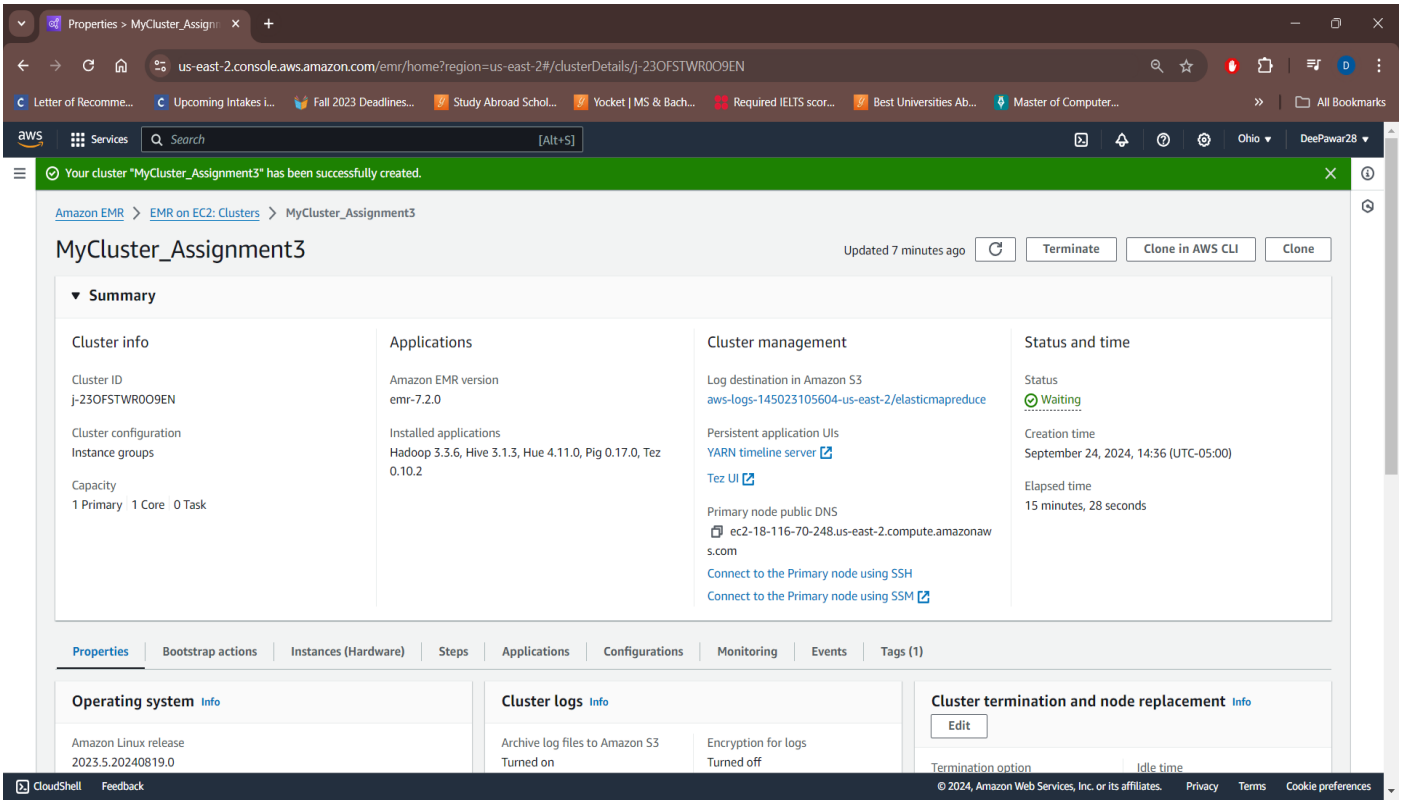

5) Install the mrjob library on your EMR primary node.

a) ssh to the primary node (/home/hadoop) as you did in assignment #2

Ans:

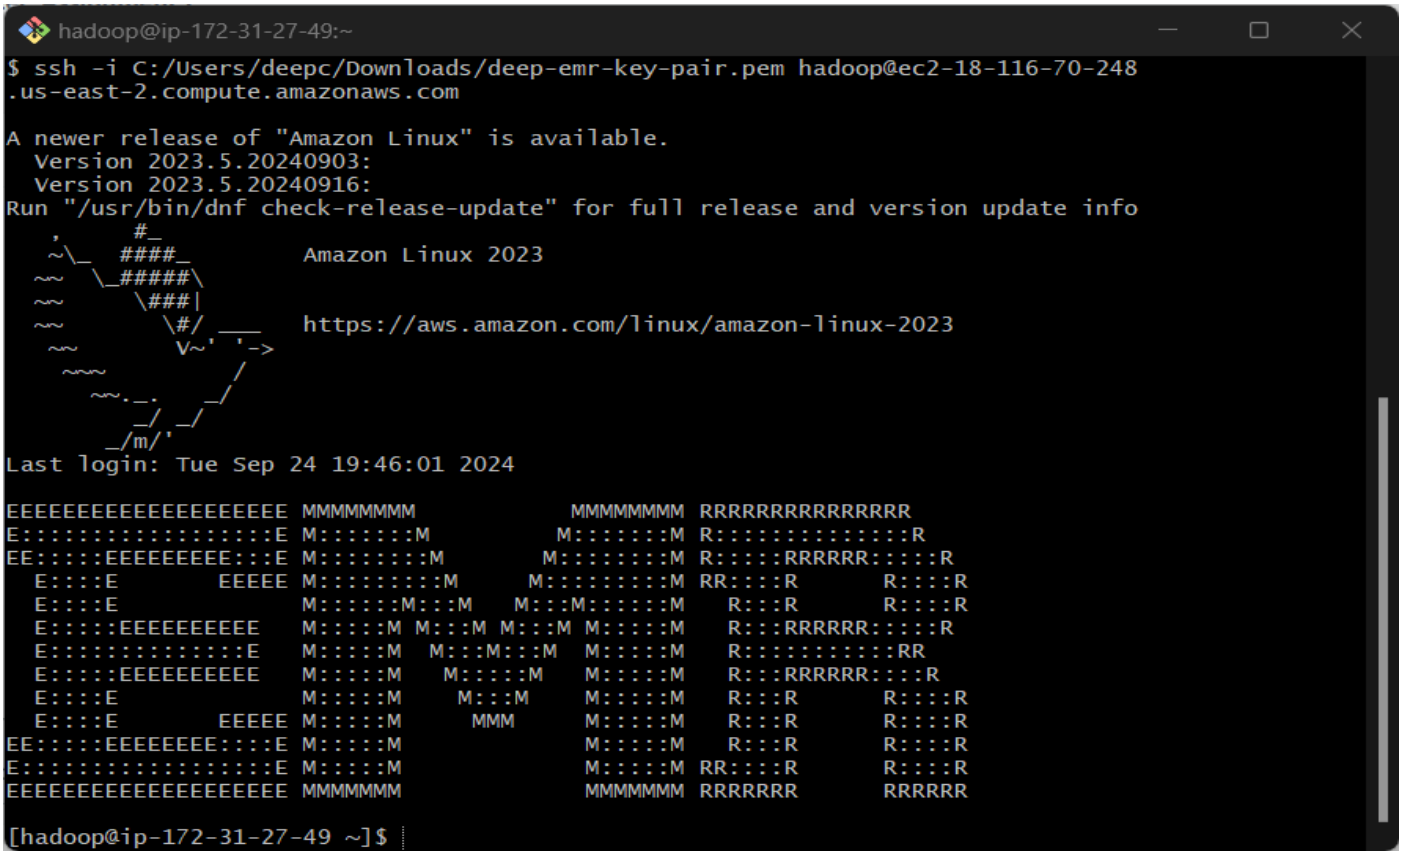

b) Enter the following (note if the first command does not work, try the second)

Ans:

- **Command:** sudo /usr/bin/pip3 install mrjob[aws]

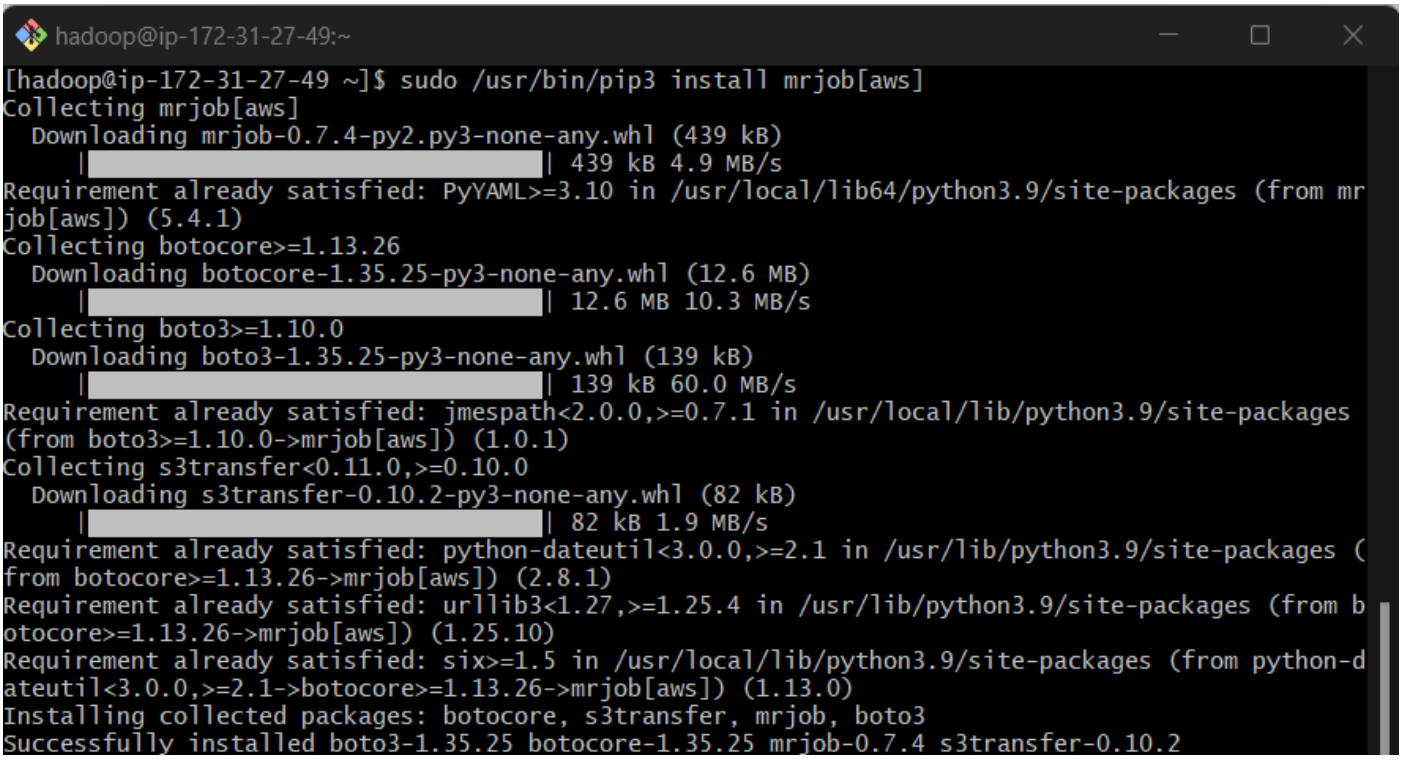

6) Next you will set up to execute the provided WordCount.py map reduce program found in the "Assignments" section of the Blackboard. This is the exact same program we saw in class.

- Step 1:

Download the two files "w.data" and "WordCount.py" to your PC or Mac. They are part of the documents included with the assignment.

- Step 2:

Ans:

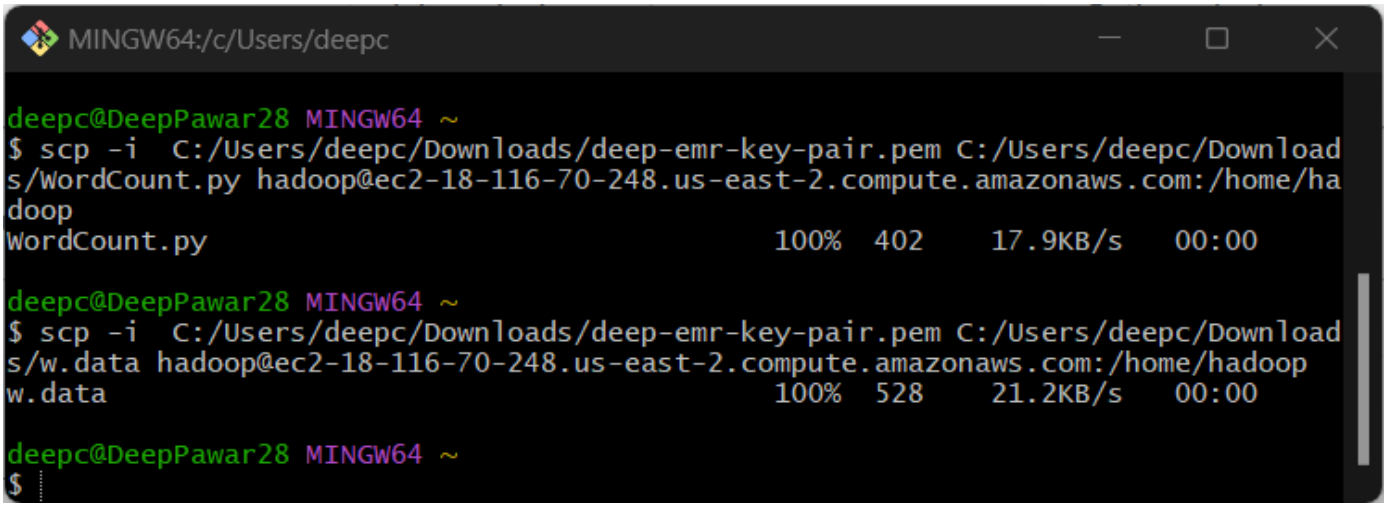

- Step 3:

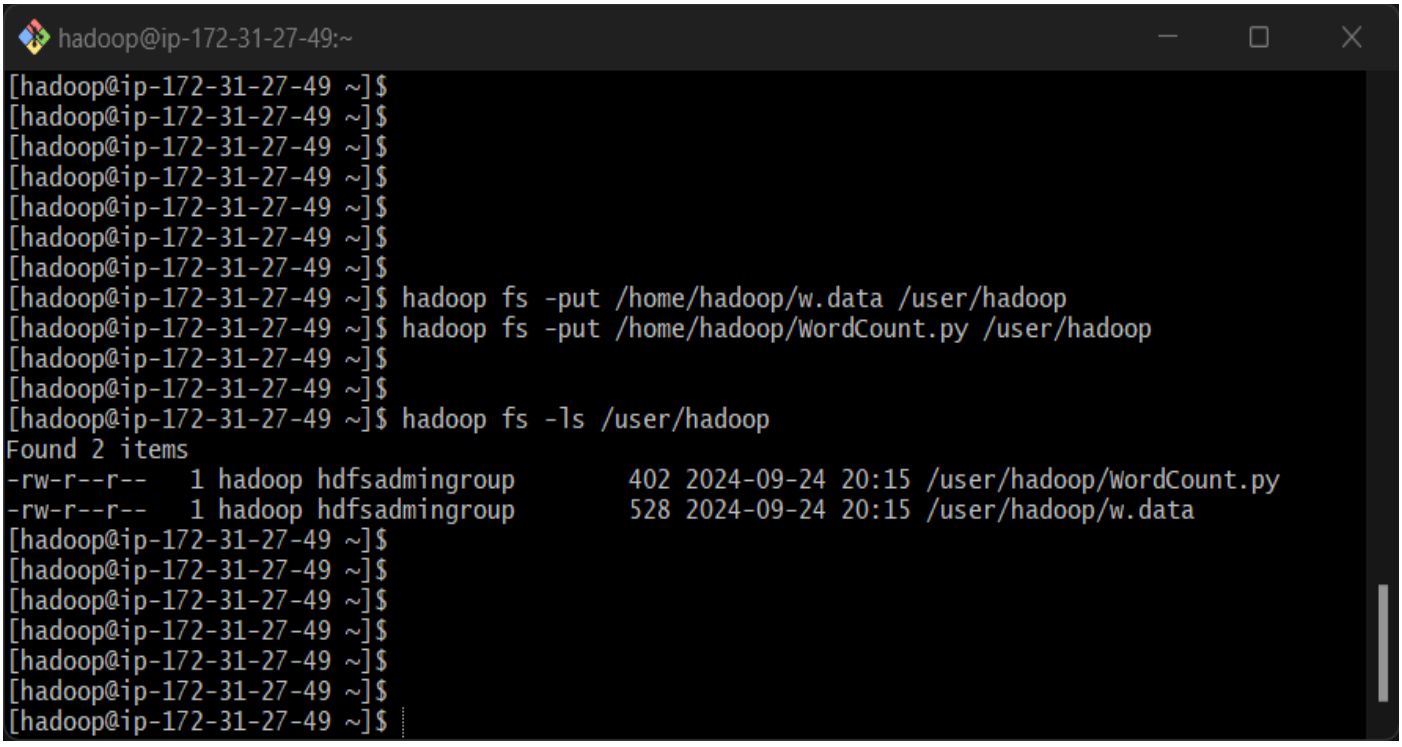

- Step 4:

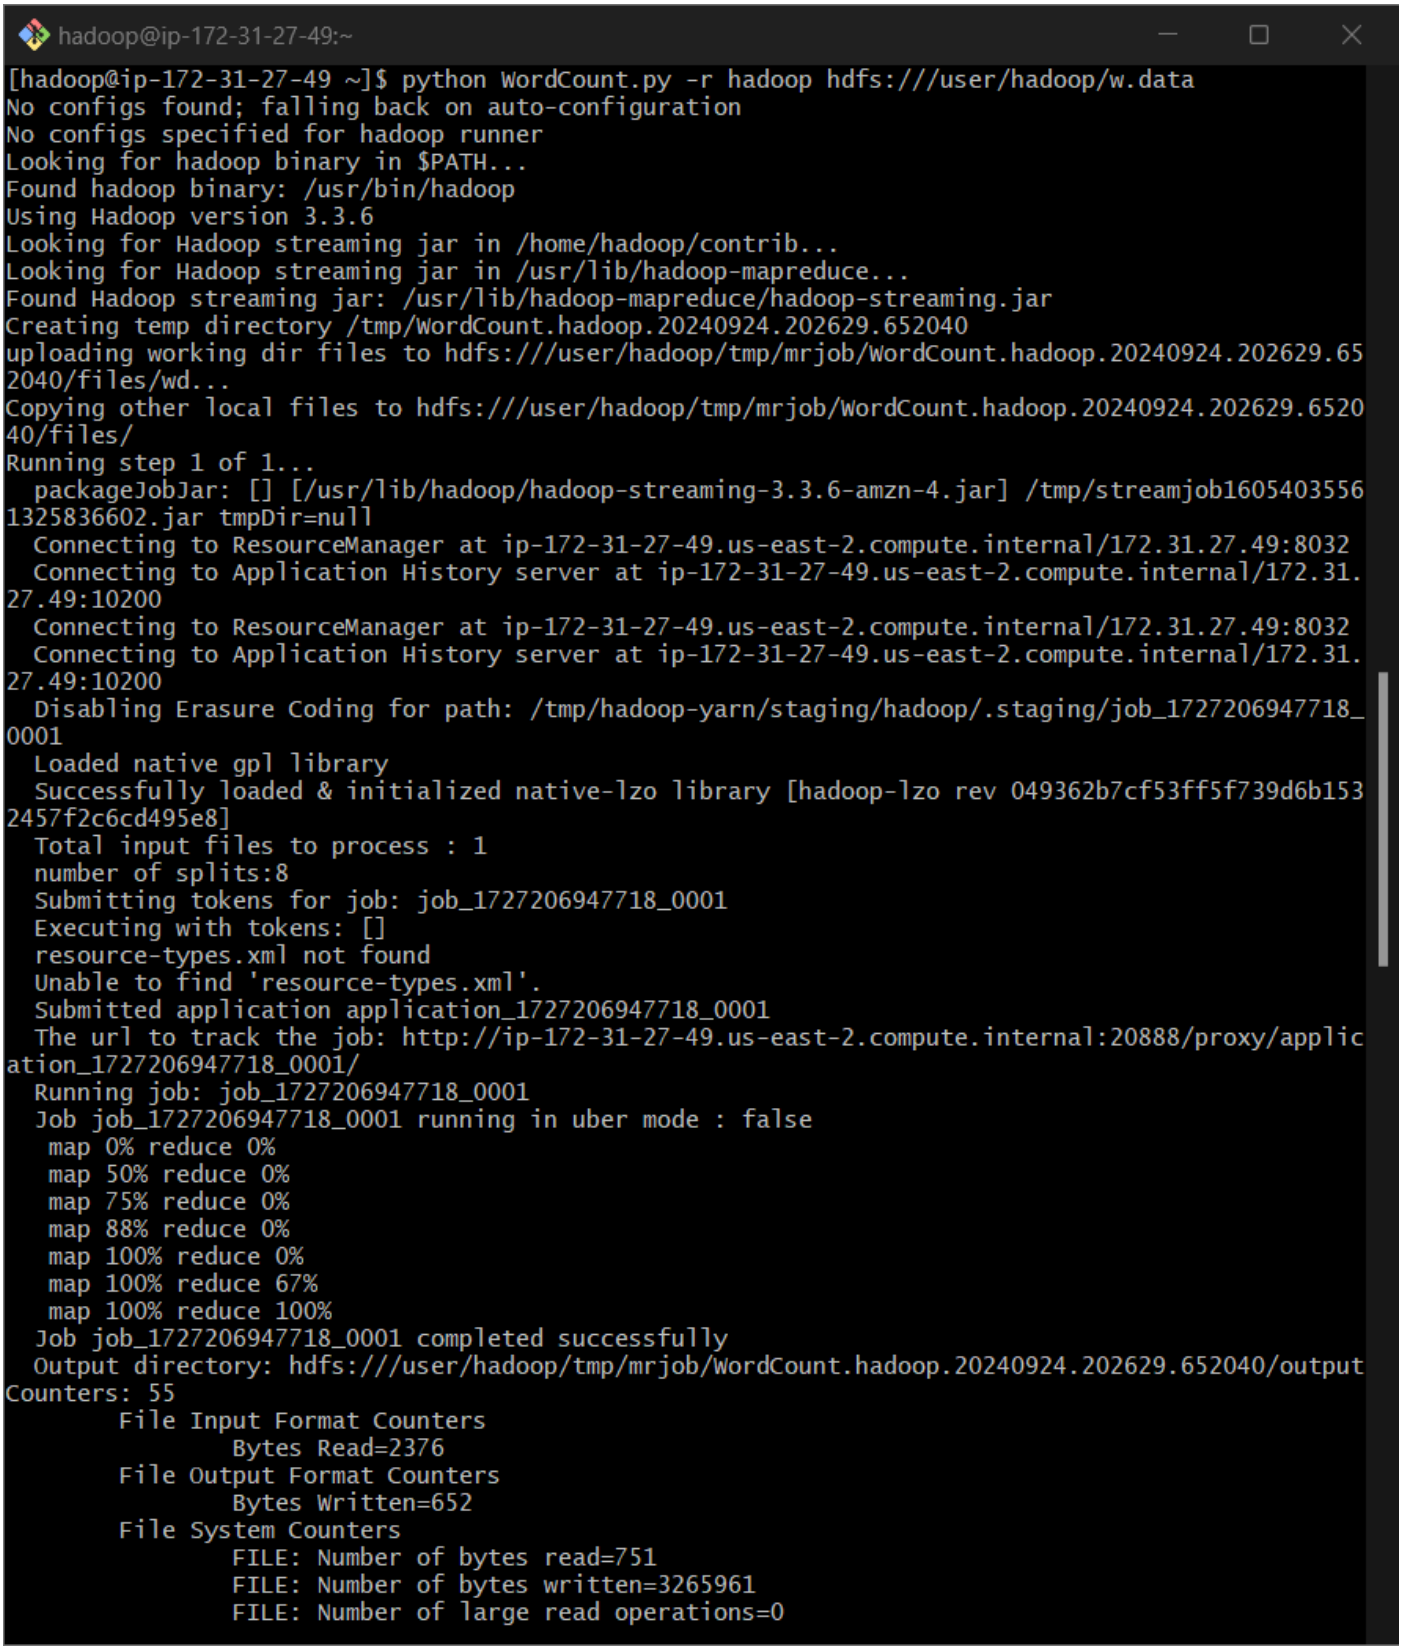
hadoop@ip-172-31-27-49:~
[hadoop@ip-172-31-27-49 ~]$ python WordCount.py -r hadoop hdfs:///user/hadoop/w.data
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 3.3.6
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/WordCount.hadoop.20240924.202629.652040
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20240924.202629.652040/files/wd...
Copying other local files to hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20240924.202629.652040/files/
Running step 1 of 1...
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6-amzn-4.jar] /tmp/streamjob16054035561325836602.jar tmpDir=null
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1727206947718_0001
Loaded native gpl library
Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b1532457f2c6cd495e8]
Total input files to process : 1
number of splits:8
Submitting tokens for job: job_1727206947718_0001
Executing with tokens: []
resource-types.xml not found
Unable to find 'resource-types.xml'.
Submitted application application_1727206947718_0001
The url to track the job: http://ip-172-31-27-49.us-east-2.compute.internal:20888/proxy/application_1727206947718_0001/
Running job: job_1727206947718_0001
Job job_1727206947718_0001 running in uber mode : false
map 0% reduce 0%
map 50% reduce 0%
map 75% reduce 0%
map 88% reduce 0%
map 100% reduce 0%
map 100% reduce 67%
map 100% reduce 100%
Job job_1727206947718_0001 completed successfully
Output directory: hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20240924.202629.652040/output
Counters: 55
File Input Format Counters
Bytes Read=2376
File Output Format Counters
Bytes Written=652
File System Counters
FILE: Number of bytes read=751
FILE: Number of bytes written=3265961
FILE: Number of large read operations=0


hadoop@ip-172-31-27-49:~
FILE: Number of read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=3376
HDFS: Number of bytes read erasure-coded=0
HDFS: Number of bytes written=652
HDFS: Number of large read operations=0
HDFS: Number of read operations=39
HDFS: Number of write operations=6
Job Counters
Data-local map tasks=8
Killed map tasks=1
Launched map tasks=8
Launched reduce tasks=3
Total megabyte-milliseconds taken by all map tasks=183937536
Total megabyte-milliseconds taken by all reduce tasks=76323840
Total time spent by all map tasks (ms)=119751
Total time spent by all maps in occupied slots (ms)=5748048
Total time spent by all reduce tasks (ms)=24845
Total time spent by all reduces in occupied slots (ms)=2385120
Total vcore-milliseconds taken by all map tasks=119751
Total vcore-milliseconds taken by all reduce tasks=24845
Map-Reduce Framework
CPU time spent (ms)=13020
Combine input records=95
Combine output records=80
Failed Shuffles=0
GC time elapsed (ms)=1011
Input split bytes=1000
Map input records=6
Map output bytes=891
Map output materialized bytes=1215
Map output records=95
Merged Map outputs=24
Peak Map Physical memory (bytes)=484982784
Peak Map Virtual memory (bytes)=3204657152
Peak Reduce Physical memory (bytes)=321642496
Peak Reduce Virtual memory (bytes)=4551507968
Physical memory (bytes) snapshot=4700872704
Reduce input groups=65
Reduce input records=80
Reduce output records=65
Reduce shuffle bytes=1215
Shuffled Maps =24
Spilled Records=160
Total committed heap usage (bytes)=4175429632
Virtual memory (bytes) snapshot=39248273408
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20240924.202629.652040/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20240924.202629.65204
0/output...
"an" 1

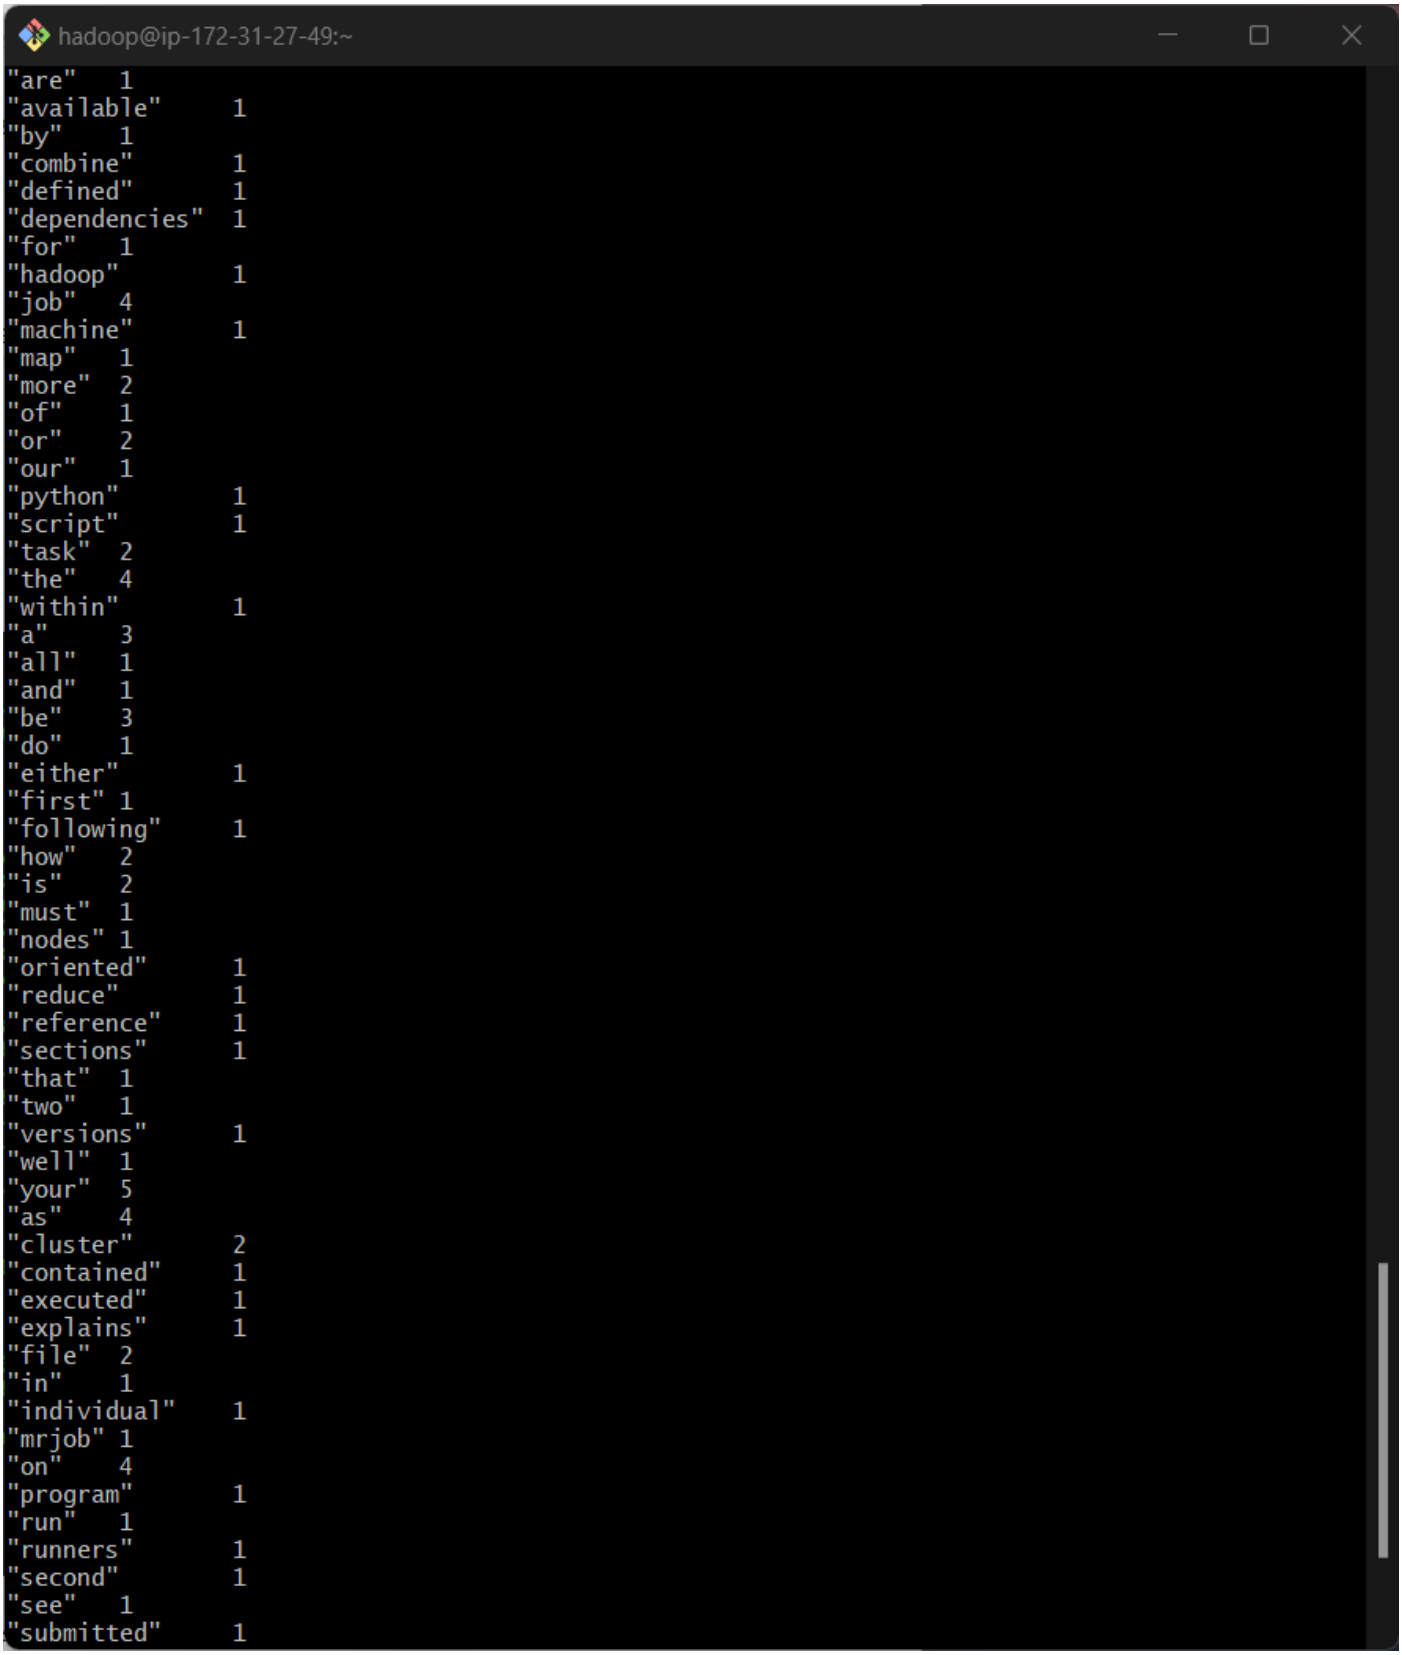
hadoop@ip-172-31-27-49:~
"are" 1
"available" 1
"by" 1
"combine" 1
"defined" 1
"dependencies" 1
"for" 1
"hadoop" 1
"job" 4
"machine" 1
"map" 1
"more" 2
"of" 1
"or" 2
"our" 1
"python" 1
"script" 1
"task" 2
"the" 4
"within" 1
"a" 3
"all" 1
"and" 1
"be" 3
"do" 1
"either" 1
"first" 1
"following" 1
"how" 2
"is" 2
"must" 1
"nodes" 1
"oriented" 1
"reduce" 1
"reference" 1
"sections" 1
"that" 1
"two" 1
"versions" 1
"well" 1
"your" 5
"as" 4
"cluster" 2
"contained" 1
"executed" 1
"explains" 1
"file" 2
"in" 1
"individual" 1
"mrjob" 1
"on" 4
"program" 1
"run" 1
"runners" 1
"second" 1
"see" 1
"submitted" 1


hadoop@ip-172-31-27-49:~
"things" 1
"those" 1
"to" 3
"uploaded" 1
"when" 1
"will" 1
"writing" 2
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount.hadoop.20240924.202629.652040...
Removing temp directory /tmp/WordCount.hadoop.20240924.202629.652040...
[hadoop@ip-172-31-27-49 ~]$

5) Now slightly modify the WordCount.py program. Call the new program WordCount2.py.

Instead of counting how many words there are in the input documents (w.data), modify the program to count how many words begin with the lower-case letters a-n (a through n inclusive) and how many begin with anything else.

The output file should look something like

a_to_n, 12

other, 21

Note, do not force words to all lower case. Now execute the program and see what happens.

6) (3 points) Submit (1) a copy of this modified program and (2) a screen shot of the results of the program's execution as the output of your assignment.

- **Code: WordCount2.py**

WordCount2.py

File Edit View

```
from mrjob.job import MRJob
import re

WORD_RE = re.compile(r"[\w']+")

class MRWordCount2(MRJob):

    def mapper(self, _, line):
        for word in WORD_RE.findall(line):
            a_to_n = "abcdefghijklmn"

            if any(word[0] ==w for w in a_to_n):
                yield "a_to_n", 1
            else:
                yield 'other', 1

    def combiner(self, word, counts):
        yield word, sum(counts)

    def reducer(self, word, counts):
        yield word, sum(counts)

if __name__ == '__main__':
    MRWordCount2.run()
```

- **Output:**

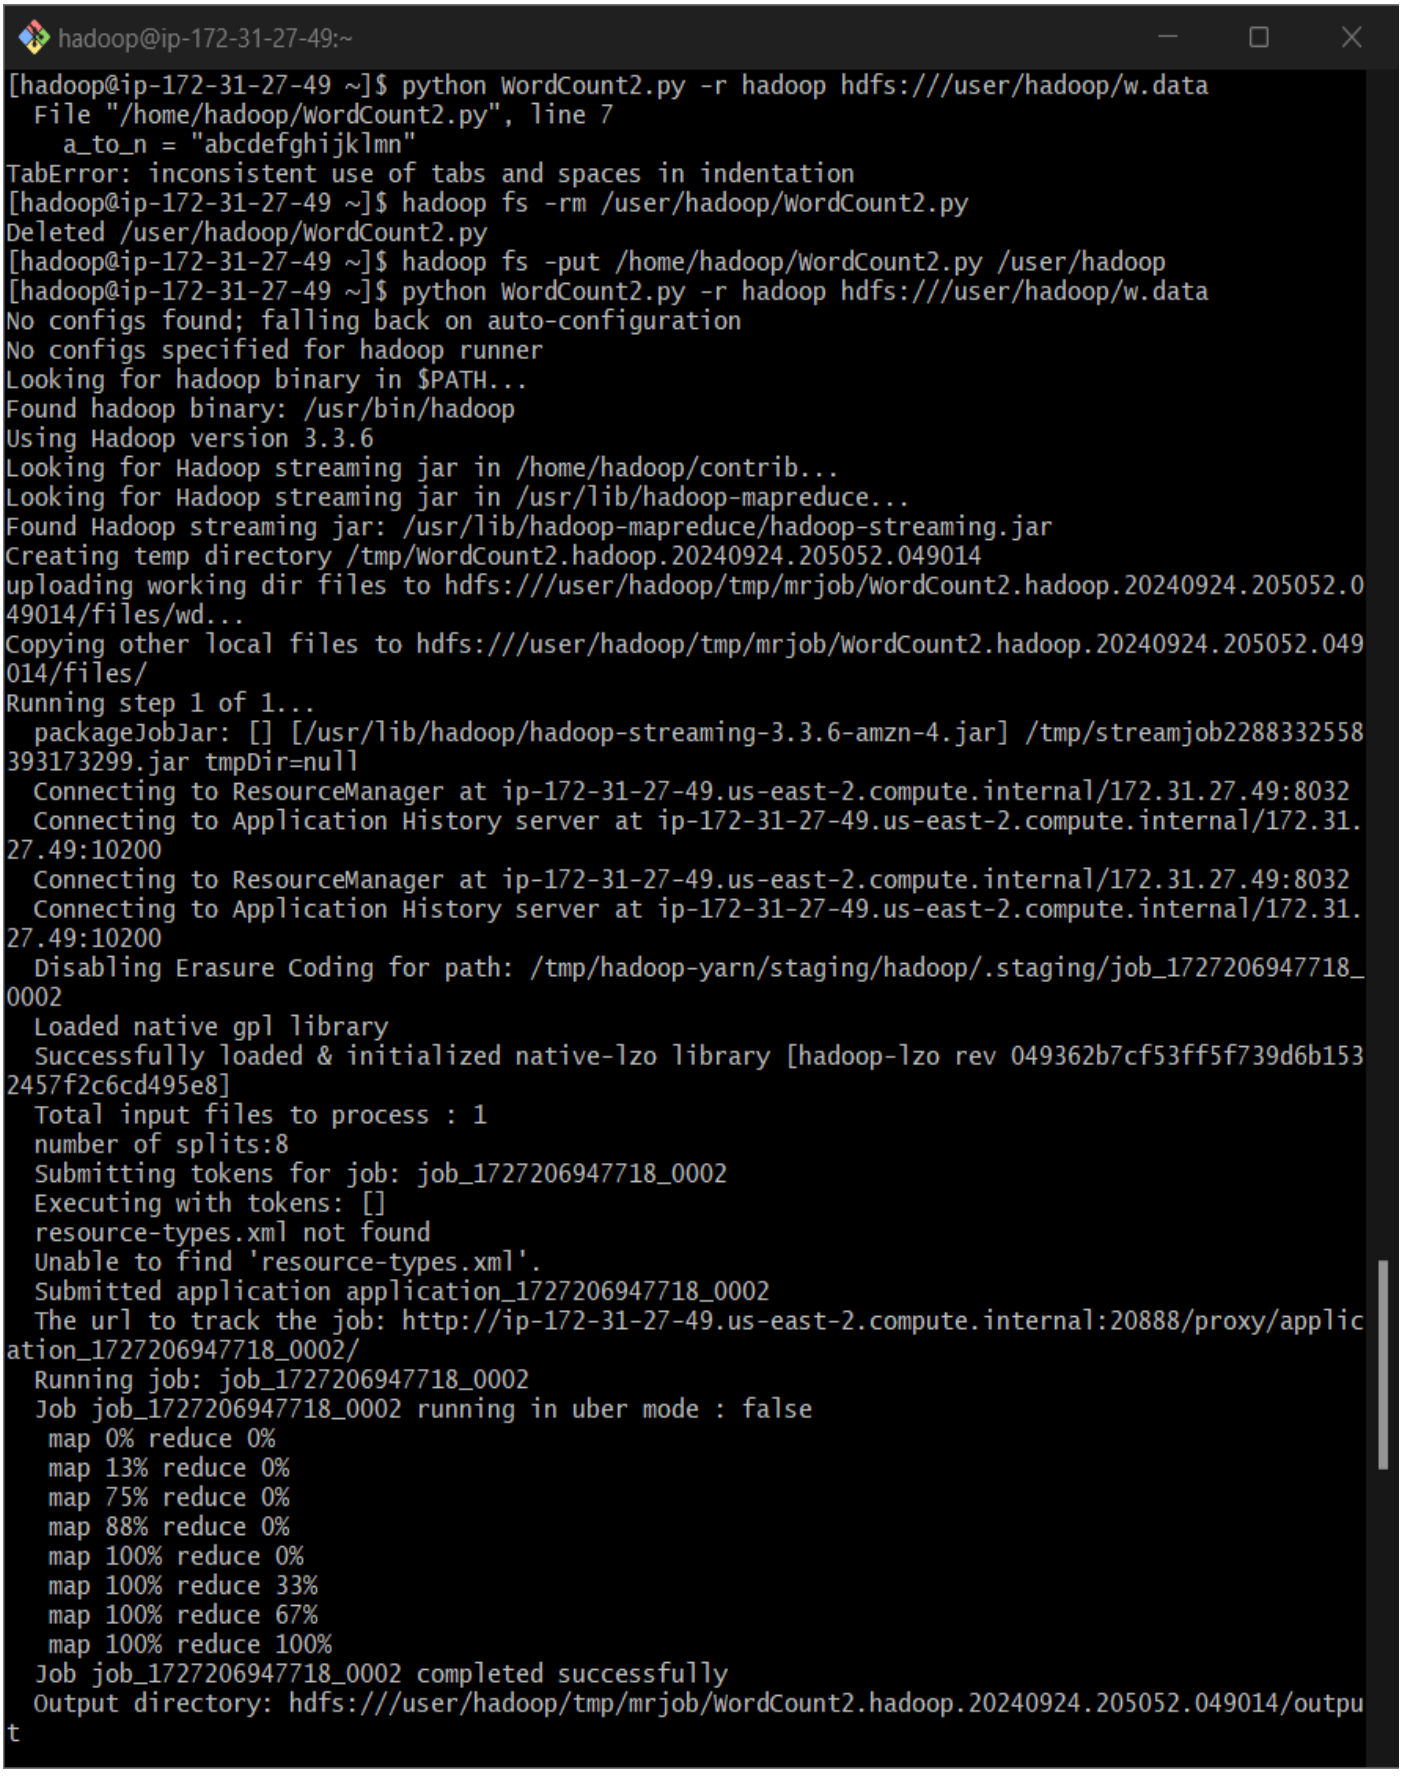
hadoop@ip-172-31-27-49:~
[hadoop@ip-172-31-27-49 ~]$ python WordCount2.py -r hadoop hdfs:///user/hadoop/w.data
File "/home/hadoop/WordCount2.py", line 7
a_to_n = "abcdefghijklmn"
TabError: inconsistent use of tabs and spaces in indentation
[hadoop@ip-172-31-27-49 ~]$ hadoop fs -rm /user/hadoop/WordCount2.py
Deleted /user/hadoop/WordCount2.py
[hadoop@ip-172-31-27-49 ~]$ hadoop fs -put /home/hadoop/WordCount2.py /user/hadoop
[hadoop@ip-172-31-27-49 ~]$ python WordCount2.py -r hadoop hdfs:///user/hadoop/w.data
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 3.3.6
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/WordCount2.hadoop.20240924.205052.049014
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20240924.205052.049014/files/wd...
Copying other local files to hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20240924.205052.049014/files/
Running step 1 of 1...
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6-amzn-4.jar] /tmp/streamjob2288332558393173299.jar tmpDir=null
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1727206947718_0002
Loaded native gpl library
Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b1532457f2c6cd495e8]
Total input files to process : 1
number of splits:8
Submitting tokens for job: job_1727206947718_0002
Executing with tokens: []
resource-types.xml not found
Unable to find 'resource-types.xml'.
Submitted application application_1727206947718_0002
The url to track the job: http://ip-172-31-27-49.us-east-2.compute.internal:20888/proxy/application_1727206947718_0002/
Running job: job_1727206947718_0002
Job job_1727206947718_0002 running in uber mode : false
map 0% reduce 0%
map 13% reduce 0%
map 75% reduce 0%
map 88% reduce 0%
map 100% reduce 0%
map 100% reduce 33%
map 100% reduce 67%
map 100% reduce 100%
Job job_1727206947718_0002 completed successfully
Output directory: hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20240924.205052.049014/output

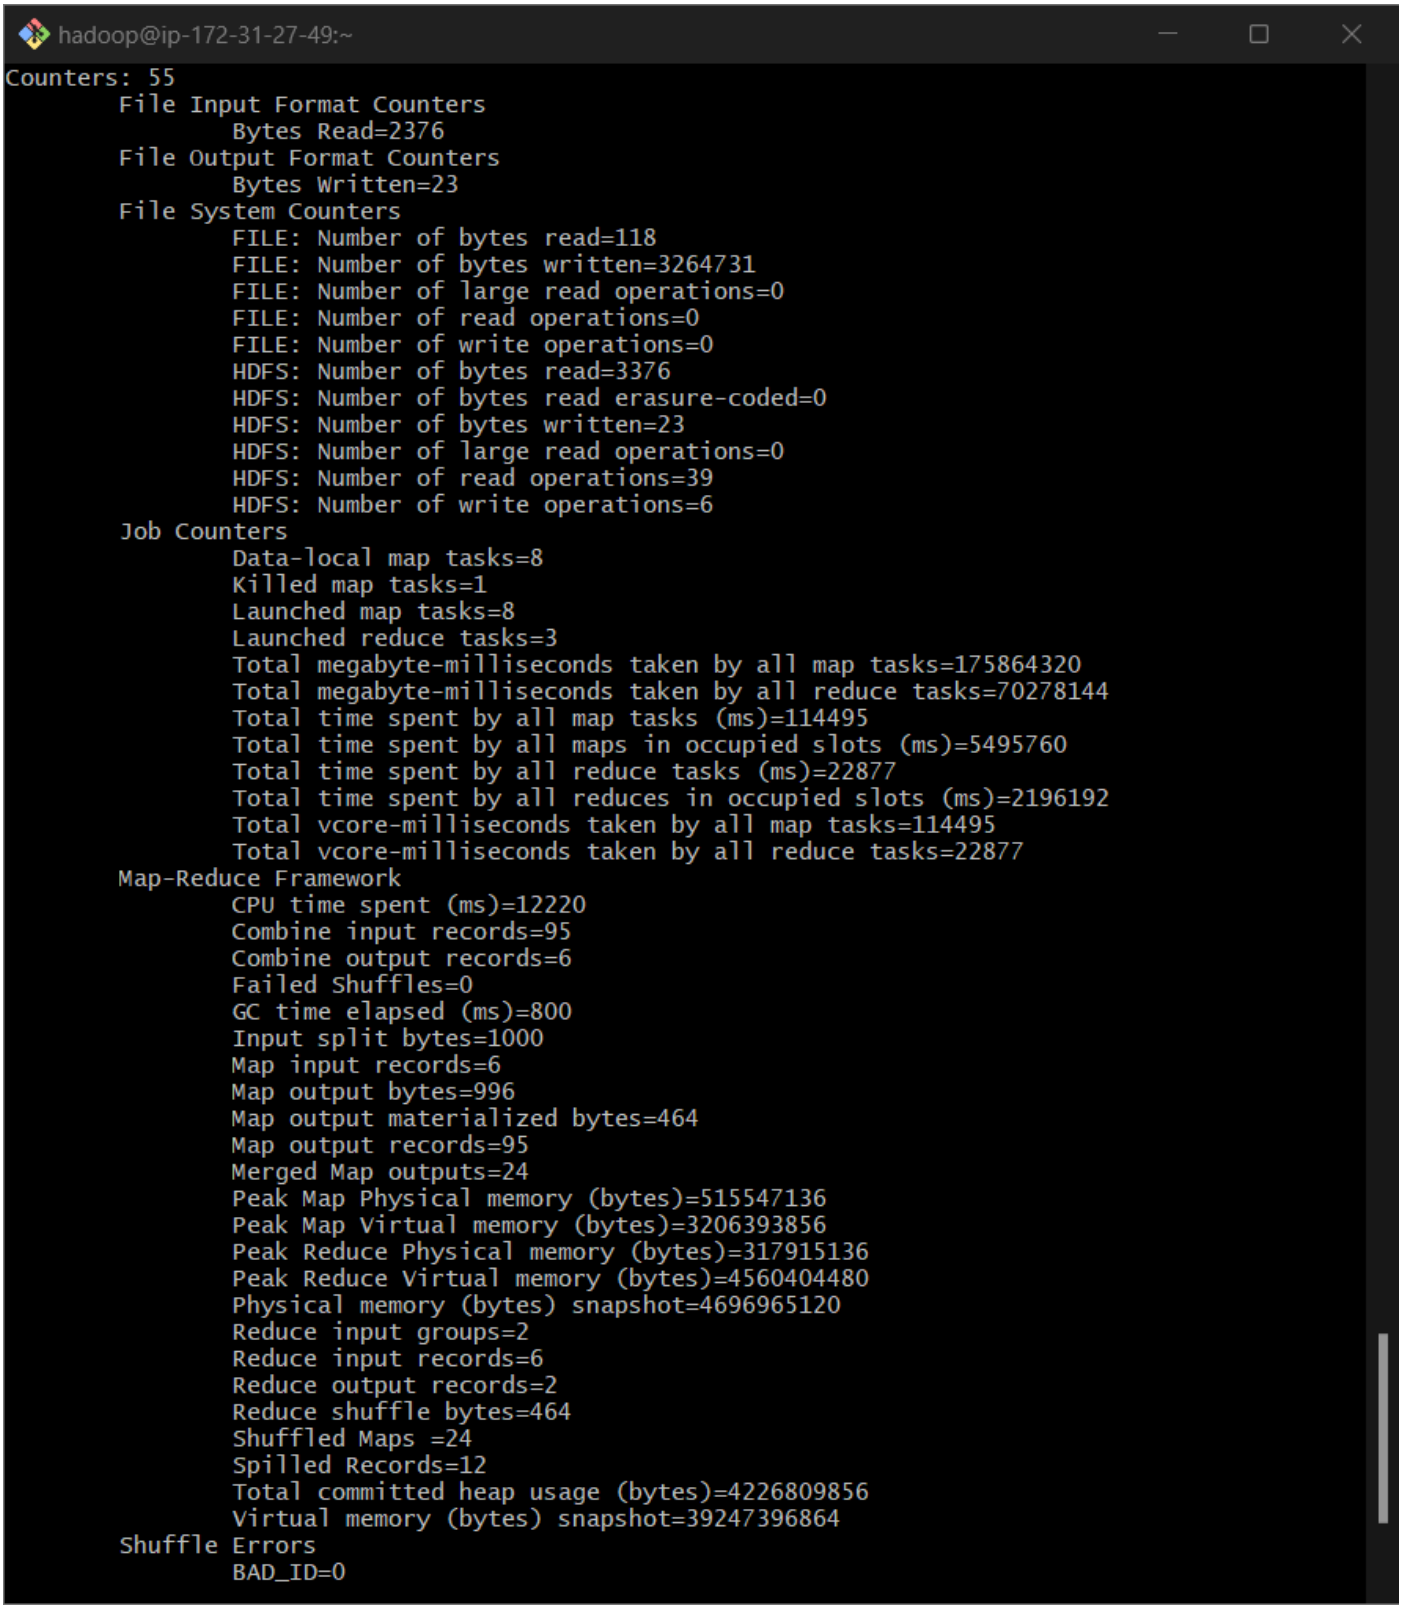
hadoop@ip-172-31-27-49:~
Counters: 55
File Input Format Counters
Bytes Read=2376
File Output Format Counters
Bytes Written=23
File System Counters
FILE: Number of bytes read=118
FILE: Number of bytes written=3264731
FILE: Number of large read operations=0
FILE: Number of read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=3376
HDFS: Number of bytes read erasure-coded=0
HDFS: Number of bytes written=23
HDFS: Number of large read operations=0
HDFS: Number of read operations=39
HDFS: Number of write operations=6
Job Counters
Data-local map tasks=8
Killed map tasks=1
Launched map tasks=8
Launched reduce tasks=3
Total megabyte-milliseconds taken by all map tasks=175864320
Total megabyte-milliseconds taken by all reduce tasks=70278144
Total time spent by all map tasks (ms)=114495
Total time spent by all maps in occupied slots (ms)=5495760
Total time spent by all reduce tasks (ms)=22877
Total time spent by all reduces in occupied slots (ms)=2196192
Total vcore-milliseconds taken by all map tasks=114495
Total vcore-milliseconds taken by all reduce tasks=22877
Map-Reduce Framework
CPU time spent (ms)=12220
Combine input records=95
Combine output records=6
Failed Shuffles=0
GC time elapsed (ms)=800
Input split bytes=1000
Map input records=6
Map output bytes=996
Map output materialized bytes=464
Map output records=95
Merged Map outputs=24
Peak Map Physical memory (bytes)=515547136
Peak Map Virtual memory (bytes)=3206393856
Peak Reduce Physical memory (bytes)=317915136
Peak Reduce Virtual memory (bytes)=4560404480
Physical memory (bytes) snapshot=4696965120
Reduce input groups=2
Reduce input records=6
Reduce output records=2
Reduce shuffle bytes=464
Shuffled Maps =24
Spilled Records=12
Total committed heap usage (bytes)=4226809856
Virtual memory (bytes) snapshot=39247396864
Shuffle Errors
BAD_ID=0
hadoop@ip-172-31-27-49:~
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20240924.205052.049014/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20240924.205052.049014/output...
"a_to_n" 46
"other" 49
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount2.hadoop.20240924.205052.049014...
Removing temp directory /tmp/WordCount2.hadoop.20240924.205052.049014...
[hadoop@ip-172-31-27-49 ~]$

7) Let's modify the WordCount.py program again. Call the new program WordCount3.py.

Instead of counting words, calculate the count of words having the same number of letters. For example, if we have a file consisting of one record of the form:

hello there joe

our job should output key value pairs similar to the following:

3, 1

5, 2

Hint, the key in a key-value pair can be an integer just as well as a string.

So, your task is to write a MrJob MapReduce program which again accepts the following file as input

hdfs:///user/hadoop/w.data

and outputs key value pairs where each one has a key with is some number of characters, and the value a count of words having that many characters. Note, please convert all words to lower case on input, so "Hello" and "hello" become the same word.

8) (4 points) When you have accomplished this, please submit the following, (1) a copy of your MRJob code and (2) a copy of the output of the execution of that code.

Ans:

- **Code: WordCount3.py**

WordCount3.py

File Edit View

```
from mrjob.job import MRJob
import re

WORD_RE = re.compile(r"[\w']+")

class MRWordCount3(MRJob):

    def mapper(self, _, line):
        for word in WORD_RE.findall(line):
            # Yield the length of the word in lowercase
            yield len(word.lower()), 1

    def combiner(self, word_len, counts):
        yield word_len, sum(counts)

    def reducer(self, word_len, counts):
        yield word_len, sum(counts)

if __name__ == '__main__':
    MRWordCount3.run()
```

- **Output:**

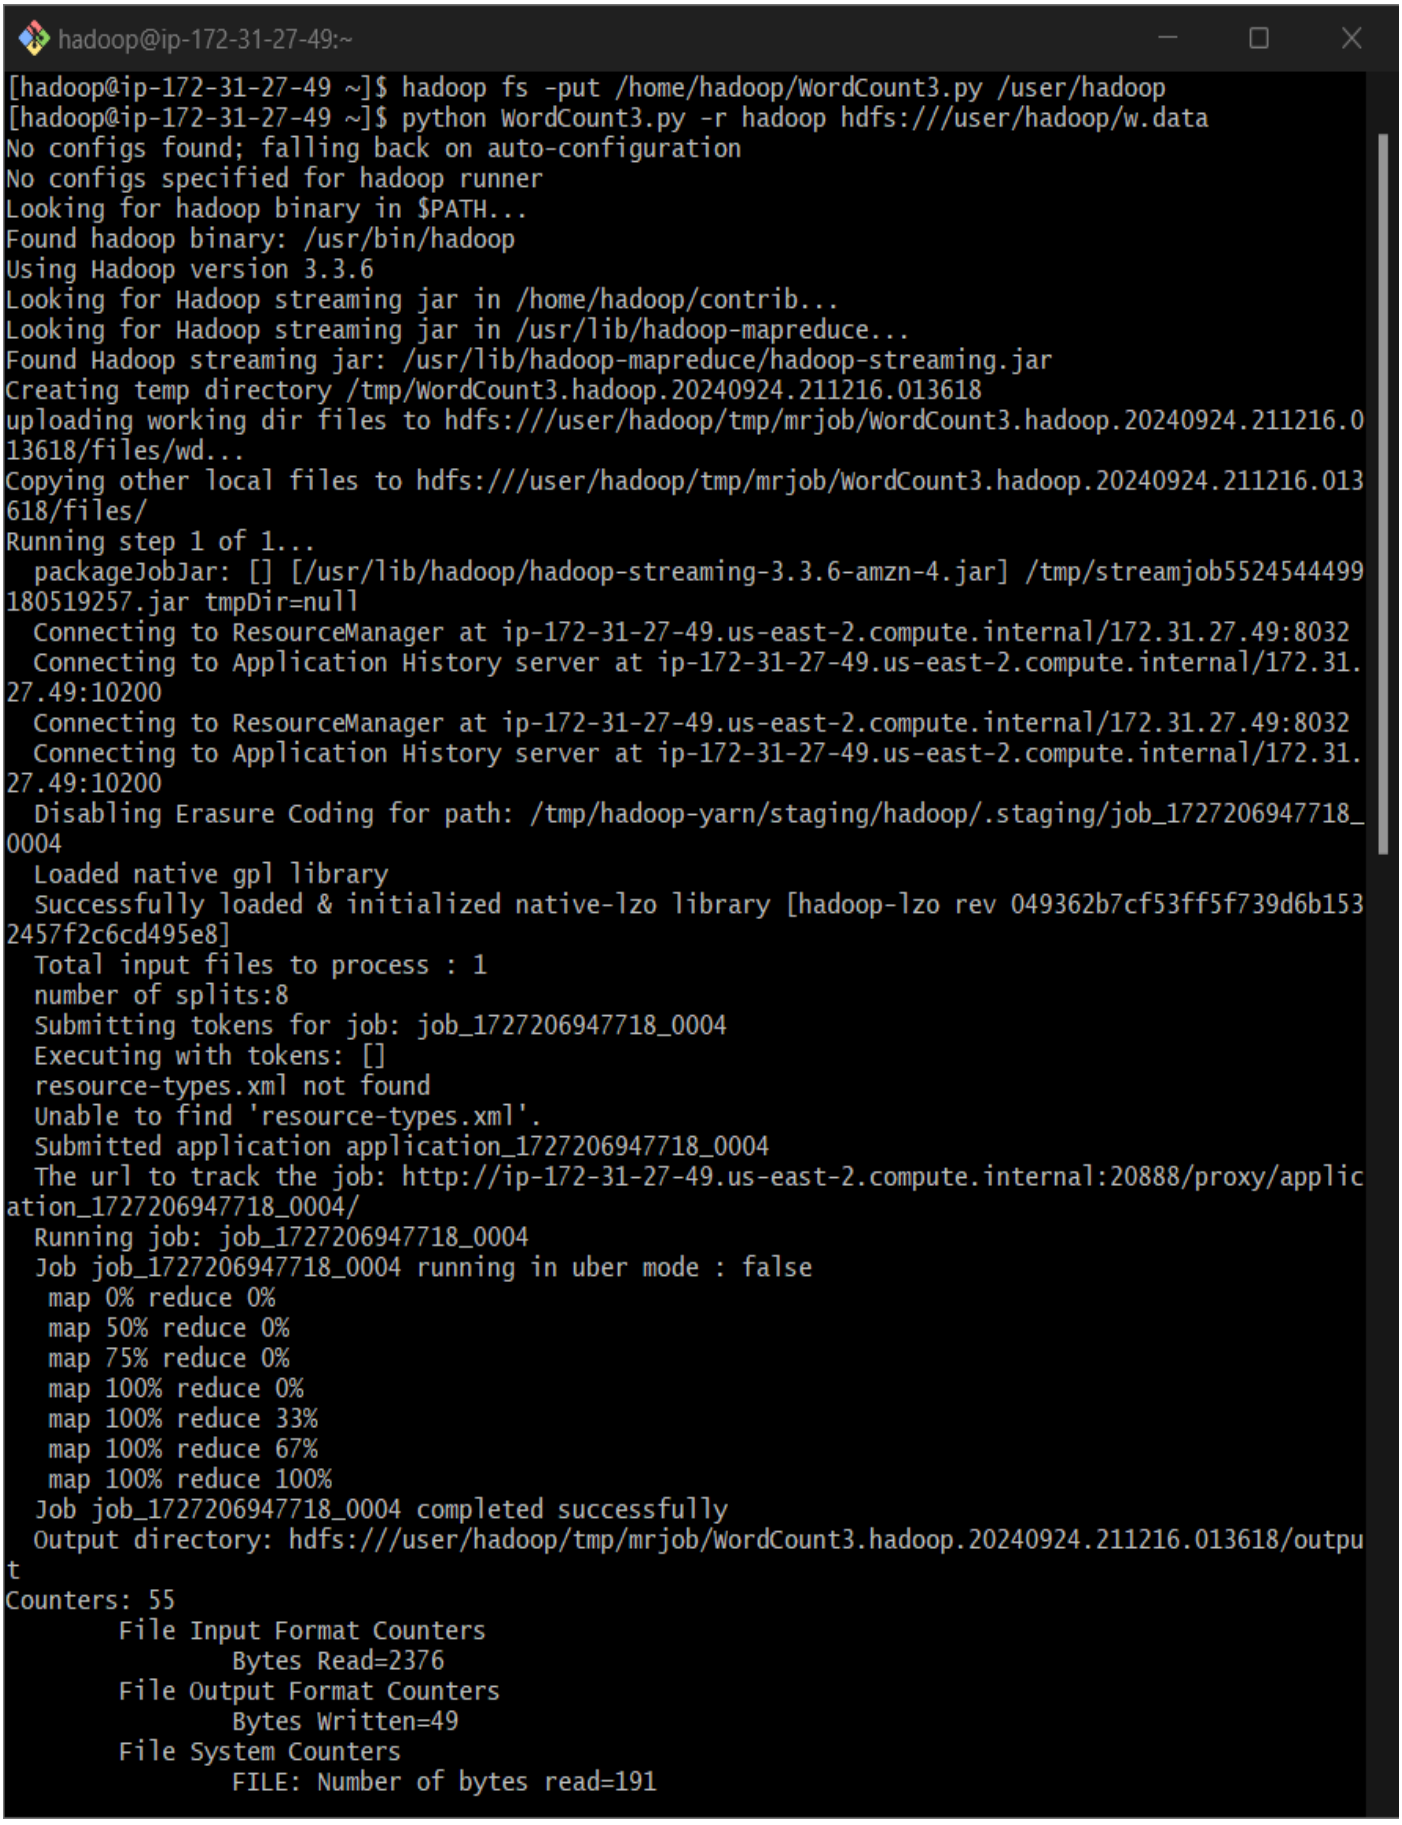
hadoop@ip-172-31-27-49:~
[hadoop@ip-172-31-27-49 ~]$ hadoop fs -put /home/hadoop/WordCount3.py /user/hadoop
[hadoop@ip-172-31-27-49 ~]$ python WordCount3.py -r hadoop hdfs:///user/hadoop/w.data
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 3.3.6
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/WordCount3.hadoop.20240924.211216.013618
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/WordCount3.hadoop.20240924.211216.013618/files/wd...
Copying other local files to hdfs:///user/hadoop/tmp/mrjob/WordCount3.hadoop.20240924.211216.013618/files/
Running step 1 of 1...
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6-amzn-4.jar] /tmp/streamjob5524544499180519257.jar tmpDir=null
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1727206947718_0004
Loaded native gpl library
Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b1532457f2c6cd495e8]
Total input files to process : 1
number of splits:8
Submitting tokens for job: job_1727206947718_0004
Executing with tokens: []
resource-types.xml not found
Unable to find 'resource-types.xml'.
Submitted application application_1727206947718_0004
The url to track the job: http://ip-172-31-27-49.us-east-2.compute.internal:20888/proxy/application_1727206947718_0004/
Running job: job_1727206947718_0004
Job job_1727206947718_0004 running in uber mode : false
map 0% reduce 0%
map 50% reduce 0%
map 75% reduce 0%
map 100% reduce 0%
map 100% reduce 33%
map 100% reduce 67%
map 100% reduce 100%
Job job_1727206947718_0004 completed successfully
Output directory: hdfs:///user/hadoop/tmp/mrjob/WordCount3.hadoop.20240924.211216.013618/output
Counters: 55
File Input Format Counters
Bytes Read=2376
File Output Format Counters
Bytes Written=49
File System Counters
FILE: Number of bytes read=191


hadoop@ip-172-31-27-49:~
File System Counters
FILE: Number of bytes read=191
FILE: Number of bytes written=3264860
FILE: Number of large read operations=0
FILE: Number of read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=3376
HDFS: Number of bytes read erasure-coded=0
HDFS: Number of bytes written=49
HDFS: Number of large read operations=0
HDFS: Number of read operations=39
HDFS: Number of write operations=6
Job Counters
Data-local map tasks=8
Killed map tasks=1
Launched map tasks=8
Launched reduce tasks=3
Total megabyte-milliseconds taken by all map tasks=179152896
Total megabyte-milliseconds taken by all reduce tasks=67974144
Total time spent by all map tasks (ms)=116636
Total time spent by all maps in occupied slots (ms)=5598528
Total time spent by all reduce tasks (ms)=22127
Total time spent by all reduces in occupied slots (ms)=2124192
Total vcore-milliseconds taken by all map tasks=116636
Total vcore-milliseconds taken by all reduce tasks=22127
Map-Reduce Framework
CPU time spent (ms)=12770
Combine input records=95
Combine output records=25
Failed Shuffles=0
GC time elapsed (ms)=700
Input split bytes=1000
Map input records=6
Map output bytes=382
Map output materialized bytes=537
Map output records=95
Merged Map outputs=24
Peak Map Physical memory (bytes)=523509760
Peak Map Virtual memory (bytes)=3215515648
Peak Reduce Physical memory (bytes)=320057344
Peak Reduce Virtual memory (bytes)=4551786496
Physical memory (bytes) snapshot=4808130560
Reduce input groups=11
Reduce input records=25
Reduce output records=11
Reduce shuffle bytes=537
Shuffled Maps =24
Spilled Records=50
Total committed heap usage (bytes)=4284481536
Virtual memory (bytes) snapshot=39276015616
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

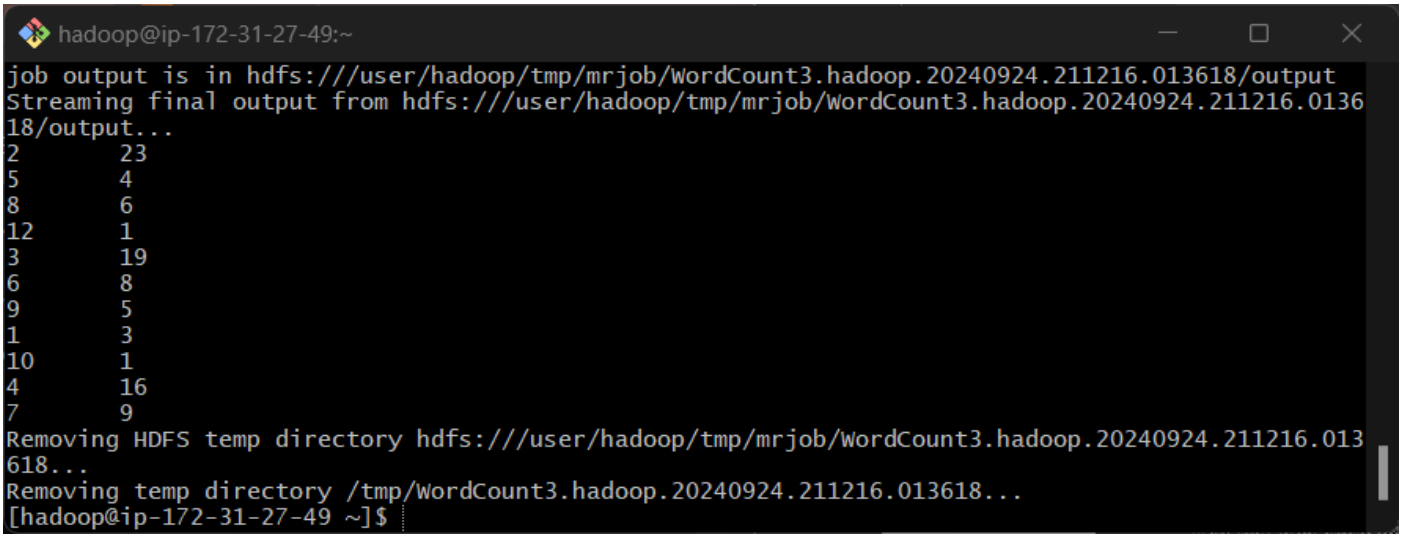
hadoop@ip-172-31-27-49:~
job output is in hdfs:///user/hadoop/tmp/mrjob/WordCount3.hadoop.20240924.211216.013618/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/WordCount3.hadoop.20240924.211216.013618/output...
2 23
5 4
8 6
12 1
3 19
6 8
9 5
1 3
10 1
4 16
7 9
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount3.hadoop.20240924.211216.013618...
Removing temp directory /tmp/WordCount3.hadoop.20240924.211216.013618...
[hadoop@ip-172-31-27-49 ~]$

9) Again, modify the WordCount.py program. Call the new program WordCount4.py.

Now we will write a MapReduce job to calculate the count of unique per record word bigrams. A word bigram is a two word sequence. For example, if we have a file consisting of records of the form:

hello there joe
hi there
there joe go
joe

Bigrams for these records are create by sliding a two word "window" across the words of the record.

hello there joe => "hello there", "there joe"
hi there => "hi there"
there joe there => "there joe", "joe there"
joe => *Note, this record has no bigrams*

Notice, that there are 2 instances of the word bigram "there Joe".

So, your task is to write a MrJob MapReduce program which accepts the following file as input

hdfs:///user/hadoop/w.data

and outputs key value pairs where each one has a key which is some word bigram string, and the value a count of the number of occurrences of that word bigram. Note, please convert all words to lower case on input, so Hello and hello become the same word.

Our job should output key value pairs similar to the following:

"hello there", 1
"hi there", 1
"joe there", 1
"there joe", 2

10) (5 points) When you have accomplished this, please submit the following, (1) a copy of your MRJob code and (2) a copy of the output of the execution of that code for at least the first 10 bigram key value pairs.

Ans:

- **Code: WordCount4.py**

WordCount4.py

File Edit View

```
from mrjob.job import MRJob
import re

WORD_RE = re.compile(r"[\w']+")

class MRWordCount4(MRJob):

    def mapper(self, _, line):
        # Find all words in the line and convert it to lowercase
        words = WORD_RE.findall(line.lower())

        for i in range(len(words) - 1):
            bigram = f"{words[i]} {words[i+1]}"
            yield bigram, 1

    def combiner(self, bigram, counts):
        yield bigram, sum(counts)

    def reducer(self, bigram, counts):
        yield bigram, sum(counts)

if __name__ == '__main__':
    MRWordCount4.run()
```

- **Output:**

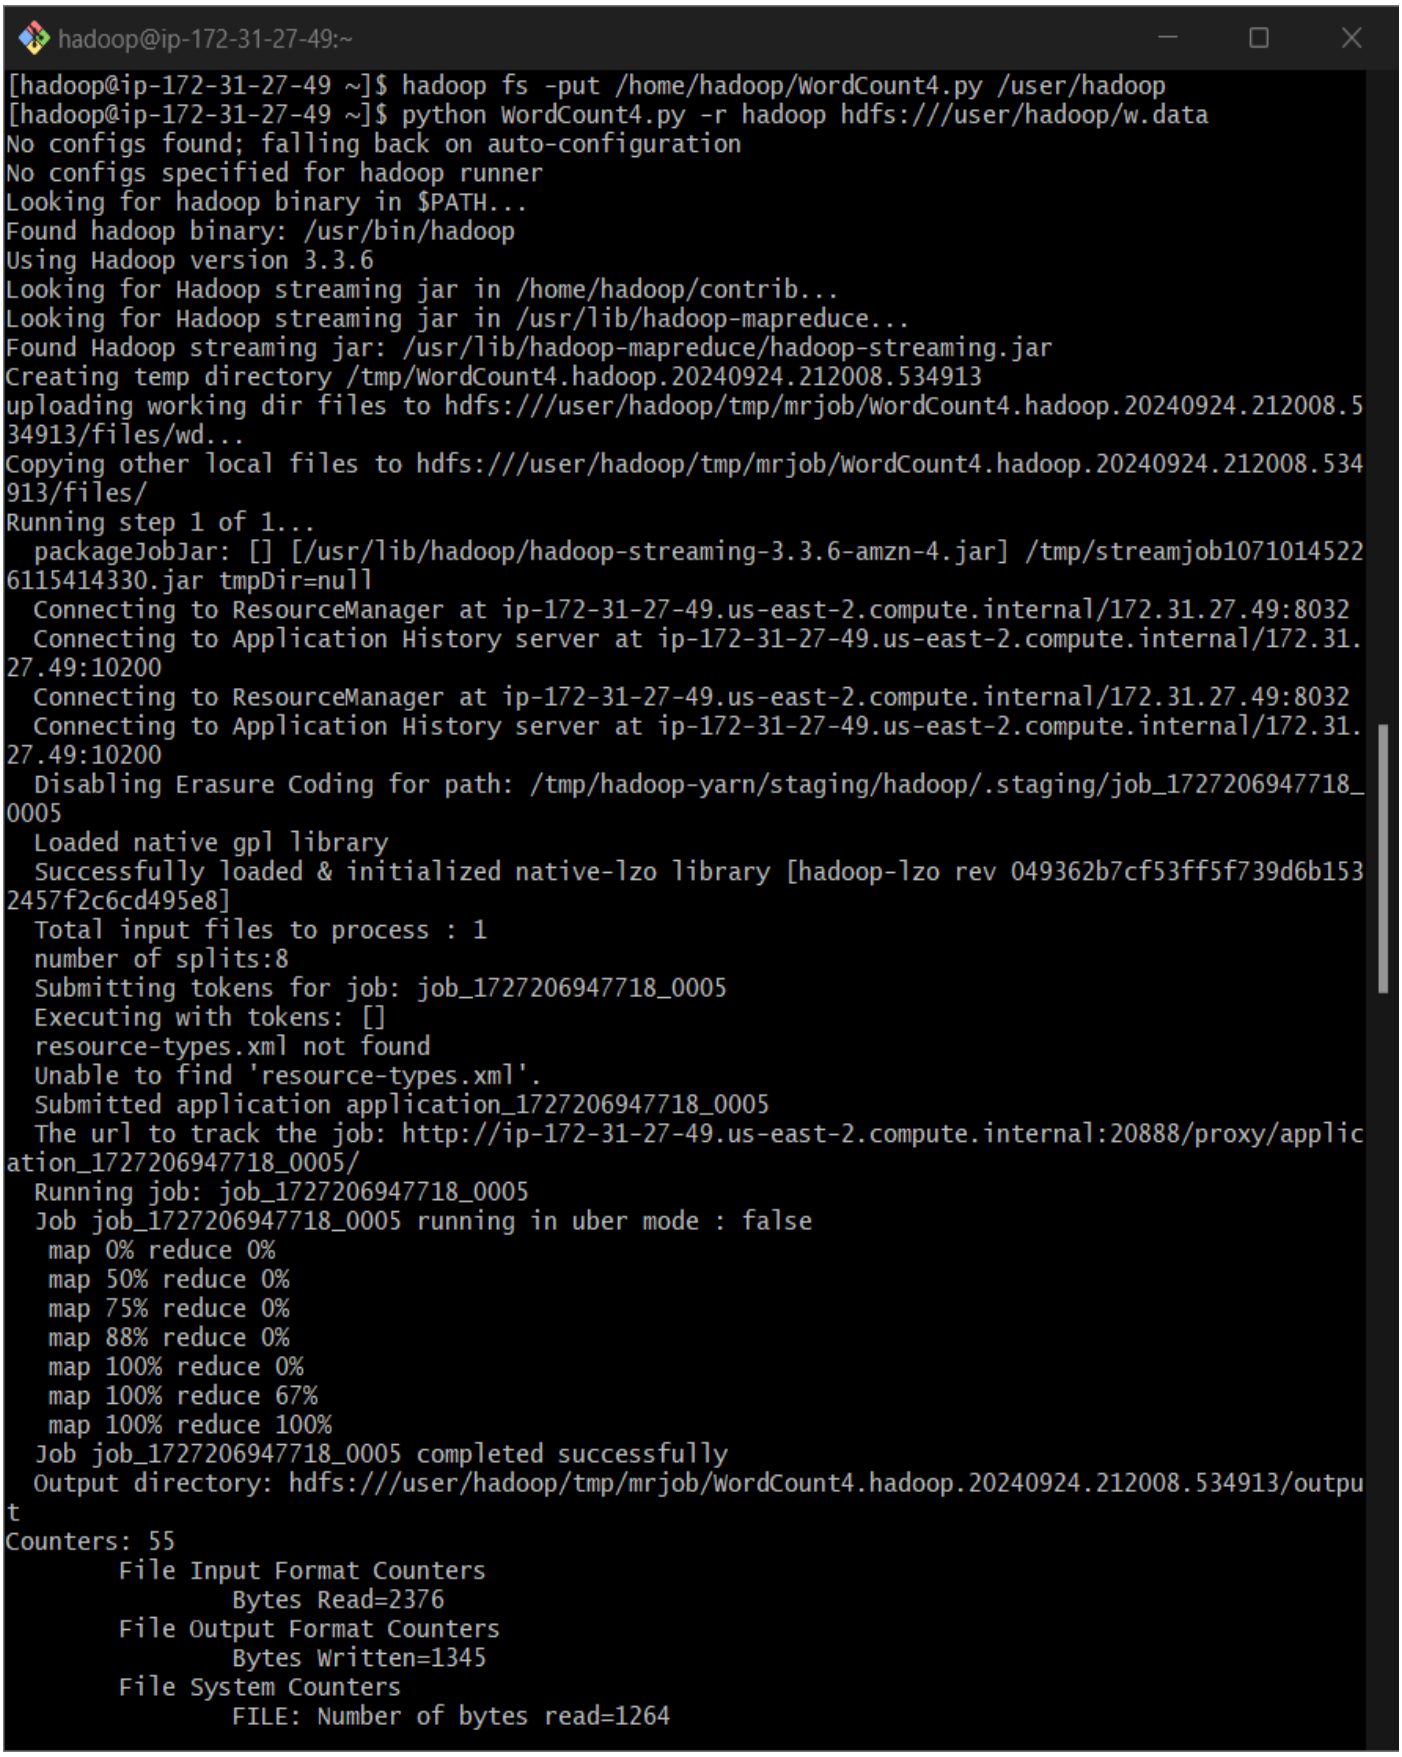
hadoop@ip-172-31-27-49:~
[hadoop@ip-172-31-27-49 ~]$ hadoop fs -put /home/hadoop/WordCount4.py /user/hadoop
[hadoop@ip-172-31-27-49 ~]$ python WordCount4.py -r hadoop hdfs:///user/hadoop/w.data
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 3.3.6
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/WordCount4.hadoop.20240924.212008.534913
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/WordCount4.hadoop.20240924.212008.534913/files/wd...
Copying other local files to hdfs:///user/hadoop/tmp/mrjob/WordCount4.hadoop.20240924.212008.534913/files/
Running step 1 of 1...
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6-amzn-4.jar] /tmp/streamjob10710145226115414330.jar tmpDir=null
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1727206947718_0005
Loaded native gpl library
Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b1532457f2c6cd495e8]
Total input files to process : 1
number of splits:8
Submitting tokens for job: job_1727206947718_0005
Executing with tokens: []
resource-types.xml not found
Unable to find 'resource-types.xml'.
Submitted application application_1727206947718_0005
The url to track the job: http://ip-172-31-27-49.us-east-2.compute.internal:20888/proxy/application_1727206947718_0005/
Running job: job_1727206947718_0005
Job job_1727206947718_0005 running in uber mode : false
map 0% reduce 0%
map 50% reduce 0%
map 75% reduce 0%
map 88% reduce 0%
map 100% reduce 0%
map 100% reduce 67%
map 100% reduce 100%
Job job_1727206947718_0005 completed successfully
Output directory: hdfs:///user/hadoop/tmp/mrjob/WordCount4.hadoop.20240924.212008.534913/output
Counters: 55
File Input Format Counters
Bytes Read=2376
File Output Format Counters
Bytes Written=1345
File System Counters
FILE: Number of bytes read=1264


hadoop@ip-172-31-27-49:~
FILE: Number of bytes written=3267154
FILE: Number of large read operations=0
FILE: Number of read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=3376
HDFS: Number of bytes read erasure-coded=0
HDFS: Number of bytes written=1345
HDFS: Number of large read operations=0
HDFS: Number of read operations=39
HDFS: Number of write operations=6
Job Counters
Data-local map tasks=8
Killed map tasks=1
Launched map tasks=8
Launched reduce tasks=3
Total megabyte-milliseconds taken by all map tasks=197950464
Total megabyte-milliseconds taken by all reduce tasks=70785024
Total time spent by all map tasks (ms)=128874
Total time spent by all maps in occupied slots (ms)=6185952
Total time spent by all reduce tasks (ms)=23042
Total time spent by all reduces in occupied slots (ms)=2212032
Total vcore-milliseconds taken by all map tasks=128874
Total vcore-milliseconds taken by all reduce tasks=23042
Map-Reduce Framework
CPU time spent (ms)=12610
Combine input records=92
Combine output records=91
Failed Shuffles=0
GC time elapsed (ms)=733
Input split bytes=1000
Map input records=6
Map output bytes=1362
Map output materialized bytes=1724
Map output records=92
Merged Map outputs=24
Peak Map Physical memory (bytes)=512450560
Peak Map Virtual memory (bytes)=3204685824
Peak Reduce Physical memory (bytes)=314343424
Peak Reduce Virtual memory (bytes)=4559392768
Physical memory (bytes) snapshot=4775915520
Reduce input groups=91
Reduce input records=91
Reduce output records=91
Reduce shuffle bytes=1724
Shuffled Maps =24
Spilled Records=182
Total committed heap usage (bytes)=4283432960
Virtual memory (bytes) snapshot=39259901952
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/WordCount4.hadoop.20240924.212008.534913/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/WordCount4.hadoop.20240924.212008.5349

hadoop@ip-172-31-27-49:~

```
13/output...
"all dependencies"      1
"and writing"   1
"are more"      1
"as well"       1
"combine or"    1
"contained within"      1
"executed on"   1
"explains how"  1
"following two" 1
"how to"        1
"how your"      1
"is run"        1
"is submitted"  1
"of writing"    1
"on that"       1
"on your"       1
"or reduce"     1
"runners explains"      1
"see how"       1
"submitted runners"     1
"those things"  1
"to be" 1
"to do" 1
"within the"    1
"your machine"  1
"your program"  1
"your second"   1
"a hadoop"      1
"as on" 1
"be contained"  1
"be defined"    1
"be executed"   1
"by mrjob"      1
"cluster as"    1
"defined in"    1
"dependencies must"     1
"file to"       1
"job and"       1
"map combine"   1
"mrjob when"    1
"nodes or"      1
"our job"       1
"program is"    1
"second job"    1
"the file"      1
"the following" 1
"to the"        1
"two sections"  1
"uploaded to"   1
"versions of"   1
"well as"       1
"when your"     1
"writing your"  2
"your job"      1
"a file"        1
"a python"      1
```

hadoop@ip-172-31-27-49:~

```
"an individual" 1
"as a"  1
"as an" 1
"available on"  1
"cluster by"    1
"do those"      1
"either be"     1
"file available"        1
"first job"     1
"for more"      1
"hadoop cluster"        1
"in a"  1
"individual map"        1
"job is"        1
"job will"      1
"machine as"    1
"more on"       1
"more reference"        1
"must either"   1
"on a"  1
"on the"        1
"or uploaded"   1
"oriented versions"     1
"python script" 1
"reduce task"   1
"reference oriented"    1
"run for"       1
"script as"     1
"sections are"  1
"task nodes"    1
"task see"      1
"the cluster"   1
"the task"      1
"will be"       1
"your first"    1
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/WordCount4.hadoop.20240924.212008.534
913...
Removing temp directory /tmp/WordCount4.hadoop.20240924.212008.534913...
[hadoop@ip-172-31-27-49 ~]$
```

11) Now do the same as the above for the files Salaries.py and Salaries.tsv. The ".tsv" file holds department and salary information for Baltimore municipal workers. Have a look at Salaries.py for the layout of the ".tsv" file and how to read it in to our map reduce program.

12) Execute the Salaries.py program to make sure it works. It should print out how many workers share each job title.

Ans:

- **Code: Salaries.py**

Salaries.py

File Edit View

```
from mrjob.job import MRJob

class MRSalaries(MRJob):

    def mapper(self, _, line):
        (name,jobTitle,agencyID,agency,hireDate,annualSalary,grossPay) = line.split('\t')
        yield jobTitle, 1

    def combiner(self, jobTitle, counts):
        yield jobTitle, sum(counts)

    def reducer(self, jobTitle, counts):
        yield jobTitle, sum(counts)


if __name__ == '__main__':
    MRSalaries.run()
```

- **Output:**

hadoop@ip-172-31-27-49:~

```
[hadoop@ip-172-31-27-49 ~]$ hadoop fs -put /home/hadoop/Salaries.py /user/hadoop
[hadoop@ip-172-31-27-49 ~]$ hadoop fs -put /home/hadoop/Salaries.tsv /user/hadoop
[hadoop@ip-172-31-27-49 ~]$ python Salaries.py -r hadoop hdfs:///user/hadoop/Salaries.tsv
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 3.3.6
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/Salaries.hadoop.20240924.212925.983915
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20240924.212925.983
915/files/wd...
Copying other local files to hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20240924.212925.98391
5/files/
Running step 1 of 1...
  packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6-amzn-4.jar] /tmp/streamjob4164776637
965734357.jar tmpDir=null
  Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
  Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.
27.49:10200
  Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
  Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.
27.49:10200
  Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1727206947718_
0006
  Loaded native gpl library
  Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b153
2457f2c6cd495e8]
  Total input files to process : 1
  number of splits:8
  Submitting tokens for job: job_1727206947718_0006
  Executing with tokens: []
  resource-types.xml not found
  Unable to find 'resource-types.xml'.
  Submitted application application_1727206947718_0006
  The url to track the job: http://ip-172-31-27-49.us-east-2.compute.internal:20888/proxy/applic
ation_1727206947718_0006/
  Running job: job_1727206947718_0006
  Job job_1727206947718_0006 running in uber mode : false
   map 0% reduce 0%
   map 25% reduce 0%
   map 75% reduce 0%
   map 100% reduce 0%
   map 100% reduce 67%
   map 100% reduce 100%
  Job job_1727206947718_0006 completed successfully
  Output directory: hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20240924.212925.983915/output
Counters: 55
        File Input Format Counters
                Bytes Read=1567508
        File Output Format Counters
                Bytes Written=29260
        File System Counters
                FILE: Number of bytes read=27045
                FILE: Number of bytes written=3355786
```

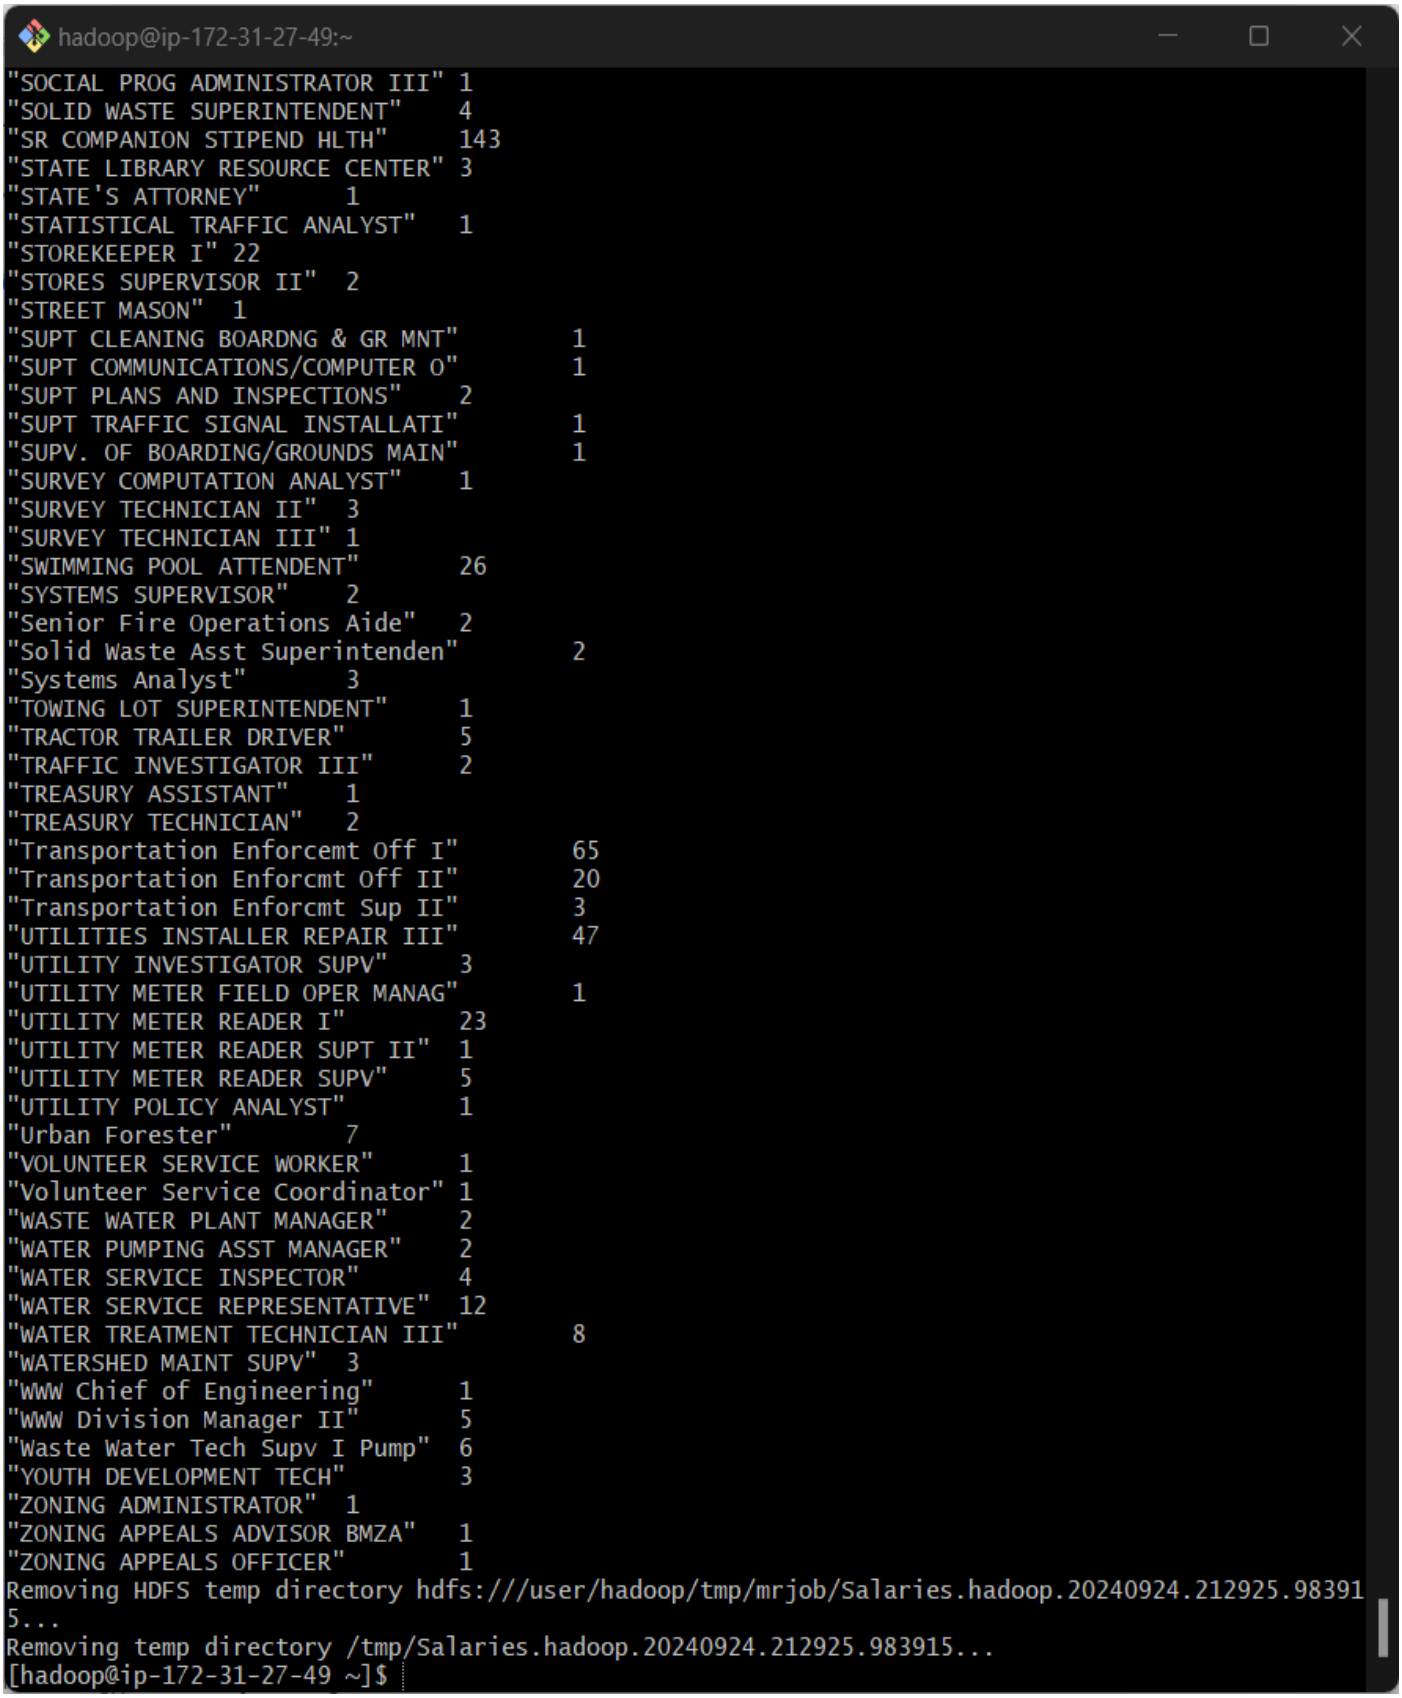
hadoop@ip-172-31-27-49:~
"SOCIAL PROG ADMINISTRATOR III" 1
"SOLID WASTE SUPERINTENDENT" 4
"SR COMPANION STIPEND HLTH" 143
"STATE LIBRARY RESOURCE CENTER" 3
"STATE'S ATTORNEY" 1
"STATISTICAL TRAFFIC ANALYST" 1
"STOREKEEPER I" 22
"STORES SUPERVISOR II" 2
"STREET MASON" 1
"SUPT CLEANING BOARDNG & GR MNT" 1
"SUPT COMMUNICATIONS/COMPUTER O" 1
"SUPT PLANS AND INSPECTIONS" 2
"SUPT TRAFFIC SIGNAL INSTALLATI" 1
"SUPV. OF BOARDING/GROUNDS MAIN" 1
"SURVEY COMPUTATION ANALYST" 1
"SURVEY TECHNICIAN II" 3
"SURVEY TECHNICIAN III" 1
"SWIMMING POOL ATTENDENT" 26
"SYSTEMS SUPERVISOR" 2
"Senior Fire Operations Aide" 2
"Solid Waste Asst Superintenden" 2
"Systems Analyst" 3
"TOWING LOT SUPERINTENDENT" 1
"TRACTOR TRAILER DRIVER" 5
"TRAFFIC INVESTIGATOR III" 2
"TREASURY ASSISTANT" 1
"TREASURY TECHNICIAN" 2
"Transportation Enforcemt Off I" 65
"Transportation Enforcmt Off II" 20
"Transportation Enforcmt Sup II" 3
"UTILITIES INSTALLER REPAIR III" 47
"UTILITY INVESTIGATOR SUPV" 3
"UTILITY METER FIELD OPER MANAG" 1
"UTILITY METER READER I" 23
"UTILITY METER READER SUPT II" 1
"UTILITY METER READER SUPV" 5
"UTILITY POLICY ANALYST" 1
"Urban Forester" 7
"VOLUNTEER SERVICE WORKER" 1
"Volunteer Service Coordinator" 1
"WASTE WATER PLANT MANAGER" 2
"WATER PUMPING ASST MANAGER" 2
"WATER SERVICE INSPECTOR" 4
"WATER SERVICE REPRESENTATIVE" 12
"WATER TREATMENT TECHNICIAN III" 8
"WATERSHED MAINT SUPV" 3
"WWW Chief of Engineering" 1
"WWW Division Manager II" 5
"Waste Water Tech Supv I Pump" 6
"YOUTH DEVELOPMENT TECH" 3
"ZONING ADMINISTRATOR" 1
"ZONING APPEALS ADVISOR BMZA" 1
"ZONING APPEALS OFFICER" 1
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/Salaries.hadoop.20240924.212925.98391
5...
Removing temp directory /tmp/Salaries.hadoop.20240924.212925.983915...
[hadoop@ip-172-31-27-49 ~]$

13) Now modify the Salaries.py program. Call it Salaries2.py

Instead of counting the number of workers per department, change the program to provide the number of workers having High, Medium or Low annual salaries. This is defined as follows:

High	100,000.00 and above
Medium	50,000.00 to 99,999.99
Low	0.00 to 49,999.99

The output of the program should be something like the following (in any order):

High 20
Medium 30
Low 10

Some important hints:

- The annual salary is a string that will need to be converted to a float.
- The mapper should output tuples with one of three keys depending on the annual salary: High, Medium and Low
- The value part of the tuple is not a salary. (What should it be?)

Now execute the program and see what happens.

14) (3 points) Submit (1) a copy of this modified program and (2) a screen shot of the results of the program's execution as the output of your assignment.

Ans:

- **Code: Salaries2.py**

Salaries2.py

File Edit View

```
from mrjob.job import MRJob

class MRSalaries2(MRJob):

    def mapper(self, _, line):
        (name, jobTitle, agencyID, agency, hireDate, annualSalary, grossPay) = line.split('\t')

        try:
            salary = float(annualSalary)
        except ValueError:
            return

        if salary >= 100000.00:
            yield "High", 1
        elif 50000.00 <= salary < 100000.00:
            yield "Medium", 1
        else:
            yield "Low", 1

    def combiner(self, salary_level, counts):
        yield salary_level, sum(counts)

    def reducer(self, salary_level, counts):
        yield salary_level, sum(counts)


if __name__ == '__main__':
    MRSalaries2.run()
```

- **Output:**

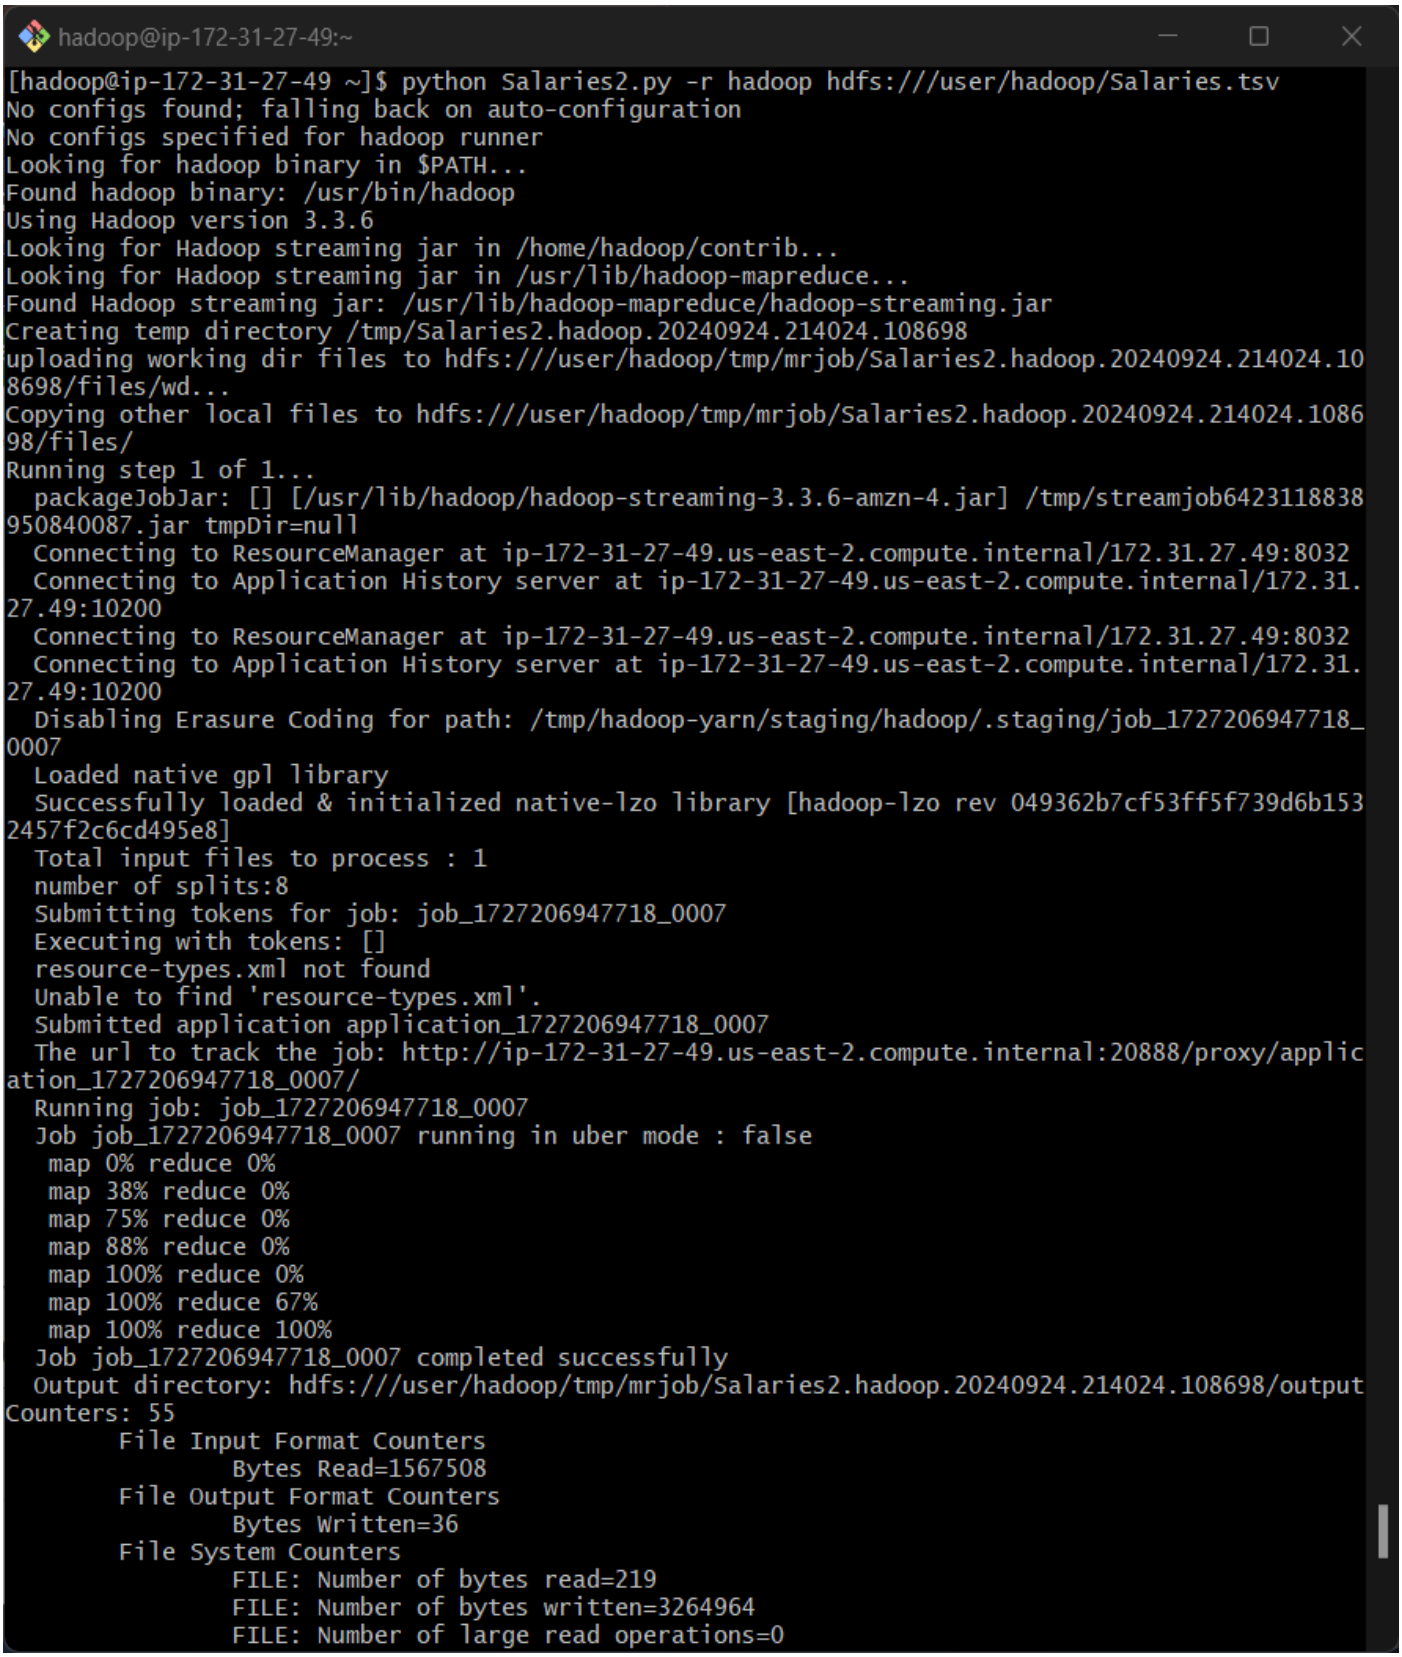
hadoop@ip-172-31-27-49:~
[hadoop@ip-172-31-27-49 ~]$ python Salaries2.py -r hadoop hdfs:///user/hadoop/Salaries.tsv
No configs found; falling back on auto-configuration
No configs specified for hadoop runner
Looking for hadoop binary in $PATH...
Found hadoop binary: /usr/bin/hadoop
Using Hadoop version 3.3.6
Looking for Hadoop streaming jar in /home/hadoop/contrib...
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar
Creating temp directory /tmp/Salaries2.hadoop.20240924.214024.108698
uploading working dir files to hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20240924.214024.108698/files/wd...
Copying other local files to hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20240924.214024.108698/files/
Running step 1 of 1...
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-3.3.6-amzn-4.jar] /tmp/streamjob6423118838950840087.jar tmpDir=null
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Connecting to ResourceManager at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:8032
Connecting to Application History server at ip-172-31-27-49.us-east-2.compute.internal/172.31.27.49:10200
Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1727206947718_0007
Loaded native gpl library
Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b1532457f2c6cd495e8]
Total input files to process : 1
number of splits:8
Submitting tokens for job: job_1727206947718_0007
Executing with tokens: []
resource-types.xml not found
Unable to find 'resource-types.xml'.
Submitted application application_1727206947718_0007
The url to track the job: http://ip-172-31-27-49.us-east-2.compute.internal:20888/proxy/application_1727206947718_0007/
Running job: job_1727206947718_0007
Job job_1727206947718_0007 running in uber mode : false
map 0% reduce 0%
map 38% reduce 0%
map 75% reduce 0%
map 88% reduce 0%
map 100% reduce 0%
map 100% reduce 67%
map 100% reduce 100%
Job job_1727206947718_0007 completed successfully
Output directory: hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20240924.214024.108698/output
Counters: 55
File Input Format Counters
Bytes Read=1567508
File Output Format Counters
Bytes Written=36
File System Counters
FILE: Number of bytes read=219
FILE: Number of bytes written=3264964
FILE: Number of large read operations=0

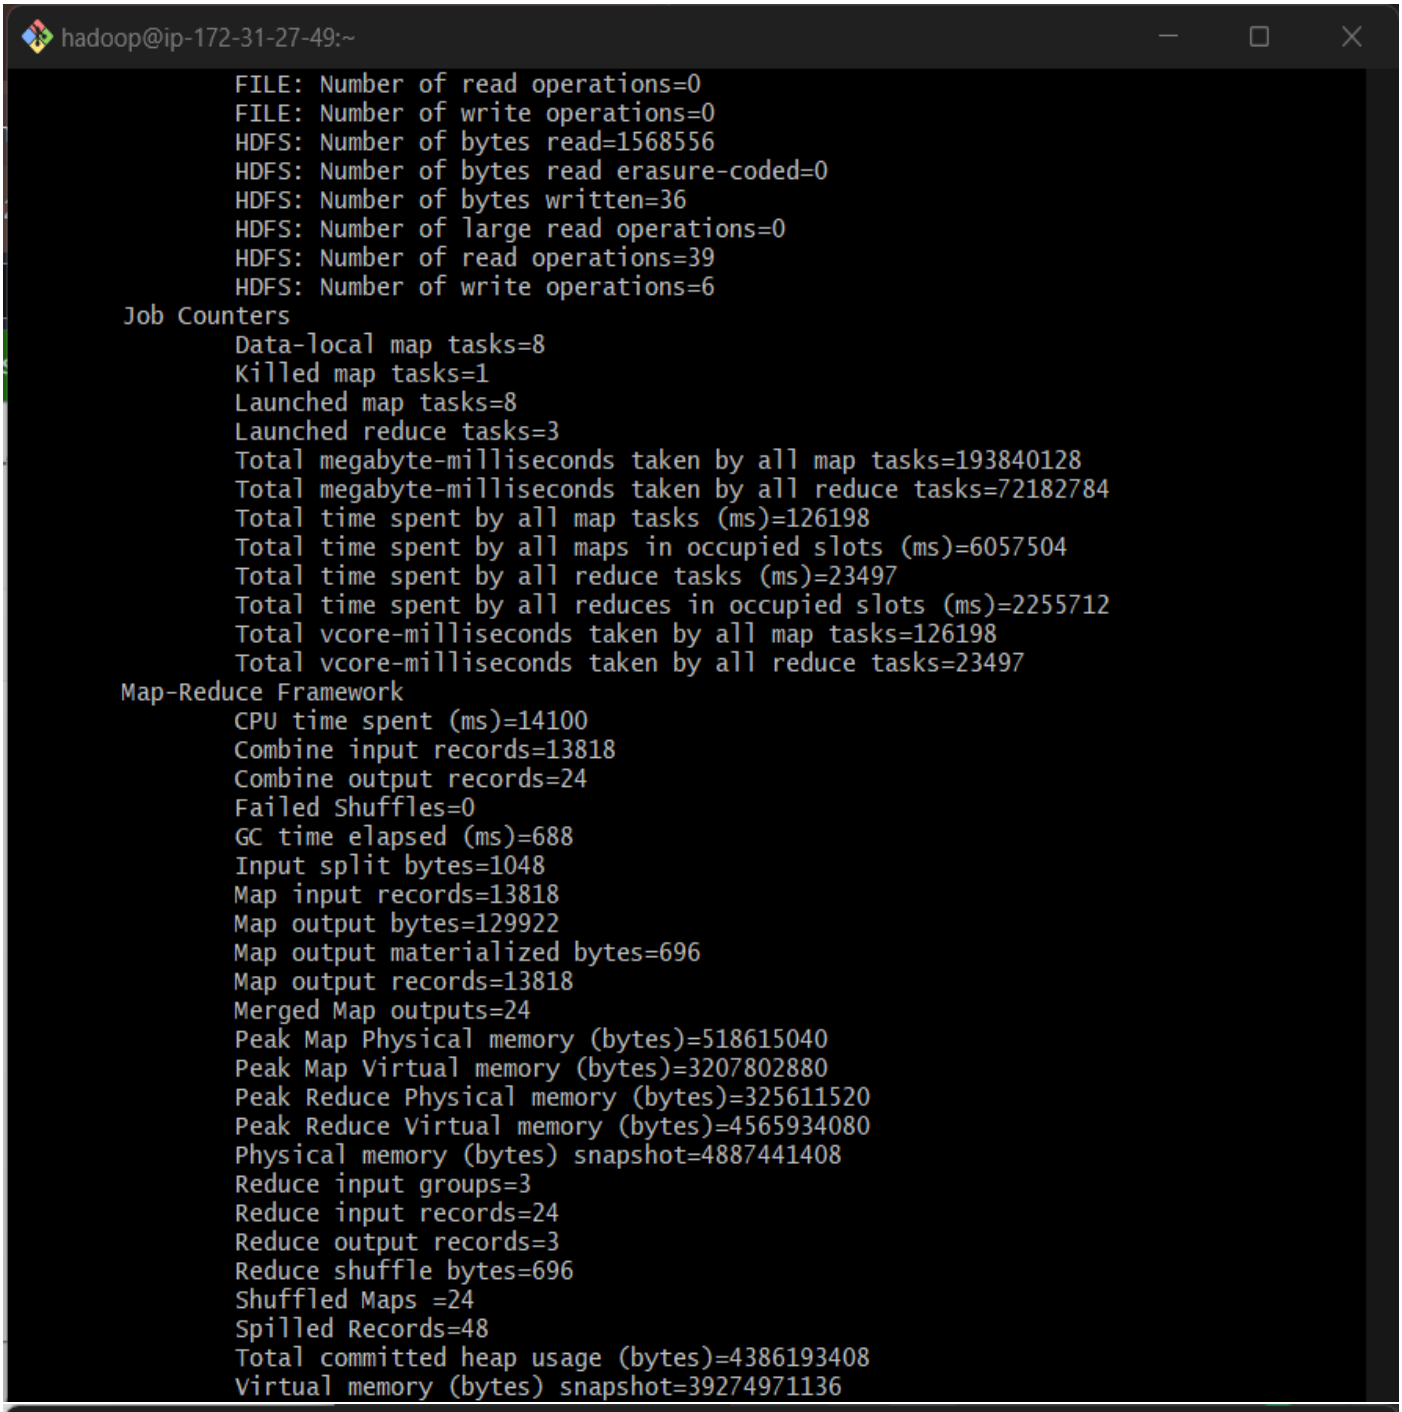
hadoop@ip-172-31-27-49:~
FILE: Number of read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=1568556
HDFS: Number of bytes read erasure-coded=0
HDFS: Number of bytes written=36
HDFS: Number of large read operations=0
HDFS: Number of read operations=39
HDFS: Number of write operations=6
Job Counters
Data-local map tasks=8
Killed map tasks=1
Launched map tasks=8
Launched reduce tasks=3
Total megabyte-milliseconds taken by all map tasks=193840128
Total megabyte-milliseconds taken by all reduce tasks=72182784
Total time spent by all map tasks (ms)=126198
Total time spent by all maps in occupied slots (ms)=6057504
Total time spent by all reduce tasks (ms)=23497
Total time spent by all reduces in occupied slots (ms)=2255712
Total vcore-milliseconds taken by all map tasks=126198
Total vcore-milliseconds taken by all reduce tasks=23497
Map-Reduce Framework
CPU time spent (ms)=14100
Combine input records=13818
Combine output records=24
Failed Shuffles=0
GC time elapsed (ms)=688
Input split bytes=1048
Map input records=13818
Map output bytes=129922
Map output materialized bytes=696
Map output records=13818
Merged Map outputs=24
Peak Map Physical memory (bytes)=518615040
Peak Map Virtual memory (bytes)=3207802880
Peak Reduce Physical memory (bytes)=325611520
Peak Reduce Virtual memory (bytes)=4565934080
Physical memory (bytes) snapshot=4887441408
Reduce input groups=3
Reduce input records=24
Reduce output records=3
Reduce shuffle bytes=696
Shuffled Maps =24
Spilled Records=48
Total committed heap usage (bytes)=4386193408
Virtual memory (bytes) snapshot=39274971136

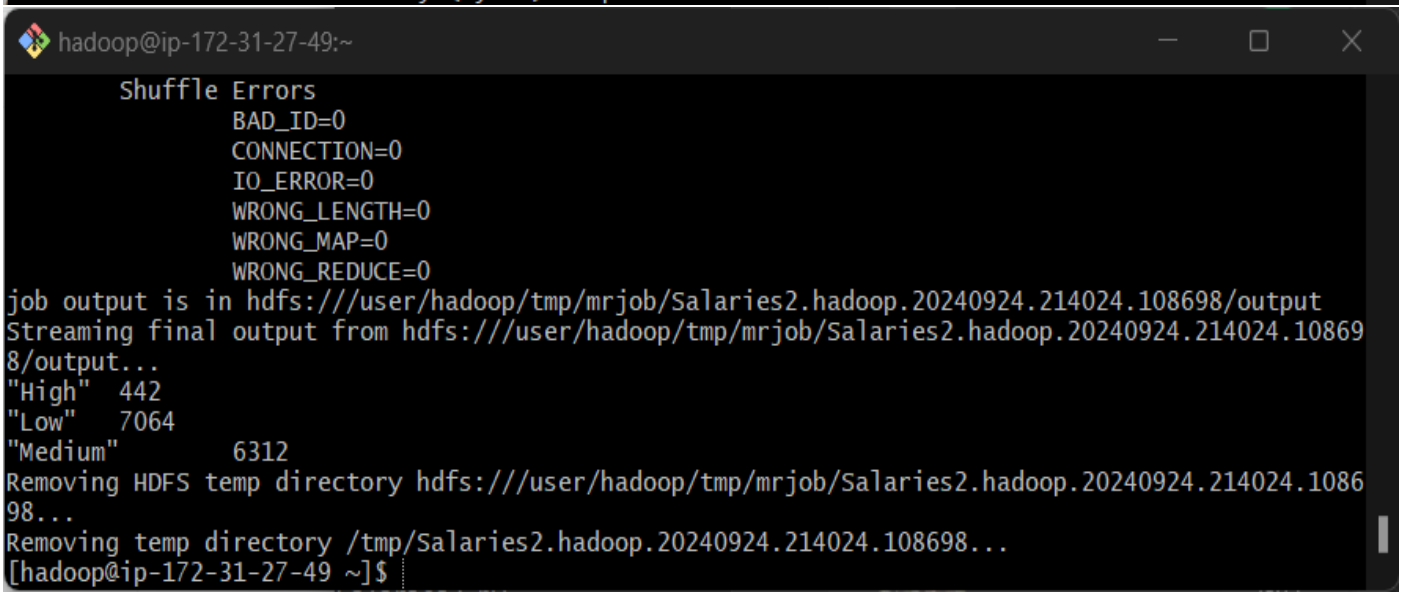
hadoop@ip-172-31-27-49:~
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20240924.214024.108698/output
Streaming final output from hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20240924.214024.10869
8/output...
"High" 442
"Low" 7064
"Medium" 6312
Removing HDFS temp directory hdfs:///user/hadoop/tmp/mrjob/Salaries2.hadoop.20240924.214024.1086
98...
Removing temp directory /tmp/Salaries2.hadoop.20240924.214024.108698...
[hadoop@ip-172-31-27-49 ~]$

15) Remember to terminate your EMR cluster and remove your S3 bucket!

Ans: