

Name: Deep Pawar(A20545137)
Professor: Joseph Rosen
Institute: Illinois Institute of Technology

CSP 554: Big Data Technologies

Fall 2024 - Assignment 4

- **Questions and Answers:**
- Configure a Hadoop (AWS EMR) environment as you have done for previous assignments:

The screenshot shows the AWS Management Console interface for an Amazon EMR cluster. The cluster name is 'Mycluster_Assignment_4' and its ID is 'j-3BZH1ORMUB676'. The cluster is in a 'Waiting' state. The console displays various tabs for the cluster, including Summary, Bootstrap actions, Instances (Hardware), Steps, Applications, Configurations, Monitoring, Events, and Tags (1). The Summary tab is currently selected, showing details such as Cluster ID, Amazon EMR version (emr-7.2.0), Installed applications (Hadoop 3.3.6, Hive 3.1.3, Hue 4.11.0, Pig 0.17.0, Tez 0.10.2), and Capacity (1 Primary, 1 Core, 0 Task). A notification banner at the top states 'Your cluster "Mycluster_Assignment_4" has been successfully created.' A tooltip is visible over the 'Primary node public DNS copied' link, showing the public DNS address: 'ec2-3-129-16-208.us-east-2.compute.amazonaws.com'.

The screenshot shows a terminal window with a Hadoop command prompt. The prompt is 'hadoop@ip-172-31-30-159:~'. The terminal displays a large ASCII art graphic of a Hadoop logo, which is a stylized 'H' made of characters. Below the logo, there is a URL: 'https://aws.amazon.com/linux/amazon-linux-2023'. The terminal also shows a large block of text consisting of a grid of characters (E, M, R) arranged in a pattern that resembles a Hadoop logo. The terminal prompt is '[hadoop@ip-172-31-30-159 ~]\$'.

- Log on to your VM using ssh and execute the file using “java TestDataGen”

[illegible]

- This will output a magic number which you should copy down and provide with the results of your assignment.

Ans: Magic Number = 38214

- It will also place the files foodratings<magic number>.txt and foodplaces<magic number>.txt in your VM home directory

Ans:

```

hadoop@ip-172-31-30-159:~
~~~~~
_/_m/'

EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRRRR
E::::::::::::::::::E M::M M::M R::::::::R
EE::::::::::::::::::E M::M M::M R::RRRRRR::R
    E::E      EEEE M::M M::M RR::R      R::R
    E::E      M::M M::M M::M R::R      R::R
    E::E      M::M M::M M::M R::RRRRR::R
    E::E      M::M M::M M::M R::::::::RR
    E::E      M::M M::M M::M R::RRRRR::R
    E::E      M::M M::M M::M R::R      R::R
    E::E      EEEE MMM   M::M R::R      R::R
EE::::::::::::::::::E M::M M::M R::R      R::R
E::::::::::::::::::E M::M M::M RR::R      R::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRR      RRRRRR

[hadoop@ip-172-31-30-159 ~]$ java TestDataGen
Magic Number = 38214
[hadoop@ip-172-31-30-159 ~]$ ls
TestDataGen.class foodplaces38214.txt foodratings38214.txt
[hadoop@ip-172-31-30-159 ~]$
```

Exercise 1) 2 points

a. Create a Hive database called “MyDb”.

Ans:

- **Command Used:** create database MyDb;

```
hive> create database MyDb;
OK
Time taken: 0.345 seconds
hive>
```

b. Now in MyDb create a table with name foodratings

Ans:

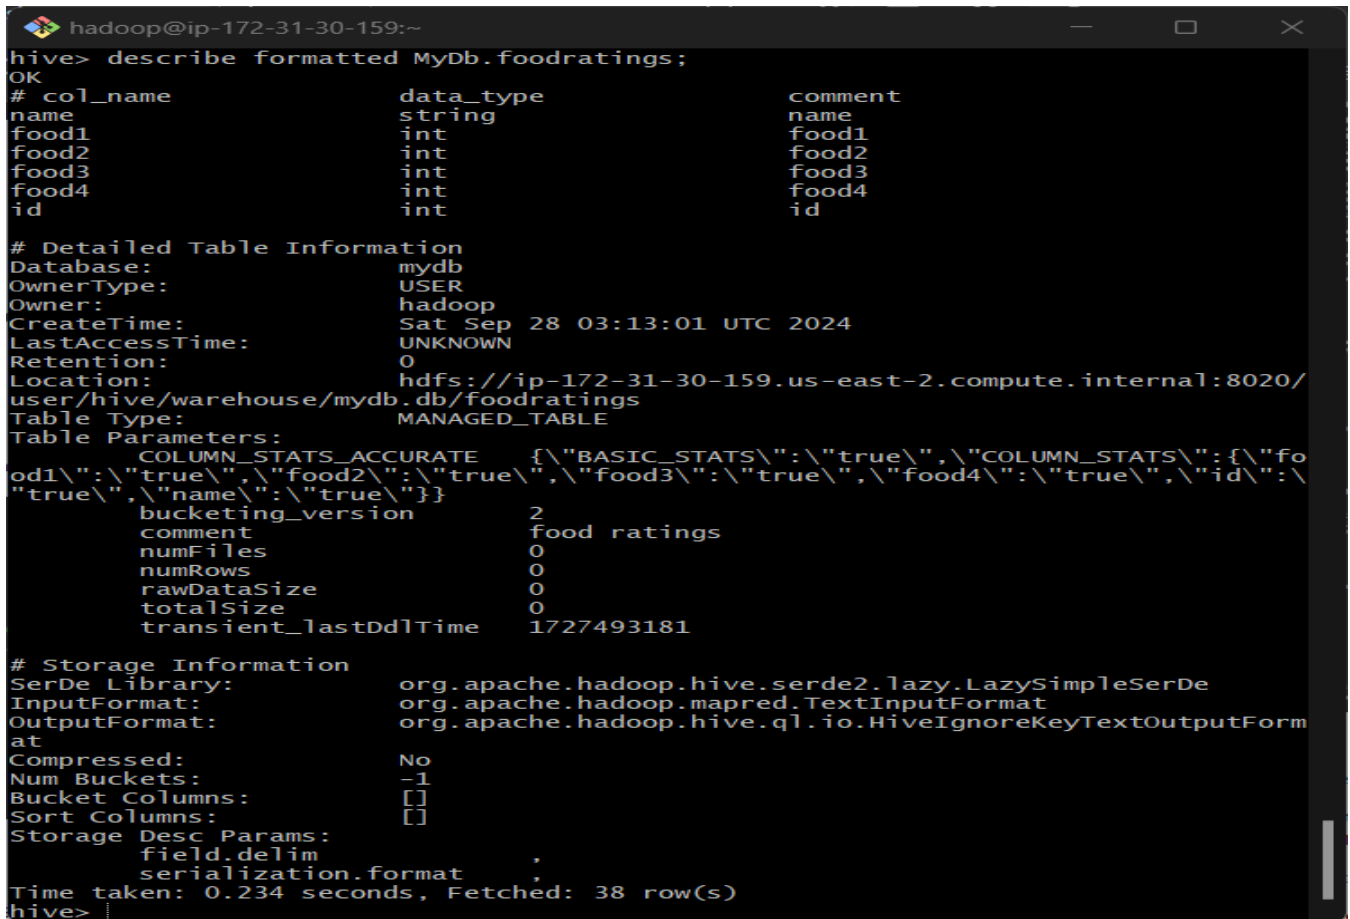
- **Command Used:** CREATE TABLE MyDb.foodratings(name STRING COMMENT 'name', food1 INT COMMENT 'food1', food2 INT COMMENT 'food2', food3 INT COMMENT 'food3', food4 INT COMMENT 'food4', id INT COMMENT 'id') COMMENT 'food ratings' ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;

```
hive> CREATE TABLE MyDb.foodratings( name STRING COMMENT 'name', food1 INT COMMENT 'food1', food2 INT COMMENT 'food2', food3 INT COMMENT 'food3', food4 INT COMMENT 'food4', id INT COMMENT 'id' ) COMMENT 'food ratings' ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;
OK
Time taken: 0.252 seconds
hive>
```

c. Execute a Hive command of ‘DESCRIBE FORMATTED MyDb.foodratings;’ and capture its output as one of the results of this exercise.

Ans:

- **Command Used:** describe formatted MyDb.foodratings;



```
hadoop@ip-172-31-30-159:~
hive> describe formatted MyDb.foodratings;
OK
# col_name      data_type      comment
name            string         name
food1           int            food1
food2           int            food2
food3           int            food3
food4           int            food4
id              int            id

# Detailed Table Information
Database:      mydb
OwnerType:     USER
Owner:         hadoop
CreateTime:    Sat Sep 28 03:13:01 UTC 2024
LastAccessTime: UNKNOWN
Retention:     0
Location:      hdfs://ip-172-31-30-159.us-east-2.compute.internal:8020/
user/hive/warehouse/mydb.db/foodratings
Table Type:    MANAGED_TABLE
Table Parameters:
  COLUMN_STATS_ACCURATE      {"BASIC_STATS\"":\"true\", \"COLUMN_STATS\"\": {\"fo
od1\": \"true\", \"food2\": \"true\", \"food3\": \"true\", \"food4\": \"true\", \"id\": \"
true\", \"name\": \"true\"}}
  bucketing_version          2
  comment                    food ratings
  numFiles                   0
  numRows                   0
  rawDataSize                0
  totalSize                  0
  transient_lastDdlTime      1727493181

# Storage Information
SerDe Library:  org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:    org.apache.hadoop.mapred.TextInputFormat
OutputFormat:   org.apache.hadoop.hive.q1.io.HiveIgnoreKeyTextOutputForm
at
Compressed:     No
Num Buckets:    -1
Bucket Columns: []
Sort Columns:   []
Storage Desc Params:
  field.delim      ,
  serialization.format
Time taken: 0.234 seconds, Fetched: 38 row(s)
hive>
```

d. Then in MyDb create a table with name foodplaces

Ans:

- **Command Used:** CREATE TABLE MyDb.foodplaces (id INT, place STRING) COMMENT 'food places' ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;

```
hive> CREATE TABLE MyDb.foodplaces ( id INT, place STRING ) COMMENT 'food places' ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;
OK
Time taken: 0.096 seconds
```

e. Execute a Hive command of 'DESCRIBE FORMATTED MyDb.foodplaces' and capture its output as another of the results of this exercise.

Ans:

- **Command Used:** describe formatted MyDb.foodplaces;

```
hive> describe formatted MyDb.foodplaces;
OK
# col_name          data_type          comment
id                  int
place               string

# Detailed Table Information
Database:           mydb
OwnerType:          USER
Owner:              hadoop
CreateTime:         Sat Sep 28 03:20:55 UTC 2024
LastAccessTime:     UNKNOWN
Retention:          0
Location:           hdfs://ip-172-31-30-159.us-east-2.compute.internal:8020/user/hive/warehouse/mydb.db/foodplaces
Table Type:         MANAGED_TABLE
Table Parameters:
  COLUMN_STATS_ACCURATE  {"BASIC_STATS\":"true","COLUMN_STATS\":"id":"true","place":"true"}}
  bucketing_version      2
  comment                 food places
  numFiles                0
  numRows                0
  rawDataSize            0
  totalSize              0
  transient_lastDdlTime  1727493655

# Storage Information
SerDe Library:        org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:          org.apache.hadoop.mapred.TextInputFormat
OutputFormat:         org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:           No
Num Buckets:          -1
Bucket Columns:       []
Sort Columns:         []
Storage Desc Params:
  field.delim           ,
  serialization.format  ,
Time taken: 0.058 seconds, Fetched: 34 row(s)
hive> |
```

Exercise 2) 2 points

- a. Load the foodratings<magic number>.txt file created using TestDataGen from your local file system into the foodratings table.

Ans:

- **Magic number** = 38214
- **Files Generated** = foodratings38214.txt and foodplaces38214.txt

Here in this case, we are using foodrating38214.txt file dataset to load into MyDb.foodratings table.

- **Command Used:** LOAD DATA LOCAL INPATH '/home/hadoop/foodratings38214.txt' OVERWRITE INTO TABLE MyDB.foodratings;

```
hive> LOAD DATA LOCAL INPATH '/home/hadoop/foodratings38214.txt' OVERWRITE INTO TABLE MyDB.foodratings;
Loading data to table mydb.foodratings
OK
Time taken: 0.238 seconds
hive>
```

- b. Execute a single hive command to output the min, max and average of the values of the food3 column of the foodratings table. This should be one hive command, not three separate ones.

Ans:

- **Command Used:** select min(food3),max(food3),avg(food3) from MyDb.foodratings;

```
hive> select min(food3),max(food3),avg(food3) from MyDb.foodratings;
Query ID = hadoop_20240928033010_fe3c6247-c546-47e1-9914-58606185a4c5
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1727491992167_0002)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	0

VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 5.27 s

```
OK
1      50      26.099
Time taken: 12.491 seconds, Fetched: 1 row(s)
hive>
```

Exercise 3) 2 points

- a. Execute a hive command to output the min, max and average of the values of the food1 column grouped by the first column 'name'.

Ans:

- **Magic number** = 38214
- **Command Used:** select name, min(food1),max(food1), avg(food1) from MyDb.foodratings group by name;

```
hive> Select name, min(food1),max(food1), avg(food1) from MyDb.foodratings group by name;
Query ID = hadoop_20240928033438_e08cefd8-6edb-4f41-a788-1c34262c8d11
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1727491992167_0002)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0	
Reducer 2	container	SUCCEEDED	2	2	0	0	0	0	

VERTICES: 02/02 [=====>>>] 100% ELAPSED TIME: 5.47 s

```
OK
Joy      1      50      25.854077253218883
Jill     1      50      24.387283236994218
Joe      1      50      26.545945945945945
Mel      1      50      25.63157894736842
Sam      1      50      25.445
Time taken: 5.881 seconds, Fetched: 5 row(s)
hive>
```

Exercise 4) 2 points

a. In MyDb create a partitioned table called 'foodratingspart'

Ans:

Here In MyDb, we will create a partitioned table called 'foodratingspart'. The partition field is 'name' and its type should be a string. The names of the non-partition columns should be food1, food2, food3, food4 and id and their types each an integer.

- **Command Used:** create table if not exists MyDb.foodratingspart (food1 int, food2 int, food3 int, food4 int, id int) PARTITIONED BY (name String) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;

```
hive> create table if not exists MyDb.foodratingspart (food1 int, food2 int, food3 int, food4 int, id int) PARTITIONED BY (name String) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;
OK
Time taken: 0.051 seconds
hive>
```

b. Execute a Hive command of 'DESCRIBE FORMATTED MyDb.foodratingspart;' and capture its output as the result of this exercise.

Ans:

- **Command Used:** describe formatted MyDb.foodratingspart;

```
hive> describe formatted MyDb.foodratingspart;
OK
# col_name      data_type      comment
food1           int
food2           int
food3           int
food4           int
id              int

# Partition Information
# col_name      data_type      comment
name            string

# Detailed Table Information
Database:        mydb
OwnerType:      USER
Owner:           hadoop
CreateTime:     Sat Sep 28 03:36:14 UTC 2024
LastAccessTime: UNKNOWN
Retention:      0
Location:       hdfs://ip-172-31-30-159.us-east-2.compute.internal:8020/user/hive/warehouse/mydb.db/foodratingspart
Table Type:     MANAGED_TABLE
Table Parameters:
  COLUMN_STATS_ACCURATE  {"BASIC_STATS": "true"}
  bucketing_version      2
  numFiles               0
  numPartitions          0
  numRows               0
  rawDataSize            0
  totalSize              0
  transient_lastDdlTime  1727494574

# Storage Information
SerDe Library:      org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:        org.apache.hadoop.mapred.TextInputFormat
OutputFormat:       org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:         No
Num Buckets:        -1
Bucket Columns:     []
Sort Columns:       []
Storage Desc Params:
  field.delim         ,
  serialization.format
Time taken: 0.074 seconds, Fetched: 41 row(s)
hive>
```

Exercise 5) 2 points

Q. Assume that the number of food critics is relatively small, say less than 10 and the number places to eat is very large, say more than 10,000. In a few short sentences explain why using the (critic) name is a good choice for a partition field while using the place id is not.

Ans:

Partitioning a table according to a column such as the critic's name is efficient since it reduces the number of partitions and speeds up the execution of queries that filter by critic name. In this case, the column contains fewer unique values less than 10. This is because the amount of data examined can be decreased by assigning each partition to a certain critic.

On the other hand, because there are a lot of places more than 10,000 using the location id as a partition field is not a smart idea. When a highly cardinal column (such as place id) is used for partitioning, this leads to an excessive number of partitions and suboptimal query performance. Keeping track of so many partitions can be expensive and time-consuming, and query optimization benefits drop as one grows in number.

Exercise 6) 2 points

a. Configure Hive to allow dynamic partition creation as described in the lecture.

Ans:

- **Configuration code commands used:**

```
set hive.exec.dynamic.partition = true;
set hive.exec.dynamic.partition.mode = nonstrict;
set hive.exec.max.dynamic.partitions = 1000;
set hive.exec.max.dynamic.partitions.prenode = 1000;
```

```
hive> set hive.exec.dynamic.partition = true;
hive>
> ;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.exec.max.dynamic.partitions = 1000;
hive> set hive.exec.max.dynamic.partitions.prenode = 1000;
Query returned non-zero code: 1, cause: hive configuration hive.exec.max.dynamic.partitions.prenode does not exists.
hive> |
```

b. Now, use a hive command to copy from MyDB.foodratings into MyDB.foodratingspart to create a partitioned table from a non-partitioned one.

Ans:

- **Command Used:** insert overwrite table MyDb.foodratingspart PARTITION (name) select food1, food2, food3, food4, id, name from Mydb.foodratings;

```
hive> insert overwrite table MyDb.foodratingspart PARTITION (name) select food1, food2, food3, food4, id, name from Mydb.foodratings;
Query ID = hadoop_20240928034149_c3d880c1-7caf-46ab-b7f6-fe894ea5322a
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1727491992167_0003)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2	container	SUCCEEDED	2	2	0	0	0	0
Reducer 3	container	SUCCEEDED	2	2	0	0	0	0

```
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 5.28 s
Loading data to table mydb.foodratingspart partition (name=null)
Loaded : 5/5 partitions.
Time taken to load dynamic partitions: 0.3 seconds
Time taken for adding to write entity : 0.001 seconds
OK
Time taken: 11.476 seconds
hive> |
```

- c. Execute a hive command to output the min, max and average of the values of the food2 column of MyDB.foodratingspart where the food critic 'name' is either Mel or Jill.

Ans:

- **Command Used:** select name, min(food2), max(food2), avg(food2) from MyDb.foodratingspart where name = 'Mel' or name = 'Jill' group by name;

```
hive> select name, min(food2), max(food2), avg(food2) from MyDb.foodratingspart where name = 'Mel' or name = 'Jill' group by name;
Query ID = hadoop_20240928034259_a33642bc-a2d7-4b15-b41d-3e6775fa0fec
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1727491992167_0003)

-----
      VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    2         2         0         0         0         0
-----
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 5.47 s
-----
OK
Jill  1      50      25.76300578034682
Mel   1      50      25.25358851674641
Time taken: 6.032 seconds, Fetched: 2 row(s)
hive>
```

Exercise 7) 2 points

- a. Load the foodplaces<.magic number>.txt file created using TestDataGen from your local file system into the foodplaces table.

Ans:

- **Command Used:** load data local inpath '/home/hadoop/foodplaces38214.txt' overwrite into table MyDb.foodplaces;

```
hive> load data local inpath '/home/hadoop/foodplaces38214.txt' overwrite into table MyDb.foodplaces;
Loading data to table mydb.foodplaces
OK
Time taken: 0.17 seconds
hive> |
```

- b. Use a join operation between the two tables (foodratings and foodplaces) to provide the average rating for field food4 for the restaurant 'Soup Bowl'

Ans:

- **Command Used:** select places.place, avg(rating.food4) from MyDb.foodratings rating JOIN MyDb.foodplaces places ON rating.id = places.id where places.place = 'Soup Bowl' group by places.place;

```
hive> select places.place, avg(rating.food4) from MyDb.foodratings rating JOIN MyDb.foodplaces places ON rating.id = places.id where places.place = 'Soup Bowl' group by places.place;
Query ID = hadoop_20240928034535_3d75743c-eba0-4e34-9cef-01cc394cd72a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1727491992167_0003)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0	0
Map 2	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 3	container	SUCCEEDED	2	2	0	0	0	0	0

```
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 6.83 s
OK
Soup Bowl      25.4020618556701
Time taken: 7.422 seconds, Fetched: 1 row(s)
hive> |
```

Exercise 8) 4 points

Read the article “An Introduction to Big Data Formats” found on the blackboard in section “Articles” and provide short (2 to 4 sentence) answers to the following questions:

- a) When is the most important consideration when choosing a row format and when a column format for your big data file?

Ans:

Row Format works well if you need to access most or all of the columns for each row in your data. It works well with transactional data or situations where many columns are requested at once. For example, retrieving every field in a record is necessary when presenting a customer's whole transaction history. Row-based storage is usually done with formats like Avro, especially for workloads that are write-heavy.

Column Format works well for analytics, particularly in workloads that require a lot of reading, and is only necessary for a subset of columns. Column-based storage enables effective retrieval by scanning only the pertinent columns when your query spans a huge dataset but only requires a small number of columns. For these kinds of situations, formats like Parquet or ORC are ideal since they provide faster reads and higher compression.

- b) What is “splittability” for a column file format and why is it important when processing large volumes of data?

Ans:

Splittability determines the ability to process parts of a file independently which in turn enables parallel processing in Hadoop. Because it allows for parallel processing, in which several processors tackle different portions of the file simultaneously, this feature is especially crucial for processing huge datasets because it dramatically speeds up data operations.

Splittability improves performance in columnar formats by enabling systems to separate data by pertinent columns instead of reading entire rows. This minimizes resource utilization in distributed computing settings like Hadoop, where data must be handled across numerous nodes, and speeds up query execution by removing unnecessary data. Big files would create performance bottlenecks without splittability, resulting in longer processing times and more resource usage.

- c) What can files stored in column format achieve better compression than those stored in row format?

Ans:

Files stored in a column format can achieve better compression than those stored in a row format because similar types of data are consecutively placed in columns. Columnar formats make it easier for compression algorithms to find patterns and redundancy in the data since each column in the format only contains values of one type of data (e.g., all integers or dates in one column). When many data types (e.g., dates, numbers, and texts) are stored next to each other in a row-based format, the potential for compression is limited. In contrast, a column containing repeated values or ranges of comparable numbers can be compressed much more efficiently.

d) Under what circumstances would it be the best choice to use the “Parquet” column file format?

Ans:

When working with huge datasets, where queries usually concentrate on evaluating individual columns rather than full rows, the Parquet column file format is the ideal option. Parquet is quite effective for workloads that require a heavy reading, such as analytical queries for business intelligence, data warehousing, or machine learning applications.

Parquet's column-based data organization makes it possible for computers to read just the relevant columns, excluding irrelevant data and greatly enhancing read performance. This makes it especially helpful for large datasets with lots of columns, as analysis typically only requires a subset of the columns. Parquet is a great option as well when compression and storage efficiency are top concerns. It can achieve high compression ratios while keeping rapid read performance by organizing similar data types into columns. This lowers storage expenses.