Name: Deep Pawar(A20545137)
Professor: Joseph Rosen
Institute: Illinois Institute of Technology

CSP 554: Big Data Technologies

Fall 2024 - Assignment 7

- **Questions and Answers:**

1. Exercise 1)

Step A: Start up an EMR cluster as previously, but instead of choosing the "Core Hadoop" configuration chose the "Spark" configuration (see below), otherwise proceed as before.

Ans:

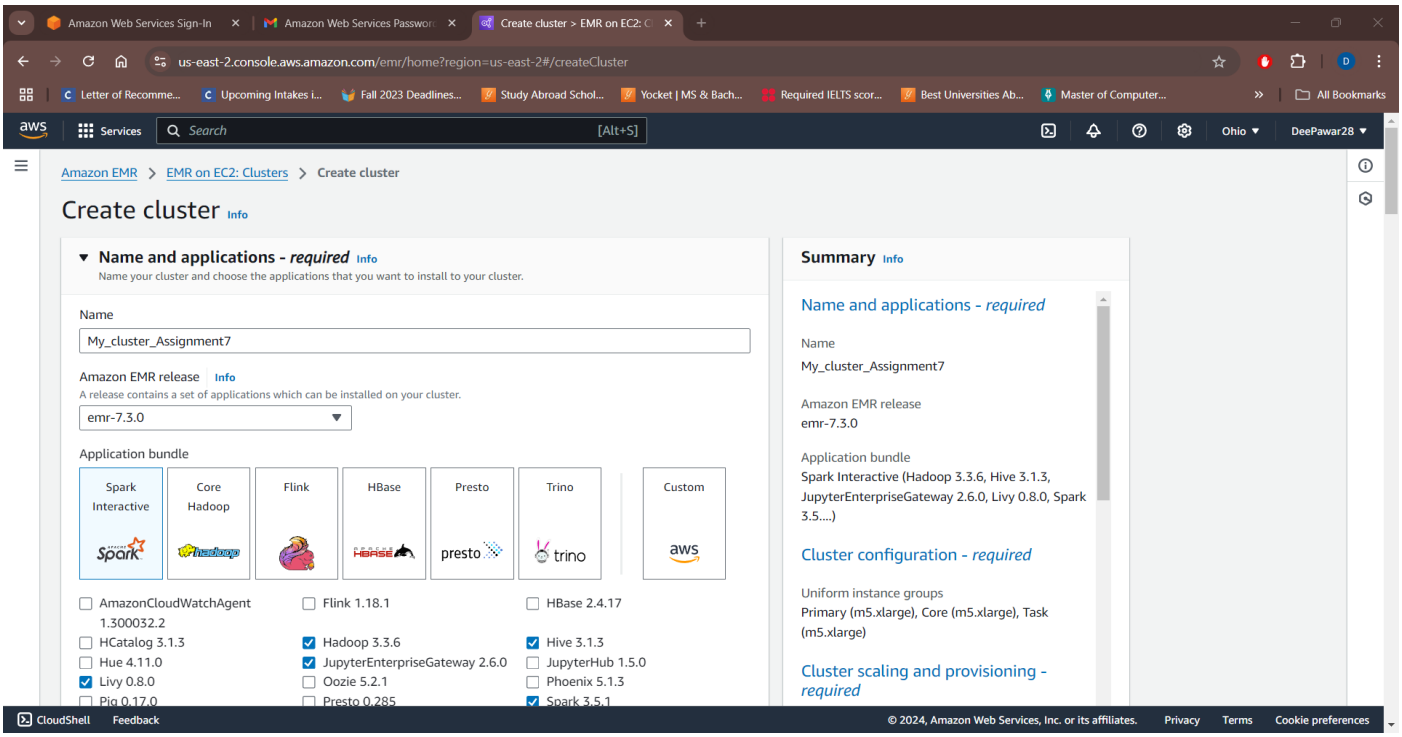

Step B: Use the TestDataGen program from previous assignments to generate new data files.

Copy both generated files to the HDFS directory "/user/hadoop"

Ans:

Magic Number: 232556

- **Command Used:**

java testDataGen

ls

hadoop fs -copyFromLocal /home/hadoop/foodplaces232556.txt

hadoop fs -copyFromLocal /home/hadoop/foodratings232556.txt

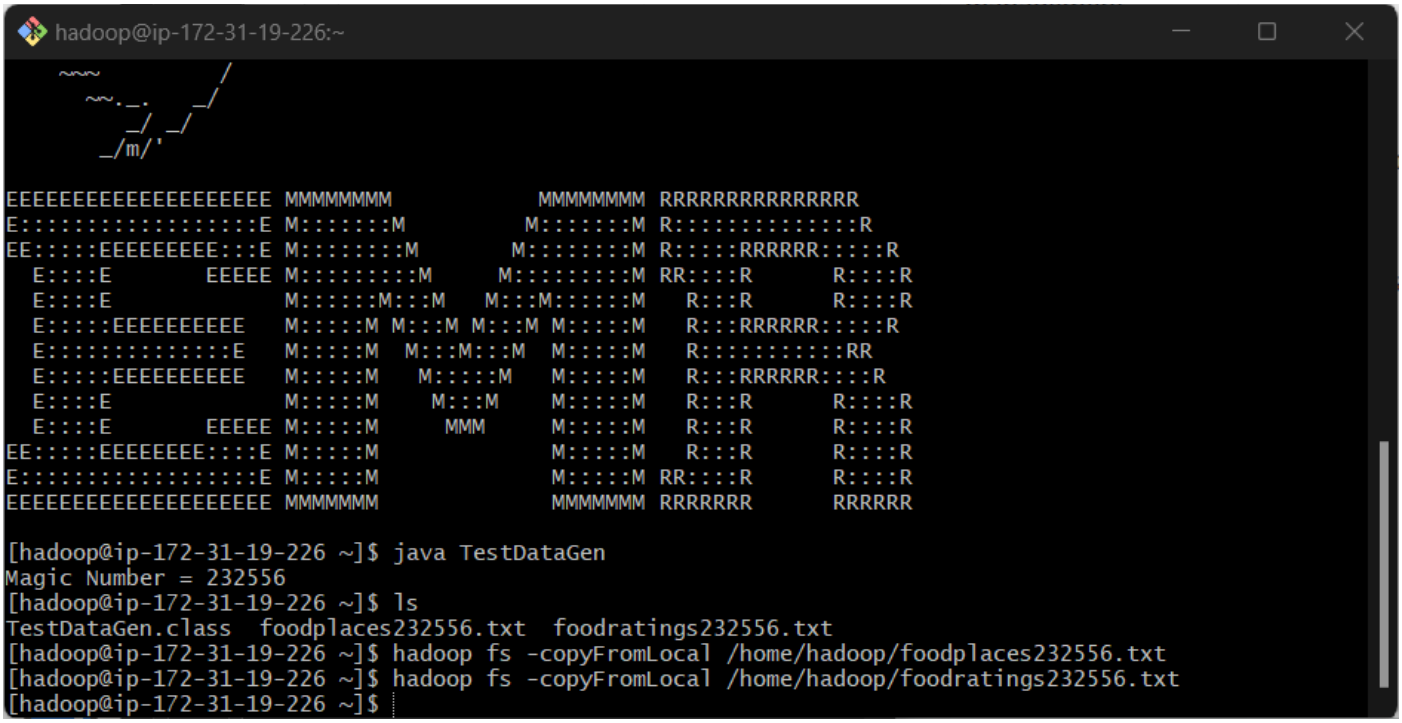

Step C: Load the 'foodratings' file as a 'csv' file into a DataFrame called foodratings.

Ans:

- **Command Used:**

pyspark

from pyspark.sql.types import *

struct1 = StructType().add("name", StringType(), True).add("food1",IntegerType(), True).add("food2",IntegerType(), True).add("food3",IntegerType(), True).add("food4",IntegerType(), True).add("placeid",IntegerType(), True)

foodratings = spark.read.schema(struct1).csv(' foodratings232556.txt')

foodratings.printSchema()

foodratings.show(5)

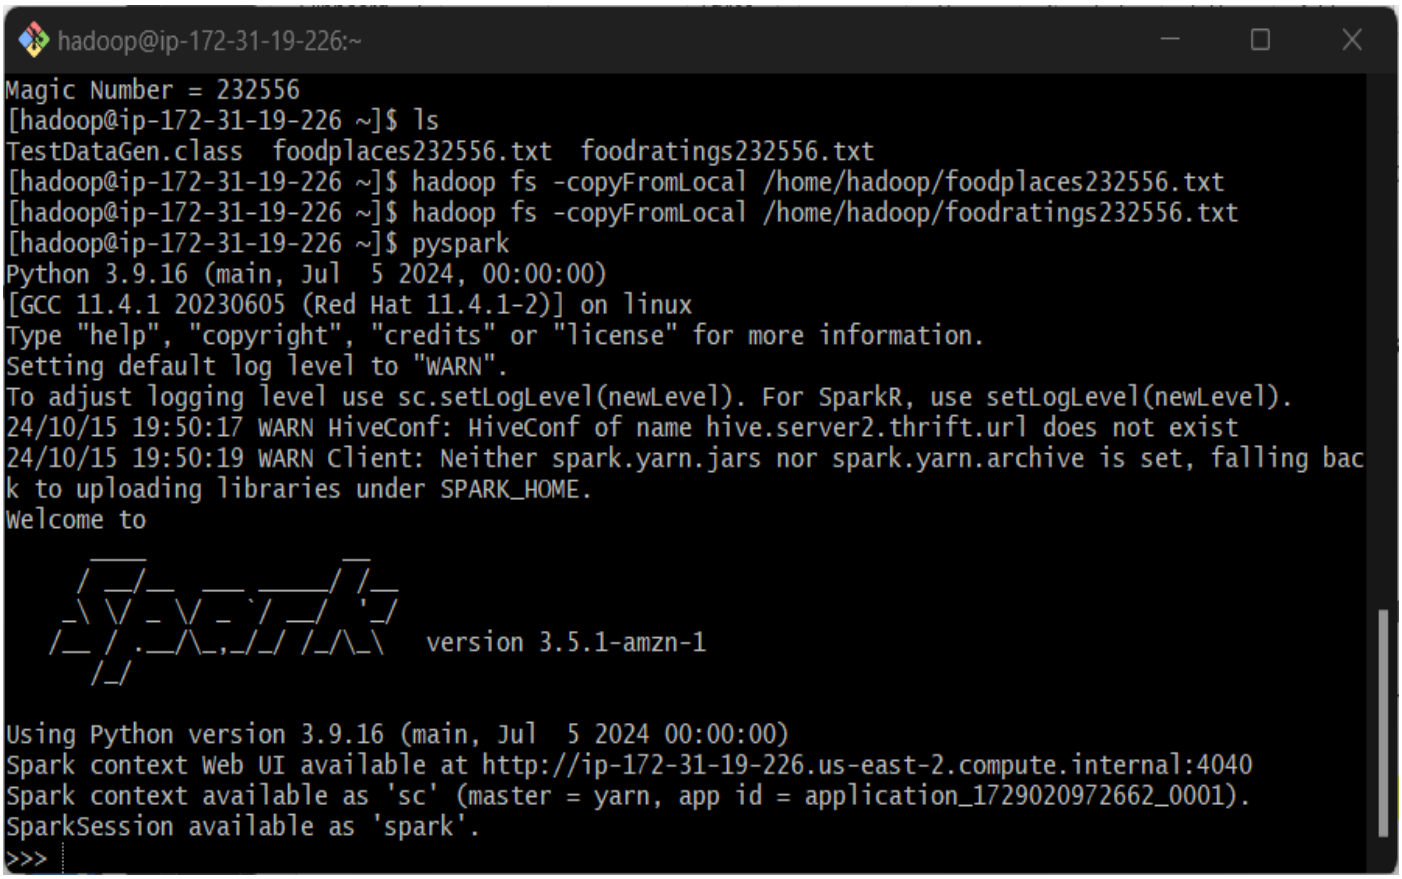
hadoop@ip-172-31-19-226:~
Magic Number = 232556
[hadoop@ip-172-31-19-226 ~]$ ls
TestDataGen.class foodplaces232556.txt foodratings232556.txt
[hadoop@ip-172-31-19-226 ~]$ hadoop fs -copyFromLocal /home/hadoop/foodplaces232556.txt
[hadoop@ip-172-31-19-226 ~]$ hadoop fs -copyFromLocal /home/hadoop/foodratings232556.txt
[hadoop@ip-172-31-19-226 ~]$ pyspark
Python 3.9.16 (main, Jul 5 2024, 00:00:00)
[GCC 11.4.1 20230605 (Red Hat 11.4.1-2)] on linux
Type "help", "copyright", "credits" or "license" for more information.
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
24/10/15 19:50:17 WARN HiveConf: HiveConf of name hive.server2.thrift.url does not exist
24/10/15 19:50:19 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling bac
k to uploading libraries under SPARK_HOME.
Welcome to
version 3.5.1-amzn-1
Using Python version 3.9.16 (main, Jul 5 2024 00:00:00)
Spark context Web UI available at http://ip-172-31-19-226.us-east-2.compute.internal:4040
Spark context available as 'sc' (master = yarn, app id = application_1729020972662_0001).
SparkSession available as 'spark'.
>>>

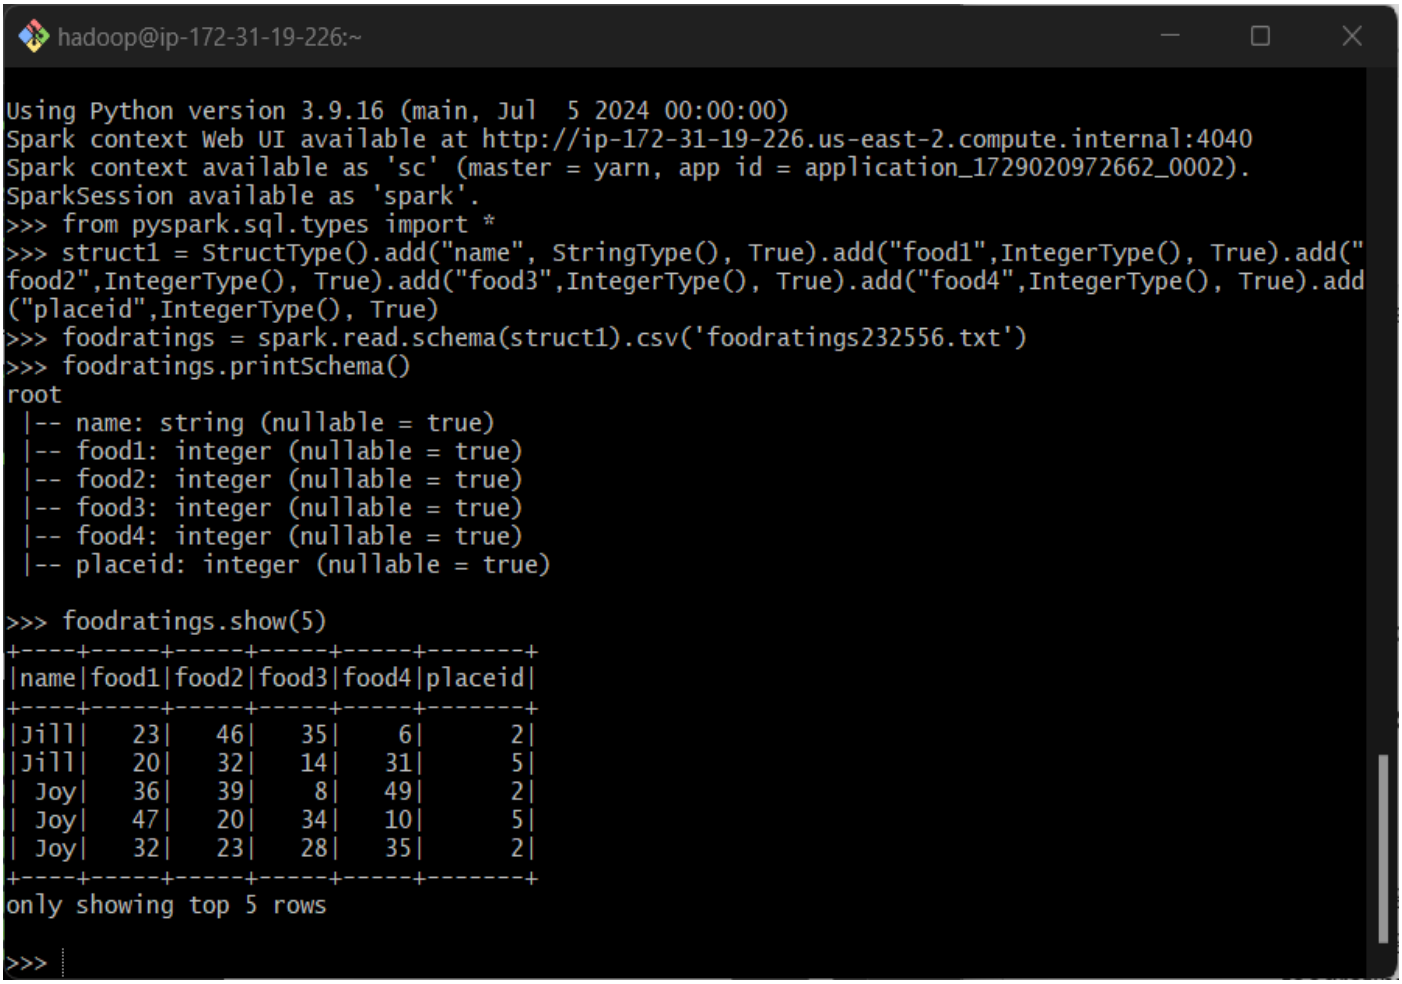
hadoop@ip-172-31-19-226:~
Using Python version 3.9.16 (main, Jul 5 2024 00:00:00)
Spark context Web UI available at http://ip-172-31-19-226.us-east-2.compute.internal:4040
Spark context available as 'sc' (master = yarn, app id = application_1729020972662_0002).
SparkSession available as 'spark'.
>>> from pyspark.sql.types import *
>>> struct1 = StructType().add("name", StringType(), True).add("food1",IntegerType(), True).add("
food2",IntegerType(), True).add("food3",IntegerType(), True).add("food4",IntegerType(), True).add
("placeid",IntegerType(), True)
>>> foodratings = spark.read.schema(struct1).csv('foodratings232556.txt')
>>> foodratings.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)
>>> foodratings.show(5)
+----+-----+-----+-----+-----+-------+
|name|food1|food2|food3|food4|placeid|
+----+-----+-----+-----+-----+-------+
|Jill| 23| 46| 35| 6| 2|
|Jill| 20| 32| 14| 31| 5|
Joy	36	39	8	49	2
Joy	47	20	34	10	5
Joy	32	23	28	35	2
+----+-----+-----+-----+-----+-------+
only showing top 5 rows
>>>

2. Exercise 2)

Load the ‘foodplaces’ file as a ‘csv’ file into a DataFrame called foodplaces.

Ans:

- **Command Used:**

```
from pyspark.sql.types import *
struct1 = StructType().add("placeid", IntegerType(), True).add("placename", StringType(), True)
foodplaces = spark.read.schema(struct1).csv('foodplaces232556.txt')
foodplaces.printSchema()
foodplaces.show(5)
```

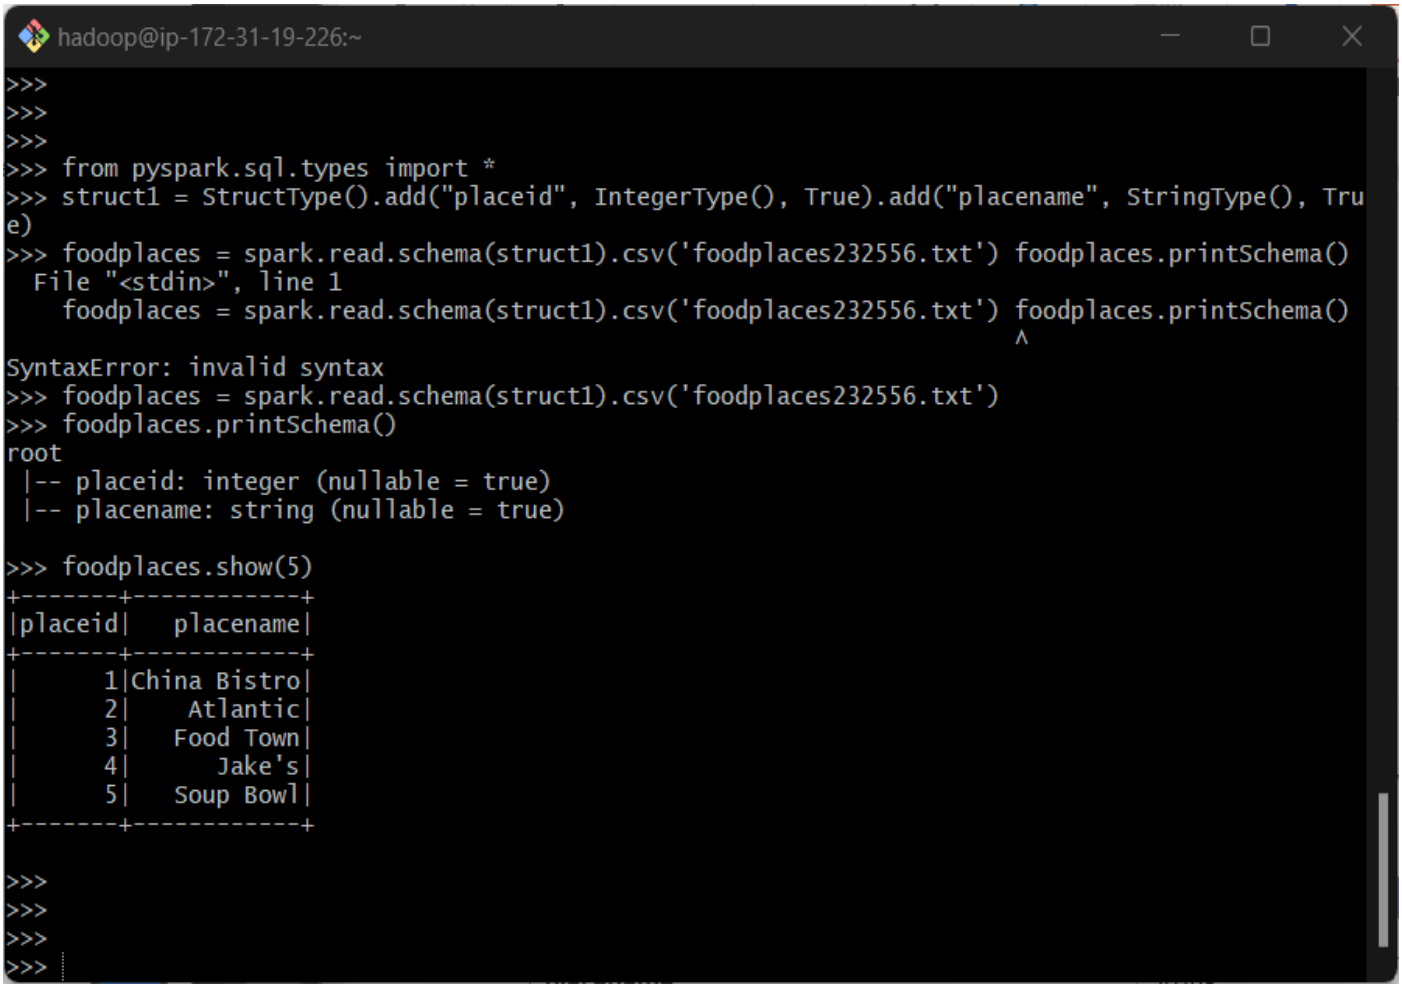

3. Exercise 3)

Step A: Register the DataFrames created in exercise 1 and 2 as tables called “foodratingsT” and “foodplacesT”

Ans:

- **Command Used:**

foodratings.createOrReplaceTempView("foodratingsT")
foodplaces.createOrReplaceTempView("foodplacesT")

```
>>>
>>>
>>>
>>> foodratings.createOrReplaceTempView("foodratingsT")
>>> foodplaces.createOrReplaceTempView("foodplacesT")
>>>
>>>
>>>
>>>
```

Step B: Use a SQL query on the table “foodratingsT” to create a new DataFrame called foodratings_ex3a holding records which meet the following condition: food2 < 25 and food4 > 40. Remember, when defining conditions in your code use maximum parentheses.

Ans:

- **Command Used:**

foodratings_ex3a = spark.sql("SELECT * from foodratingsT where food2 < 25 and food4 > 40")

foodratings_ex3a.printSchema()

foodratings_ex3a.show(5)

```
hadoop@ip-172-31-19-226:~
>>>
>>>
>>>
>>> foodratings_ex3a = spark.sql("SELECT * from foodratingsT where food2 < 25 and food4 > 40")
>>> foodratings_ex3a.printSchema()
root
 |-- name: string (nullable = true)
 |-- food1: integer (nullable = true)
 |-- food2: integer (nullable = true)
 |-- food3: integer (nullable = true)
 |-- food4: integer (nullable = true)
 |-- placeid: integer (nullable = true)

>>> foodratings_ex3a.show(5)
+----+-----+-----+-----+-----+-------+
|name|food1|food2|food3|food4|placeid|
+----+-----+-----+-----+-----+-------+
| Joe|   19|   14|   20|   46|      2|
| Sam|   17|    7|   44|   48|      3|
| Mel|   12|   17|   43|   45|      5|
| Joe|   13|   12|   26|   44|      4|
| Joy|   50|   22|   21|   41|      4|
+----+-----+-----+-----+-----+-------+
only showing top 5 rows

>>>
>>>
>>>
>>>
>>>
>>>
>>>
```

Step C: Use a SQL query on the table “foodplacesT” to create a new DataFrame called foodplaces_ex3b holding records which meet the following condition: placeid > 3

Ans:

- **Command Used:**

foodplaces_ex3b = spark.sql("SELECT * from foodplacesT where placeid> 3")

foodplaces_ex3b.printSchema()

foodplaces_ex3b.show(5)

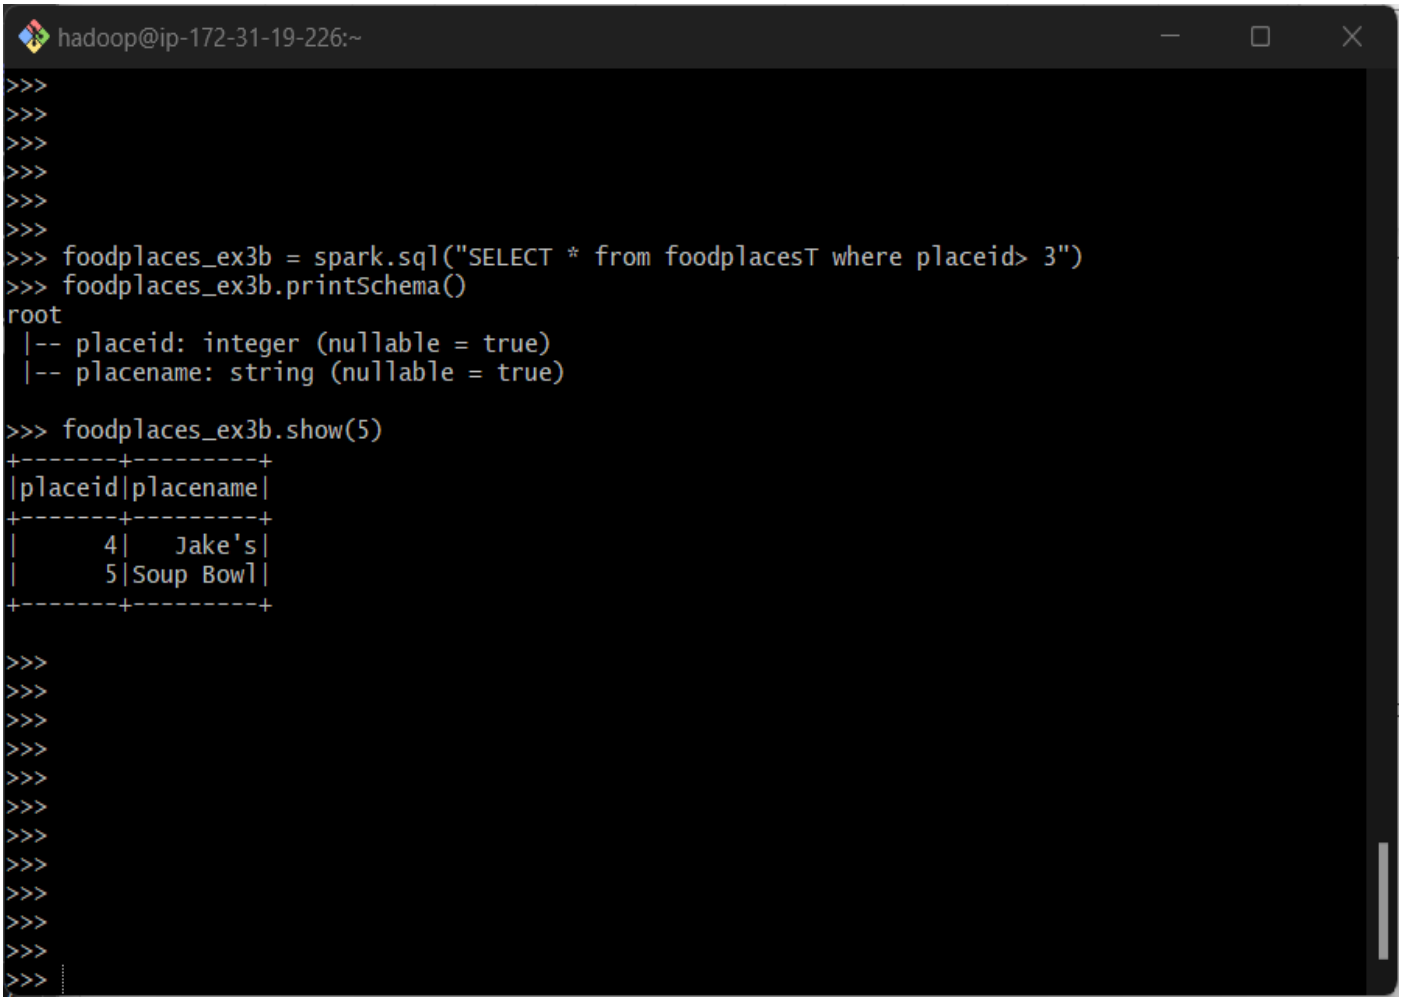

4. Exercise 4)

Use a transformation (not a SparkSQL query) on the DataFrame ‘foodratings’ created in exercise 1 to create a new DataFrame called foodratings_ex4 that includes only those records (rows) where the ‘name’ field is “Mel” and food3 < 25.

Ans:

- **Command Used:**

foodratings_ex4 = foodratings.filter(foodratings.name == "Mel").filter(foodratings.food3 < 25)

foodratings_ex4.printSchema()

foodratings_ex4.show(5)

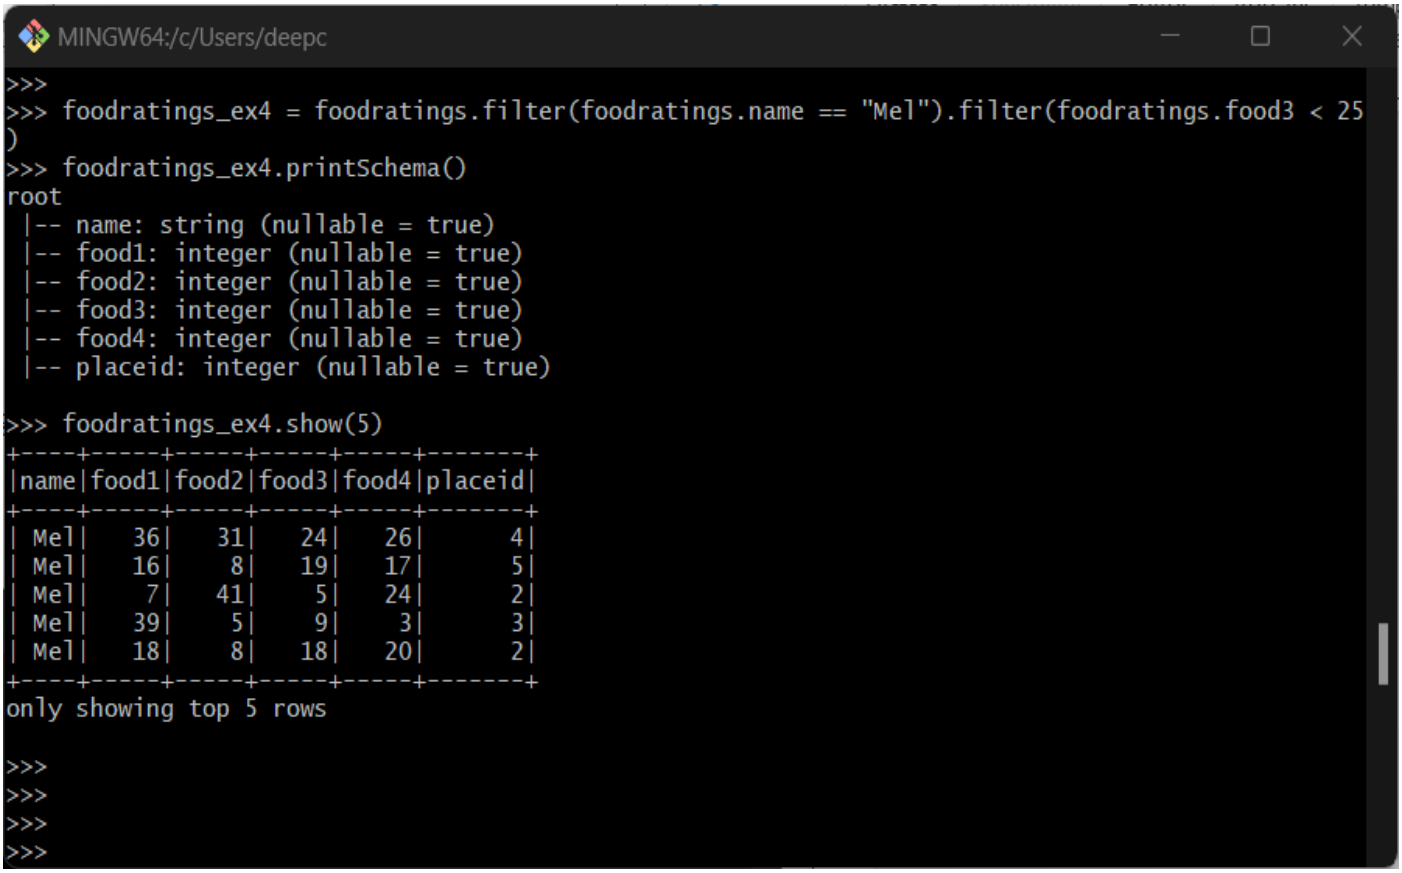

5. Exercise 5)

Use a transformation (**not a SparkSQL query**) on the DataFrame ‘foodratings’ created in exercise 1 to create a new DataFrame called foodratings_ex5 that includes only the columns (fields) ‘name’ and ‘placeid’

Ans:

- **Command Used:**

foodratings_ex5 = foodratings.select(foodratings.name, foodratings.placeid)

foodratings_ex5.printSchema()

foodratings_ex5.show(5)

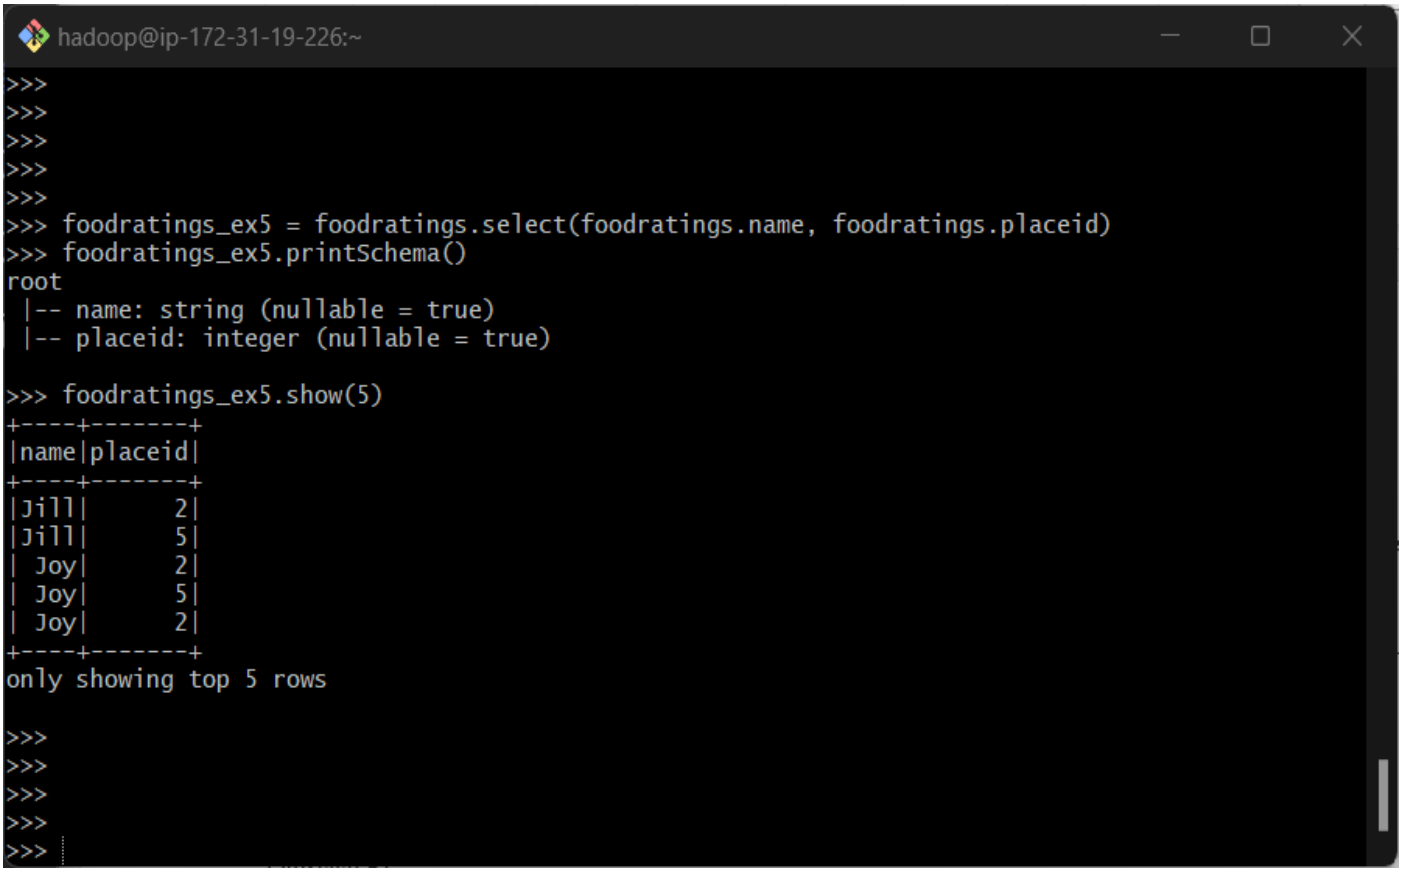

6. Exercise 6)

Use a transformation (**not a SparkSQL query**) to create a new DataFrame called ex6 which is the inner join, on placeid, of the DataFrames ‘foodratings’ and ‘foodplaces’ created in exercises 1 and 2

Ans:

- **Command Used:**

ex6 = foodratings.join(foodplaces, foodratings.placeid == foodplaces.placeid,"inner").drop(foodratings.placeid)

ex6.printSchema()

ex6.show(5)