

**Name:** Deep Pawar(A20545137)  
**Professor:** Joseph Rosen  
**Institute:** Illinois Institute of Technology

# CSP 554: Big Data Technologies

Fall 2024 - Assignment 10

## • Questions and Answers:

### Exercise 1) 5 points extra credit

Follow the document “Spark Streaming Demo Instructions” included with this assignment and execute the demo code. Provide enough screen shots to indicate you have completed the document through step 12. Then remember to terminate your VM.

- 1) Start up an EMR cluster as previously, but instead of choosing the “Core Hadoop” configuration choose the “Spark” configuration (see below), otherwise proceed as before.

The screenshot displays the AWS Management Console interface for an Amazon EMR cluster. At the top, a green notification bar states: "Your cluster 'Mycluster\_Assignment10' has been successfully created." The main header shows the cluster name "Mycluster\_Assignment10" and its status as "Waiting". Below this, a "Summary" section is expanded, showing four columns of information:

Cluster info	Applications	Cluster management	Status and time
<b>Cluster ID</b> j-14XQE1ZG13AYX	<b>Amazon EMR version</b> emr-7.4.0	<b>Log destination in Amazon S3</b> <a href="#">aws-logs-145023105604-us-east-2/elasticmapreduce</a>	<b>Status</b> Waiting
<b>Cluster configuration</b> Instance groups	<b>Installed applications</b> Hadoop 3.4.0, Hive 3.1.3, JupyterEnterpriseGateway 2.6.0, Livy 0.8.0, Spark 3.5.2	<b>Persistent application UIs</b> <a href="#">Spark History Server</a> <a href="#">YARN timeline server</a> <a href="#">Tez UI</a>	<b>Creation time</b> November 11, 2024, 18:31 (UTC-05:00)
<b>Capacity</b> 1 Primary 1 Core 0 Task		<b>Primary node public DNS</b> ec2-13-59-171-93.us-east-2.compute.amazonaws.com <a href="#">Connect to the Primary node using SSH</a> <a href="#">Connect to the Primary node using SSM</a>	<b>Elapsed time</b> 6 minutes, 57 seconds

Below the summary, a horizontal tab bar includes: Properties (selected), Bootstrap actions, Instances (Hardware), Steps, Applications, Configurations, Monitoring, Events, and Tags (1). The footer of the console shows "CloudShell", "Feedback", and copyright information for Amazon Web Services, Inc. or its affiliates.

4) There is one item you must change in consume.py. In the following line, you must replace <Master public DNS> with your own public DNS name (found as described above)

```
public class Utilities { Untitled-6 • consume.py X public class OrderItem { Untitled-5 • import java.io.IOException; Untitled-4 •
C: > Users > deepc > Downloads > consume.py
1 from pyspark import SparkContext
2 from pyspark.streaming import StreamingContext
3
4 # Create a local StreamingContext with a batch interval of 10 seconds
5 sc = SparkContext("yarn", "NetworkWordCount")
6 ssc = StreamingContext(sc, 10)
7
8 # Create a DStream
9 lines = ssc.socketTextStream("ec2-13-59-171-93.us-east-2.compute.amazonaws.com", 3333)
10
11 # Split each line into words
12 words = lines.flatMap(lambda line: line.split(" "))
13
14 # Count each word in each batch
15 pairs = words.map(lambda word: (word, 1))
16 wordCounts = pairs.reduceByKey(lambda x, y: x + y)
17
18 # Print each batch
19 wordCounts.pprint()
20
21 ssc.start() # Start the computation
22 ssc.awaitTermination() # Wait for the computation to terminate
```

5) scp this modified consume.py file to your EMR cluster primary (master) node. You may need to answer a security question with “Y/y” or “Yes”.

6) Then scp the file log4j.properties to your EMR cluster primary (master) node.

```
MINGW64:/c/Users/deepc
deepc@DeepPawar28 MINGW64 ~
$ scp -i C:/Users/deepc/Downloads/deep-emr-key-pair.pem C:/Users/deepc/Downloads/consume.py hadoop@ec2-13-59-171-93.us-east-2.compute.amazonaws.com:/home/hadoop/consume.py
100% 697 29.1KB/s 00:00

deepc@DeepPawar28 MINGW64 ~
$ scp -i C:/Users/deepc/Downloads/deep-emr-key-pair.pem C:/Users/deepc/Downloads/log4j.properties hadoop@ec2-13-59-171-93.us-east-2.compute.amazonaws.com:/home/hadoop/log4j.properties
100% 3199 132.5KB/s 00:00

deepc@DeepPawar28 MINGW64 ~
$
```

7) Open two terminal sessions to the EMR primary node. We will call one the EC2-1 window and the other the EC2-2 window.

- **EC2-1 window**

[illegible]

- **EC2-2 window**

[illegible]

```
sudo cp ./log4j.properties /etc/spark/conf/log4j.properties
```

9) In the EC2-1 window enter the following command to open a TCP (socket) connection on port 3333

```
hadoop@ip-172-31-23-44:~
#####_ Amazon Linux 2023
#####\
#####|
\#/
V~'-'>
https://aws.amazon.com/linux/amazon-linux-2023
-/m/'

EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRRR
E::::::::::::::::::E M::::::::M M::::::::M R:::::::::R
EE::::::::::::::::::E M::::::::M M::::::::M R:::::::::R
E::::E EEEEE M::::::::M M::::::::M RR::::R R::::R
E::::E M::::M:M:M M::M:M:M R::R R::::R
E::::EEEEEEEE M::::M M::M M::M M::::M R::RRRRRR::::R
E::::::::::::::::::E M::::M M::M:M:M M::::M R:::::::::RR
E::::EEEEEEEE M::::M M::M:M M::::M R::RRRRRR::::R
E::::E M::::M M::M M::::M R::R R::::R
E::::E EEEEE M::::M MMM M::::M R::R R::::R
EE::::::::::::::::::E M::::M M::::M R::R R::::R
E::::::::::::::::::E M::::M M::::M RR::::R R::::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRR RRRRRR

[hadoop@ip-172-31-23-44 ~]$ sudo cp ./log4j.properties /etc/spark/conf/log4j.properties
[hadoop@ip-172-31-23-44 ~]$ nc -lk 3333
```

10) In the EC2-2 window enter the following command:

spark-submit consume.py

This takes a while to start up. So, wait for some messages issued to the console before continuing. Note, when you do this you might see a message beginning with “WARN StreamingContext:...” which you can ignore.

```
hadoop@ip-172-31-23-44:~
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRRRRRRRRR
E::::::::::::::::::::E M::::::::M M::::::::M R:::::::::R
EE::::::::EEEEEEEEEE E M::::::::M M::::::::M R::::RRRRR::::R
E:::E EEEEE M::::::::M M::::::::M RR:::R R:::R
E:::E M::::::::M::M M::M::M R:::R R:::R
E::::EEEEEEEEEE M:::M M::M M::M M:::M R::RRRRR::::R
E::::::::::::E M:::M M::M::M M:::M R:::::::::RR
E::::EEEEEEEEEE M:::M M:::M M:::M R::RRRRR::::R
E:::E M:::M M::M M:::M R:::R R:::R
E:::E EEEEE M:::M MMM M:::M R:::R R:::R
EE::::EEEEEEEE::E M:::M M:::M R:::R R:::R
E::::::::::::E M:::M M:::M RR:::R R:::R
EEEEEEEEEEEEEEEEEEEE MMMMMMMM MMMMMMMM RRRRRRR RRRRRR

[hadoop@ip-172-31-23-44 ~]$ spark-submit consume.py
24/11/11 23:51:02 INFO EMRParamSideChannel: Setting FGAC mode to false
24/11/11 23:51:02 INFO SparkContext: Running Spark version 3.5.2-amzn-0
24/11/11 23:51:02 INFO SparkContext: OS info Linux, 6.1.112-122.189.amzn2023.x86_64, amd64
24/11/11 23:51:02 INFO SparkContext: Java version 17.0.12
24/11/11 23:51:03 INFO ResourceUtils: =====
=====
24/11/11 23:51:03 INFO ResourceUtils: No custom resources configured for spark.driver.
24/11/11 23:51:03 INFO ResourceUtils: =====
=====
24/11/11 23:51:03 INFO SparkContext: Submitted application: NetworkWordCount
24/11/11 23:51:03 INFO ResourceProfile: Default ResourceProfile created, executor resources: Map(executorType -> name: executorType, amount: 1, script: , vendor: , cores -> name: cores, amount: 4, script: , vendor: , memory -> name: memory, amount: 4743, script: , vendor: , offHeap -> name: offHeap, amount: 0, script: , vendor: ), task resources: Map(cpus -> name: cpus, amount: 1.0)
24/11/11 23:51:03 INFO ResourceProfile: Limiting resource is cpus at 4 tasks per executor
24/11/11 23:51:03 INFO ResourceProfileManager: Added ResourceProfile id: 0
24/11/11 23:51:03 INFO ResourceProfile: User executor ResourceProfile created, executor resources: Map(executorType -> name: executorType, amount: 1, script: , vendor: , cores -> name: cores, amount: 4, script: , vendor: , memory -> name: memory, amount: 4743, script: , vendor: , offHeap -> name: offHeap, amount: 0, script: , vendor: ), task resources: Map(cpus -> name: cpus, amount: 1.0)
```

11) Now in the EC2-1 window enter one or more lines of text and press Enter/Return after each one including the last. You should see the word count results scroll by in the EC2-2 window

- **EC2-1 window**

```
hadoop@ip-172-31-23-44:~  
https://aws.amazon.com/linux/amazon-linux-2023  
#####  
E::::::::::::::::::::E M::::::::M M::::::::M R::::::::::::R  
EE::::::::::::E M::::::::M M::::::::M R::::::::::::R  
E::::E EEEEE M::::::::M M::::::::M RR::::R R::::R  
E::::E M::::M M::M M::M M::::M R::R R::R  
E::::EEEEEEEE M::::M M::M M::M M::::M R::RRRRR::::R  
E::::::::::::E M::::M M::M M::M M::::M R::::::::RR  
E::::::::::::E M::::M M::M M::M M::::M R::RRRRR::::R  
E::::E M::::M M::M M::M M::::M R::R R::R  
E::::E EEEEE M::::M MMM M::::M R::R R::R  
EE::::::::::::E M::::M M::::M R::R R::R  
E::::::::::::E M::::M M::::M RR::::R R::::R  
EEEEEEEEEEEEEEEE MMMMMMM MMMMMMM RRRRRRR RRRRRR  
[hadoop@ip-172-31-23-44 ~]$ sudo cp ./log4j.properties /etc/spark/conf/log4j.properties  
[hadoop@ip-172-31-23-44 ~]$ nc -lk 3333  
I am Deep Pawar  
I am From Pune Maharashtra India  
Currently pursuing Masters in Computer Science  
This is Big Data Technologies
```

- **EC2-2 window**

```
hadoop@ip-172-31-23-44:~  
.internal:34913 in memory (size: 6.8 KiB, free: 4.8 GiB)  
24/11/11 23:53:00 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 42 (PythonRDD[83] at RDD  
at PythonRDD.scala:55) (first 15 tasks are for partitions Vector(1))  
24/11/11 23:53:00 INFO YarnScheduler: Adding task set 42.0 with 1 tasks resource profile 0  
24/11/11 23:53:00 INFO TaskSetManager: Starting task 0.0 in stage 42.0 (TID 91) (ip-172-31-26-83.us-east-  
2.compute.internal, executor 1, partition 1, NODE_LOCAL, 8977 bytes)  
24/11/11 23:53:00 INFO BlockManagerInfo: Removed broadcast_22_piece0 on ip-172-31-23-44.us-east-2.compute  
.internal:36845 in memory (size: 6.4 KiB, free: 1048.7 MiB)  
24/11/11 23:53:00 INFO BlockManagerInfo: Removed broadcast_22_piece0 on ip-172-31-26-83.us-east-2.compute  
.internal:34913 in memory (size: 6.4 KiB, free: 4.8 GiB)  
24/11/11 23:53:00 INFO BlockManagerInfo: Added broadcast_23_piece0 in memory on ip-172-31-26-83.us-east-2  
.compute.internal:34913 (size: 6.4 KiB, free: 4.8 GiB)  
24/11/11 23:53:00 INFO TaskSetManager: Finished task 0.0 in stage 42.0 (TID 91) in 73 ms on ip-172-31-26-  
83.us-east-2.compute.internal (executor 1) (1/1)  
24/11/11 23:53:00 INFO YarnScheduler: Removed TaskSet 42.0, whose tasks have all completed, from pool  
24/11/11 23:53:00 INFO DAGScheduler: ResultStage 42 (runJob at PythonRDD.scala:191) finished in 0.083 s  
24/11/11 23:53:00 INFO DAGScheduler: Job 21 is finished. Cancelling potential speculative or zombie tasks  
for this job  
24/11/11 23:53:00 INFO YarnScheduler: Killing all running tasks in stage 42: Stage finished  
24/11/11 23:53:00 INFO DAGScheduler: Job 21 finished: runJob at PythonRDD.scala:191, took 0.086174 s  
-----  
Time: 2024-11-11 23:53:00  
-----  
( 'am', 1)  
( 'Pawar', 1)  
( 'I', 1)  
( 'Deep', 1)  
24/11/11 23:53:00 INFO JobScheduler: Finished job streaming job 1731369180000 ms.0 from job set of time 1  
731369180000 ms  
24/11/11 23:53:00 INFO JobScheduler: Total delay: 0.383 s for time 1731369180000 ms (execution: 0.362 s)  
24/11/11 23:53:00 INFO PythonRDD: Removing RDD 73 from persistence list  
24/11/11 23:53:00 INFO BlockRDD: Removing RDD 68 from persistence list  
24/11/11 23:53:00 INFO BlockManager: Removing RDD 73  
24/11/11 23:53:00 INFO SocketInputStream: Removing blocks of RDD BlockRDD[68] at socketTextStream at Nat  
iveMethodAccessorImpl.java:0 of time 1731369180000 ms  
24/11/11 23:53:00 INFO ReceivedBlockTracker: Deleting batches: 1731369160000 ms  
24/11/11 23:53:00 INFO InputInfoTracker: remove old batch metadata: 1731369160000 ms  
24/11/11 23:53:00 INFO BlockManager: Removing RDD 68
```



```

hadoop@ip-172-31-23-44:~
24/11/11 23:57:50 INFO BlockManagerInfo: Removed broadcast_80_piece0 on ip-172-31-26-83.us-east-2.compute
.internal:34913 in memory (size: 6.8 KiB, free: 4.8 GiB)
24/11/11 23:57:50 INFO TaskSetManager: Starting task 0.0 in stage 158.0 (TID 150) (ip-172-31-26-83.us-eas
t-2.compute.internal, executor 1, partition 1, NODE_LOCAL, 8977 bytes)
24/11/11 23:57:50 INFO BlockManagerInfo: Removed broadcast_81_piece0 on ip-172-31-23-44.us-east-2.compute
.internal:36845 in memory (size: 6.4 KiB, free: 1048.7 MiB)
24/11/11 23:57:50 INFO BlockManagerInfo: Removed broadcast_81_piece0 on ip-172-31-26-83.us-east-2.compute
.internal:34913 in memory (size: 6.4 KiB, free: 4.8 GiB)
24/11/11 23:57:50 INFO BlockManagerInfo: Added broadcast_82_piece0 in memory on ip-172-31-26-83.us-east-2
.compute.internal:34913 (size: 6.4 KiB, free: 4.8 GiB)
24/11/11 23:57:50 INFO TaskSetManager: Finished task 0.0 in stage 158.0 (TID 150) in 74 ms on ip-172-31-2
6-83.us-east-2.compute.internal (executor 1) (1/1)
24/11/11 23:57:50 INFO YarnScheduler: Removed TaskSet 158.0, whose tasks have all completed, from pool
24/11/11 23:57:50 INFO DAGScheduler: ResultStage 158 (runJob at PythonRDD.scala:191) finished in 0.082 s
24/11/11 23:57:50 INFO DAGScheduler: Job 79 is finished. Cancelling potential speculative or zombie tasks
for this job
24/11/11 23:57:50 INFO YarnScheduler: Killing all running tasks in stage 158: Stage finished
24/11/11 23:57:50 INFO DAGScheduler: Job 79 finished: runJob at PythonRDD.scala:191, took 0.085046 s
-----
Time: 2024-11-11 23:57:50
-----
('am', 1)
('Maharashtra', 1)
('India', 1)
('I', 1)
('From', 1)
('Pune', 1)
24/11/11 23:57:50 INFO JobScheduler: Finished job streaming job 1731369470000 ms.0 from job set of time 1
731369470000 ms
24/11/11 23:57:50 INFO JobScheduler: Total delay: 0.310 s for time 1731369470000 ms (execution: 0.294 s)
24/11/11 23:57:50 INFO PythonRDD: Removing RDD 305 from persistence list
24/11/11 23:57:50 INFO BlockRDD: Removing RDD 300 from persistence list
24/11/11 23:57:50 INFO SocketInputDStream: Removing blocks of RDD BlockRDD[300] at socketTextStream at Na
tiveMethodAccessorImpl.java:0 of time 1731369470000 ms
24/11/11 23:57:50 INFO ReceivedBlockTracker: Deleting batches: 1731369450000 ms
24/11/11 23:57:50 INFO InputInfoTracker: remove old batch metadata: 1731369450000 ms
24/11/11 23:57:50 INFO BlockManager: Removing RDD 305
24/11/11 23:57:50 INFO BlockManager: Removing RDD 300

```

```

hadoop@ip-172-31-23-44:~
DD at PythonRDD.scala:55) (first 15 tasks are for partitions Vector(1))
24/11/11 23:58:50 INFO YarnScheduler: Adding task set 182.0 with 1 tasks resource profile 0
24/11/11 23:58:50 INFO TaskSetManager: Starting task 0.0 in stage 182.0 (TID 163) (ip-172-31-26-83.us-eas
t-2.compute.internal, executor 1, partition 1, NODE_LOCAL, 8977 bytes)
24/11/11 23:58:50 INFO BlockManagerInfo: Removed broadcast_94_piece0 on ip-172-31-23-44.us-east-2.compute
.internal:36845 in memory (size: 6.4 KiB, free: 1048.7 MiB)
24/11/11 23:58:50 INFO BlockManagerInfo: Removed broadcast_94_piece0 on ip-172-31-26-83.us-east-2.compute
.internal:34913 in memory (size: 6.4 KiB, free: 4.8 GiB)
24/11/11 23:58:50 INFO BlockManagerInfo: Added broadcast_95_piece0 in memory on ip-172-31-26-83.us-east-2
.compute.internal:34913 (size: 6.4 KiB, free: 4.8 GiB)
24/11/11 23:58:50 INFO TaskSetManager: Finished task 0.0 in stage 182.0 (TID 163) in 59 ms on ip-172-31-2
6-83.us-east-2.compute.internal (executor 1) (1/1)
24/11/11 23:58:50 INFO YarnScheduler: Removed TaskSet 182.0, whose tasks have all completed, from pool
24/11/11 23:58:50 INFO DAGScheduler: ResultStage 182 (runJob at PythonRDD.scala:191) finished in 0.068 s
24/11/11 23:58:50 INFO DAGScheduler: Job 91 is finished. Cancelling potential speculative or zombie tasks
for this job
24/11/11 23:58:50 INFO YarnScheduler: Killing all running tasks in stage 182: Stage finished
24/11/11 23:58:50 INFO DAGScheduler: Job 91 finished: runJob at PythonRDD.scala:191, took 0.069545 s
-----
Time: 2024-11-11 23:58:50
-----
('Currently', 1)
('in', 1)
('Science', 1)
('pursuing', 1)
('Masters', 1)
('Computer', 1)
24/11/11 23:58:50 INFO JobScheduler: Finished job streaming job 1731369530000 ms.0 from job set of time 1
731369530000 ms
24/11/11 23:58:50 INFO JobScheduler: Total delay: 0.272 s for time 1731369530000 ms (execution: 0.257 s)
24/11/11 23:58:50 INFO PythonRDD: Removing RDD 353 from persistence list
24/11/11 23:58:50 INFO BlockRDD: Removing RDD 348 from persistence list
24/11/11 23:58:50 INFO SocketInputDStream: Removing blocks of RDD BlockRDD[348] at socketTextStream at Na
tiveMethodAccessorImpl.java:0 of time 1731369530000 ms
24/11/11 23:58:50 INFO ReceivedBlockTracker: Deleting batches: 1731369510000 ms
24/11/11 23:58:50 INFO InputInfoTracker: remove old batch metadata: 1731369510000 ms
24/11/11 23:58:50 INFO BlockManager: Removing RDD 353
24/11/11 23:58:50 INFO BlockManager: Removing RDD 348

```

```

hadoop@ip-172-31-23-44:~
6.4 KiB, actual size: 6.4 KiB, free 1048.5 MiB)
24/11/12 00:00:20 INFO BlockManagerInfo: Added broadcast_114_piece0 in memory on ip-172-31-23-44.us-east-2.compute.internal:36845 (size: 6.4 KiB, free: 1048.7 MiB)
24/11/12 00:00:20 INFO SparkContext: Created broadcast 114 from broadcast at DAGScheduler.scala:1664
24/11/12 00:00:20 INFO DAGScheduler: Submitting 1 missing tasks from ResultStage 218 (PythonRDD[435] at RDD at PythonRDD.scala:55) (first 15 tasks are for partitions Vector(1))
24/11/12 00:00:20 INFO YarnScheduler: Adding task set 218.0 with 1 tasks resource profile 0
24/11/12 00:00:20 INFO TaskSetManager: Starting task 0.0 in stage 218.0 (TID 182) (ip-172-31-26-83.us-east-2.compute.internal, executor 1, partition 1, NODE_LOCAL, 8977 bytes)
24/11/12 00:00:20 INFO BlockManagerInfo: Added broadcast_114_piece0 in memory on ip-172-31-26-83.us-east-2.compute.internal:34913 (size: 6.4 KiB, free: 4.8 GiB)
24/11/12 00:00:20 INFO TaskSetManager: Finished task 0.0 in stage 218.0 (TID 182) in 62 ms on ip-172-31-26-83.us-east-2.compute.internal (executor 1) (1/1)
24/11/12 00:00:20 INFO YarnScheduler: Removed TaskSet 218.0, whose tasks have all completed, from pool
24/11/12 00:00:20 INFO DAGScheduler: ResultStage 218 (runJob at PythonRDD.scala:191) finished in 0.066 s
24/11/12 00:00:20 INFO DAGScheduler: Job 109 is finished. Cancelling potential speculative or zombie tasks for this job
24/11/12 00:00:20 INFO YarnScheduler: Killing all running tasks in stage 218: Stage finished
24/11/12 00:00:20 INFO DAGScheduler: Job 109 finished: runJob at PythonRDD.scala:191, took 0.068330 s
-----
Time: 2024-11-12 00:00:20
-----
('is', 1)
('This', 1)
('Big', 1)
('Data', 1)
('Technologies', 1)

24/11/12 00:00:20 INFO JobScheduler: Finished job streaming job 1731369620000 ms.0 from job set of time 1731369620000 ms
24/11/12 00:00:20 INFO JobScheduler: Total delay: 0.240 s for time 1731369620000 ms (execution: 0.225 s)
24/11/12 00:00:20 INFO PythonRDD: Removing RDD 425 from persistence list
24/11/12 00:00:20 INFO BlockRDD: Removing RDD 420 from persistence list
24/11/12 00:00:20 INFO BlockManager: Removing RDD 425
24/11/12 00:00:20 INFO SocketInputDStream: Removing blocks of RDD BlockRDD[420] at socketTextStream at NativeMethodAccessorImpl.java:0 of time 1731369620000 ms
24/11/12 00:00:20 INFO ReceivedBlockTracker: Deleting batches: 1731369600000 ms
24/11/12 00:00:20 INFO BlockManager: Removing RDD 420
24/11/12 00:00:20 INFO InputInfoTracker: remove old batch metadata: 1731369600000 ms

```

12) Remember to terminate your EMR instance after you are done!

The screenshot shows the AWS Management Console for an Amazon EMR cluster named 'Mycluster\_Assignment10'. The cluster is in a 'Terminated' state. The console displays various tabs like Summary, Properties, Bootstrap actions, Instances (Hardware), Steps, Applications, Configurations, Monitoring, Events, and Tags (1). The Summary tab is active, showing cluster info, applications, cluster management, and status and time.

Cluster info	Applications	Cluster management	Status and time
Cluster ID j-14XQE1ZG13AYX	Amazon EMR version emr-7.4.0	Log destination in Amazon S3 <a href="#">aws-logs-145023105604-us-east-2/elasticmapreduce</a>	Status Terminated
Cluster configuration Instance groups	Installed applications Hadoop 3.4.0, Hive 3.1.3, JupyterEnterpriseGateway 2.6.0, Livy 0.8.0, Spark 3.5.2	Persistent application UIs <a href="#">Spark History Server</a> <a href="#">YARN timeline server</a> <a href="#">Tez UI</a>	Creation time November 11, 2024, 18:31 (UTC-05:00)
Capacity 1 Primary 1 Core 0 Task		Primary node public DNS <a href="#">ec2-13-59-171-93.us-east-2.compute.amazonaws.com</a> <a href="#">Connect to the Primary node using SSH</a>	Elapsed time 29 minutes, 38 seconds
			End time November 11, 2024, 19:01 (UTC-05:00)

Operating system: Amazon Linux release  
Cluster logs: Archive log files to Amazon S3  
Cluster termination and node replacement: Edit